

Women’s Autonomy and Abortion Decision-Making*

Tom Zohar  Nina Brooks  Cauê Dobbin

September 23, 2026

Abstract

We study how abortion subsidies affect abortion take-up using administrative data covering the universe of legal abortions in Israel. Leveraging a reform that expanded eligibility for government funding of abortion, we find that the subsidy increased the share of pregnancies ending in abortion by 6.8 percentage points, with the effect concentrated among young women from backgrounds with strict views on abortion (14.8 percentage points). We propose a model of the family in which the daughter is liquidity constrained: she earns little at the time of pregnancy, so even a small price can put an abortion out of reach unless her parents pay for it. The subsidy therefore removes the need to rely on parental financial support. This changes behavior only when the daughter wants an abortion that her parents would refuse to finance, a mismatch that is more common in strict communities. Survey evidence on intergenerational mismatch in abortion attitudes supports this mechanism, and corroborating evidence from the United States points to its external validity. Our findings show that public funding expands young women’s autonomy over abortion decisions, placing abortion subsidies in the same policy space as parental consent laws.

*Brooks: School for Environment and Sustainability, University of Michigan, ninarb@umich.edu. Dobbin: Economics Department, Georgetown University, caue.dobbin@georgetown.edu. Zohar: Economics Department, CEMFI, tom.zohar@cemfi.es. We graciously acknowledge the Israeli Central Bureau of Statistics for providing access to the data. We are grateful to Theodore Joyce, Robert Kaestner, and Jason Ward for sharing replication data and code for their paper. We thank Ran Abramitzky, Pascaline Dupas, Liran Einav, Hedva Eyal, Nezih Guner, Caitlin Myers, Sharon Orshalimy, Petra Persson, and Isaac Sorkin, as well as participants at APPAM, ASSA, Stanford University Economics seminars, the 2021 NBER Summer Institute, and the 2023 BSE Summer Forum, for helpful comments. We thank Gioia Blayer, Giulia Camera, Daniel Fernandez, David Herskovits, Luca Natalucci, Moritz Osterhuber, Carles Pare Ogg, and Pol Vila Pinol for excellent research assistance. We gratefully acknowledge financial support from the Leonard W. Ely and Shirley R. Ely Graduate Student Fellowship, the Stanford Earth Dean’s Fellowship, the David and Lucile Packard SGF Fellowship, the Shultz Fellowship, the Freeman Spogli Institute for International Studies, Proyecto de Generación de Conocimiento 2021, and the Stanford Graduate Research Opportunity Grant for this work. All views and any errors are our own. Corresponding author: Tom Zohar, tom.zohar@cemfi.es.

1 Introduction

The legal status of abortion remains contested worldwide, and even where abortion is legal, women often face substantial barriers to access (Singh et al., 2018). A central policy question is whether the financial cost of the procedure meaningfully affects utilization. Standard economic reasoning suggests it should not: the cost of an abortion is small relative to the monetary cost of carrying a pregnancy to term and raising a child. Yet a large body of evidence documents that even modest out-of-pocket costs impede access.¹ Understanding why and for whom small financial barriers produce large behavioral responses is essential for evaluating whether abortion should be publicly subsidized.

This paper answers that question by overcoming common data limitations and empirical challenges in the abortion literature. We combine individual-level administrative data covering the universe of births and legal abortions in Israel, linked to rich demographic and economic information, and a reform that changed the price of abortion for a well-defined group of women. Before 2014, women aged 19 and younger were fully covered by government funding, while those aged 20 and above paid the full cost out of pocket; the reform extended full coverage to women aged 20 to 32. Our difference-in-differences design compares women who conceived just above and just below age 20, before and after the reform.

Two empirical patterns emerge. First, the subsidy increased the share of pregnancies ending in abortion by 6.8 percentage points. Second, the effect is concentrated among young women from ethno-religious backgrounds with strict abortion attitudes: 14.8 percentage points, compared with 5.6 percentage points among women from lenient backgrounds.² The large effects among women from strict backgrounds are robust across a series of validation exercises, whereas for women from lenient backgrounds we cannot distinguish a small effect from no effect at all. Together, these findings are difficult to reconcile with standard consumer-choice explanations. A frictionless price response would require implausibly large elasticities. Liquidity constraints could explain why the subsidy has an effect, but not why the effect is substantially larger in strict communities.

To explain these patterns, we build a two-period model of abortion decisions with parental altruism and liquidity constraints. A pregnant young woman values consumption over two periods and earns little today relative to her future income. If she obtains an abortion, she bears a direct utility cost that depends on her abortion attitudes and summarizes all costs and benefits of abortion relative to carrying the pregnancy to term. Her parents are altruistic toward her but hold their own attitudes

¹Health-policy research documents how financial costs prevent or delay access (e.g., Fried, 2000; Doran and Nancarrow, 2015; Dickman et al., 2022; Upadhyay, 2022; Grossman et al., 2016). In economics, the evidence comes primarily from expansions and restrictions of Medicaid funding for abortion in the United States (e.g., Blank et al., 1996; Levine et al., 1996; Cook et al., 1999).

²Israel exhibits substantial ethno-religious heterogeneity in abortion attitudes: secular Jews tend to hold lenient views, whereas religious Jews and Arabs generally hold stricter views. Behavior mirrors these attitudes: before the reform, 12.6% of pregnancies among lenient-background women ended in abortion, compared with 2.3% among strict-background women (Figure 1).

toward abortion. We study the subsidy's effects in three layers. First, with only a lifetime budget constraint and no liquidity restriction, the subsidy is a pure price change, and its effect is small because the procedure costs little relative to lifetime income. Second, adding a liquidity constraint makes the subsidy potentially powerful: most young pregnant women in our data do not have stable employment in the year before conception, so an out-of-pocket abortion may be infeasible despite ample lifetime resources. Third, we add family transfers. The daughter decides whether to disclose the pregnancy to her parents, who in turn decide whether to cover the abortion. If her parents would refuse coverage, she can obtain an abortion only by financing it herself or through the subsidy.

The model illuminates how autonomy and intergenerational mismatch shape the subsidy effect. A daughter has *autonomy* when the abortion outcome depends only on her own attitude, not on her parents'.³ By making the procedure free, the subsidy grants autonomy to daughters who cannot self-finance by removing the financial need to involve parents. This changes behavior only in *mismatch* families, where the daughter wants an abortion that her parents would refuse to finance. Families whose abortion outcome changes because of the subsidy, the compliers, are therefore defined by the conjunction of two factors: a subsidy-induced increase in autonomy and intergenerational mismatch.

We then take the model to the data and use it to understand the larger subsidy effects in strict communities. We show that the community-level effect can be decomposed into an *abortion pool*, the share of pregnancies that end in abortion if the procedure is free, and a *subsidy-induced share*, the fraction of that pool that only happens in the presence of the subsidy. Strict communities have a smaller pool (daughters are less likely to want an abortion) but a larger induced share (parents are more likely to oppose abortion and refuse coverage), making the subsidy instrumental. The total effect is therefore theoretically ambiguous; which force dominates is an empirical question.

The data confirm both margins, but the abortion pool is only about 20% smaller in strict communities, while the induced share is more than three times larger, generating the much larger total effect we estimate. We propose one explanation for this asymmetry: greater intergenerational mismatch in abortion attitudes. In lenient communities, parents and daughters hold similarly permissive views, so many daughters want an abortion but many parents are also willing to cover it, leaving few compliers. In strict communities, daughters are only somewhat more abortion-averse than daughters in lenient communities, so the abortion pool shrinks modestly, while parents are much more abortion-averse and more likely to refuse coverage. The subsidy therefore has a larger effect because it closes a wedge between daughters' demand for abortion and parents' willingness to finance it. We validate this mechanism using survey data showing that older and younger respondents hold similar abortion attitudes in secular Israeli communities but diverge sharply in religious ones. Calibrating the model to these survey

³Autonomy has no intrinsic utility value in the model; it is a conceptual tool for characterizing who controls the abortion decision and when that control affects outcomes.

data, we find that the observed intergenerational mismatch is large enough to account for the entire estimated difference in subsidy effects.

Our analysis reframes the welfare evaluation of publicly funded abortion for young women as a question of decision-making authority within the family. Because the financial cost of an abortion is small relative to the lifetime cost of raising a child, a standard comparison of the subsidy's fiscal cost with the welfare gains to recipients is unlikely to capture the policy's central trade-off. We show that subsidies allow young women to obtain an abortion without relying on parental financial support, shifting effective decision-making authority from parents toward daughters. The welfare question is therefore fundamentally about autonomy: who should decide whether a pregnancy continues, the young woman or her parents? Our framework does not take a position on this normative question, but it clarifies how the policy should be evaluated by placing abortion subsidies in the same conceptual space as parental consent laws, which have been extensively studied (e.g., Haas-Wilson, 1996; Kane and Staiger, 1996; Bitler and Zavodny, 2001; Joyce and Kaestner, 2001; Levine, 2003; Colman and Joyce, 2009; New, 2011; Myers, 2017; Joyce et al., 2020; Myers and Ladd, 2020; U.S. Government Accountability Office, 2025).

The connection to parental consent laws is reinforced by auxiliary evidence from the United States. We revisit two influential studies of parental involvement laws (Myers, 2017; Joyce et al., 2020) and show that their effects vary sharply with local religiosity, a proxy for abortion attitudes. These laws have significantly larger effects in more religious areas but have small effects in more secular ones, the mirror image of our subsidy results. We also find a parallel pattern in abortion attitudes: as in Israel, intergenerational mismatch is greater in more religious communities. Together, these findings suggest that the mechanisms of autonomy and intergenerational mismatch we identify extend beyond Israel and abortion subsidies to other policies that shift decision-making authority between young women and their parents.

Our work contributes to the literature on abortion funding (Blank et al., 1996; Kane and Staiger, 1996; Levine et al., 1996; Cook et al., 1999; Bitler and Zavodny, 2001; Hoehn-Velasco et al., 2025), where a recurring puzzle is why relatively small changes in out-of-pocket costs generate large behavioral responses. We tackle this puzzle using unusually rich data that combine individual-level information on abortion decisions and religiosity with intergenerational links. Together with our model, these data offer an explanation for this puzzle: the relevant constraint is often the young woman's lack of liquidity combined with her parents' unwillingness to finance an abortion. Economic barriers therefore operate through a social channel.

We also connect to the broader literature on barriers to abortion access, most directly to parental involvement laws discussed above, as well as to mandatory waiting periods (Altındağ and Joyce, 2022; Joyce and Kaestner, 2001; Lindo and Pineda-Torres, 2019; Bitler and Zavodny, 2001), provider restric-

tions and clinic closures (Colman and Joyce, 2011; Fischer et al., 2018; Lindo et al., 2020; Venator and Fletcher, 2021; Myers et al., 2025), and procedural requirements (Londoño-Vélez and Saravia, 2025). Our framework highlights which young women are most affected by financial barriers: those who cannot rely on family support because of intergenerational mismatch in abortion attitudes. This mechanism may also help explain the effects of other barriers, such as clinic distance, insofar as easier access reduces young women’s need to rely on parental support.

Our modeling approach builds on work that analyzes abortion decisions within broader settings of fertility, marriage, and child outcomes (e.g., Akerlof et al., 1996; Levine and Staiger, 2002, 2004; Ananat et al., 2009; Forsstrom, 2021). In this literature, the abortion decision is typically modeled as the woman’s decision alone. Forsstrom (2021) comes closest to our setting by incorporating parental involvement and stigma, but does not generate predictions for how policy effects vary with community attitudes. Our framework explains the empirical link between community attitudes and policy effects by embedding the abortion decision in a family problem in which altruistic parents hold their own abortion attitudes and the daughter can conceal the abortion from them.⁴ It also suggests that a central question for the normative evaluation of abortion subsidies is how decision-making authority should be allocated within the family.

The rest of the paper proceeds as follows. Section 2 describes the data and institutional context. Section 3 presents the empirical design and the main results. Section 4 develops the family model, and Section 5 takes it to the data. Section 6 presents the welfare analysis, and Section 7 concludes.

2 Data and Institutional Context

2.1 Data Sources

Conceptions (Abortions and Births): Our data on abortions come from administrative records of Israel’s national abortion committee and include every woman who applied for abortion approval between January 2009 and March 2016, including those who ultimately obtained abortions through private physicians. The data include each woman’s demographic characteristics, gestational age at application, and the official approval outcome (see further details on the abortion committee in Section 2.2).

We link these data to the 2017 civil registry, which provides information on all recorded live births, date of birth of the woman, and her marital status at time of conception. Together, the abortion and birth registries capture the universe of recorded pregnancies in Israel during our study period, excluding only illegal abortions and early miscarriages.⁵

⁴Stigma of the kind modeled in Forsstrom (2021) complements our mechanism. Appendix G extends our baseline model by making it costly for the daughter to disclose her pregnancy and shows that the main conclusions continue to hold.

⁵Abortions after 24 weeks of gestation are included in the abortion committee data, but these late-term cases are reviewed by a special committee and are rare (approximately 250 per year). We exclude them, along with stillbirths (defined as the

A key advantage of these data is that they allow us to measure abortions and births at the individual pregnancy level rather than relying on aggregate rates. When we report the share of pregnancies ending in abortion, we refer to this as the *abortion ratio*, to distinguish it from an *abortion rate* (abortions per 1,000 women of a given age). Both the abortion and birth records include gestational age, which we use to infer the date of conception and thereby construct a complete timeline for each observed pregnancy.

Ethno-Religious Classification: We combine data from the Census and the Ministry of Education to classify the ethno-religious background of women in our sample.

For Jewish women, we proxy religiosity using the type of secondary school they attended: secular (*mamlachti*), religious (*mamlachti-dati*), or Ultra-Orthodox (*haredi*). Although individual religiosity may evolve later in life, school type reflects parental religiosity at the time of schooling decisions and therefore provides a relevant measure of the social environment in which the woman was raised, including the fertility and abortion norms to which she was exposed.

Ethnicity is recorded when individuals receive their national identification card and is reported in the Census. We use this information to construct an indicator for whether an individual is of Arab ethnicity.

As explained below in Section 2.2, we classify religious Jews, Ultra-Orthodox Jews, and Arabs of any religion as having *strict* attitudes toward abortion. The remaining population—consisting mostly of secular Jews—is categorized as *lenient*. This classification reflects the substantially more permissive abortion attitudes among secular Jews relative to the other groups.

Socioeconomic Status (SES): To measure household economic resources, we use tax records on the young woman’s earnings and her parents’ earnings (father and mother), linked across generations using civil registry identifiers with a 96% match rate. Our main measure is the combined earnings of the woman’s parents in the year of conception (household earnings). We classify households as low- or high-SES based on whether their earnings fall below or above the median earnings in our sample.

Table A.1 shows that this classification captures sizable differences in economic resources: high-SES households earn nearly five times more than low-SES households. SES is also correlated with other important characteristics: high-SES daughters are almost twice as likely to have post-secondary education, and low-SES families are disproportionately represented among religious and Arab households.

Although high-SES young women earn more than their low-SES counterparts, the difference in their own earnings is modest (15%), compared to parental earnings that are 4.7 times higher. In both SES groups, daughters’ own earnings constitute only a small share of total household income, implying that

death of a fetus after 20 weeks of gestation; roughly 1,400 per year), from the analysis. Early spontaneous miscarriages are also not observed; however, because these events are not deliberate and are unlikely to vary differentially across the treated and comparison age groups, their omission may add noise but not bias.

young women in our sample depend on parental financial support regardless of SES. Indeed, research suggests that most women in our sample still live with their parents: Zionov (2021) reports that 84% of unmarried, childless women aged 18–24 live with their parents.

Education: We use data from the education registry spanning 2005–2016, which include information on vocational training programs and universities. The data cover both completed degrees and degrees that individuals started but had not yet completed.


2.2 Background on Abortion in Israel

Abortion Attitudes. Abortion attitudes and behaviors in Israel vary substantially across ethno-religious groups. The population is roughly 75% Jewish, 19% Muslim, and 2% Christian, with the Muslim and Christian populations consisting predominantly of Arabs (Central Bureau of Statistics (Israel), 2018). Judaism generally permits abortion to protect the mother’s life or health and, more broadly, when a child would be born into unstable circumstances (Amir, 2015). In practice, secular Jews—who constitute about 45% of the Jewish population—tend to hold permissive views, have high contraceptive use, and have relatively low fertility, whereas religious and Ultra-Orthodox Jews—roughly 30% combined—strongly oppose abortion, rarely use contraception, and have high fertility. Among Muslim- and Christian-Arab minorities, abortion is more heavily stigmatized and often considered taboo; Islam generally prohibits abortion except to save the mother’s life (Shapiro, 2014; Bhalotra et al., 2021), and studies of similar populations document strong social sanctions around abortion (Foster et al., 2007).

Due to these differences in abortion attitudes, we define the secular Jewish population as having *lenient* attitudes toward abortion and combine religious Jews, Ultra-Orthodox Jews, and Arabs of any religion into a single group with *strict* attitudes toward abortion.

These differences in abortion attitudes are reflected in the abortion ratios shown in Figure 1. In the lenient population, 12.6% of pregnancies end in abortion, compared with 2.3% in the strict population. Among our analytic sample of unmarried women—described in detail below in Section 3.1—abortion ratios are substantially higher, but the large gap between lenient and strict groups persists. This variation underscores that abortion access and decision-making in Israel take place within a highly segmented social and religious landscape (see Figure A.1 for a more detailed ethno-religious breakdown).

Figure 1: Abortion Ratios in Israel by Abortion Attitudes



Notes: The figure reports abortion ratios (the percentage of pregnancies that end in abortion) for two populations: all women aged 16–45 who conceived between 2009–2013 (on the left), and the analytic sample—defined in Section 3.1—of unmarried women aged 16–23 who conceived for the first time in the same period (on the right). Within each population, observations are grouped by attitudes toward abortion: “lenient” (secular Jews) and “strict” (religious Jews, Ultra-Orthodox Jews, and Israeli Arabs).

Public debate around abortion in Israel generates relatively little political polarization or media attention compared with other high-income countries, and legal reforms have historically encountered limited opposition from either secular or religious parties. As a result, the abortion system provides relatively easy legal access, but operates within the context of substantial normative heterogeneity across groups. Further details on the Israeli context are provided in Appendix B.1.

Abortion Committee. Abortion has been legal in Israel since 1977, contingent on approval from a committee composed of two medical professionals and a social worker. All abortions—public or private—require this approval. Social workers typically assist applicants throughout the process (Oberman, 2021), and, as a result, approval rates are extremely high: in our data, 99% of applications are approved, and 97% of approved applications result in an abortion.

The committee approves an abortion if at least one of the following conditions is met: (1) the woman is under 18 or over 40; (2) the pregnancy occurred outside marriage; (3) the pregnancy resulted from rape or incest; (4) the pregnancy endangers the woman’s life or health; or (5) the fetus has a congenital disorder (see approval shares in Table B.1a). Importantly, all unmarried women qualify automatically under the out-of-marriage criterion, while married women must meet another condition or report that the pregnancy is out of wedlock. More details on the committee process are provided in Appendix B.2.

Given the high approval rates and confidentiality of the process, incentives for illegal abortions are minimal. Anecdotal evidence suggests that some abortions occur outside the committee system (Oberman, 2021), but credible estimates on the quantity of these cases are unavailable.⁶ Moreover, in

⁶A newspaper article claimed that 15,000 illegal abortions occur annually in Israel (Newman, 2017), but we were unable

Section 3.3 we show that there is no evidence that the 2014 reform induced women who would have otherwise obtained an illegal abortion to instead obtain a legal one.

Abortion Cost. Once approved, women pay an out-of-pocket cost for an abortion, in contrast to most healthcare services in Israel. In 2014, the public-sector cost ranged from NIS 2,100–3,500 (USD 588–980), depending on gestational stage. Private procedures (which still require the committee’s approval) are faster but costlier, up to NIS 8,000 (USD 2,240). For comparison, average monthly earnings for working women were NIS 7,666 (USD 2,146), and for young women in our sample NIS 2,109 (USD 590).⁷

Policy Change: Eliminating the Cost of Abortion. Before 2014, only women aged 19 or younger and those meeting medical or legal criteria were eligible for fully subsidized abortions (see Table B.1a). According to abortion activists,⁸ women who were not eligible were often surprised to learn, upon arrival at the clinic, that they would need to pay between \$588 and \$980 for the procedure, since co-pays are typically rare (and small) in the Israeli healthcare system. Many women—especially religious women—struggled to secure this amount on short notice because they could not ask friends or family members for financial support for an abortion.

In January 2014, following advocacy efforts and a surplus in the national healthcare budget, the government expanded eligibility for full public funding (Amsterdamski et al., 2021). As shown in Table B.1b, the reform extended coverage from age 19 to all women up to age 32, including the committee fee (Kelner, 2013). Age was used as a proxy for income, with 32 as a budget-driven upper cutoff. After the reform, all women aged 32 or younger could obtain an abortion free of charge once approved, while the approval process itself remained unchanged. To the best of our knowledge, no other reproductive or income-related policies in Israel exhibit discontinuities at ages 19 or 32.

3 The Effect of Subsidies on Abortion Decisions

In this section, we present our main empirical analysis: estimates of the effects of the 2014 extension of abortion subsidies, with particular focus on heterogeneity by abortion attitudes. Section 3.1 defines the analytical sample. Section 3.2 presents the main results. Section 3.3 reports a series of robustness checks and shows that our findings remain stable.

to verify this figure: the organization cited in the article could not identify the underlying data, and the reporter did not respond to follow-up inquiries. A related concern is that ordering medication abortion pills online could provide a way to bypass the committee, but such practices were not widespread in Israel during the time period covered by our analysis (Oberman, 2021).

⁷See Appendix B.3 for details on childrearing costs in Israel.

⁸Interview conducted by the authors with Sharon Orshalimy, Israeli reproductive-justice activist and 2013 Young Leader with Women Deliver, Tel Aviv, Israel, July 2019.

3.1 Sample

Our main analytical sample consists of unmarried women aged 16–23 at the time of conception of their first pregnancy. We restrict the sample to *unmarried* women because pregnancies conceived within marriage can be aborted only for health reasons or if the woman is younger than 18, cases that were already subsidized prior to the 2014 reform.⁹ The policy therefore primarily affected unmarried women.¹⁰ We further restrict to first pregnancies to avoid endogeneity in marriage decisions: a prior abortion may alter subsequent marriage decisions, implying that the marital status of women with earlier pregnancies could itself be influenced by the subsidy. Limiting the sample to young women (16–23) ensures comparability between the treated (20–23) and control (16–19) groups. The sample period, 2009–2016, is chosen to avoid contamination from earlier policy changes. Appendix C provides additional details on these restrictions. After applying these restrictions, the final sample includes 40,495 pregnancies. Appendix D shows that our results are robust to alternative sample definitions.

3.2 Main Result: Effect Heterogeneity by Abortion Attitudes

To estimate the effect of the subsidy on abortion, we employ a difference-in-differences design that exploits both the timing of the 2014 Israeli policy change and the age cutoff at 19 years (see Table B.1b). Formally, we estimate the following regression:

$$abort_{it} = \delta (Post_t \times T_i) + \gamma_{a_i} + \gamma_{y_t} + \gamma_{m_t} + \beta X_{it} + \epsilon_{it} , \quad (1)$$

where the dependent variable $abort_{it}$ equals one if woman i , who conceived in month–year t , obtained an abortion. The year and month corresponding to t are denoted by y_t and m_t , respectively. The treatment indicator $T_i \equiv \mathbb{1}\{20 \leq age_i\}$ identifies women whose eligibility for the subsidy was affected by the reform, based on their age at conception, and $Post_t \equiv \mathbb{1}\{t \geq \text{Dec-2013}\}$ marks conceptions occurring after the policy took effect.¹¹ The coefficient δ captures the average effect of the subsidy on the probability of obtaining an abortion, conditional on pregnancy, expressed in percentage points.

We include age-at-conception fixed effects (γ_{a_i}) to control for systematic differences in abortion and fertility behavior across ages, conception-year fixed effects (γ_{y_t}), and conception-month fixed effects (γ_{m_t}) to account for seasonality in conception. The term X_{it} is a vector of individual controls with as-

⁹Married women can also get an abortion if they are older than 40, but those cases are excluded from our sample anyway because they are outside of our age range.

¹⁰Abortion is rare for married women: only 3.9% of pregnancies among married women aged 16–45 ended in abortion, in contrast to 32.6% among unmarried women. See Table B.1 for details on legal criteria for abortion and subsidy eligibility.

¹¹The policy was implemented in January 2014. However, women conceiving in late 2013 could apply for the subsidy if the abortion occurred in early 2014, when the reform was already in place. Because most legal abortions in Israel occur within the first eight weeks of pregnancy, we shift the treatment date one month earlier to include conceptions in December 2013 that were effectively treated under the new policy.

sociated coefficients β . In our preferred specification, X_{it} includes interactions between education and ethno-religious group fixed effects, which capture key determinants of abortion and fertility decisions (Kearney and Levine, 2012; Eckstein et al., 2019; Almond et al., 2019). Appendix D.1 discusses the rationale for these controls and examines alternative specifications. The error term is ϵ_{it} , and standard errors are clustered at the age-by-year level, corresponding to the level of treatment assignment (Abadie et al., 2022).

OLS estimates from Equation (1) are presented in Column (1) of Table 1. We find that the subsidy increased the abortion ratio by 6.8 percentage points, corresponding to 9.2% of the baseline value.

Table 1: Impact of the 2014 Subsidy Expansion on Abortion Utilization

	Full Sample	HH Socioeconomic Status		Daughter's Work Status		Abortion Attitudes	
		Low SES	High SES	Does Not Work	Works	Lenient	Strict
Treatment Effect	0.068*** (0.011) [-0.02, 0.04]	0.071*** (0.014) [-0.03, 0.07]	0.064*** (0.013) [-0.06, 0.03]	0.07*** (0.012) [-0.02, 0.05]	0.037** (0.016) [-0.11, 0.08]	0.056*** (0.01) [-0.04, 0.03]	0.148*** (0.027) [0.07, 0.29]
P-value		0.707		0.107		0.002	
Baseline Mean	0.738	0.658	0.830	0.746	0.719	0.773	0.509
N	40,495	21,558	17,030	29,439	11,056	35,231	5,264

Notes: This table reports OLS estimates of Equation 1, which identifies the effect of the 2014 reform on abortion ratios. The sample consists of all unmarried women aged 16–23 who became pregnant for the first time between 2009–2016. Columns (2) and (3) report estimates separately for women from low- and high-SES families; Columns (4) and (5) for women without and with stable employment in the year before conception, defined as own earnings above six months of the minimum wage; and Columns (6) and (7) for women from communities with lenient and strict attitudes toward abortion. P-values correspond to tests of statistical differences between Columns (2) and (3), between Columns (4) and (5), and between Columns (6) and (7). All regressions include month fixed effects and the interaction of ethno-religious-group and education fixed effects. Standard errors, clustered by age-by-year at conception, are in parentheses. Upper and lower bounds that account for differential pre-trends appear in brackets (details in Appendix D). * $p < 0.10$, ** $p < 0.05$, and *** $p < 0.01$.

Why did the subsidy lead to more abortion? If this result were driven primarily by a price effect or by easing liquidity constraints, we would expect a stronger response among low-SES women, who are typically more price sensitive and financially constrained. To test this, we estimate Equation (1) separately by household socioeconomic status (Columns (2) and (3)) and by the daughter's work status (Columns (4) and (5)). The results in Table 1 show that the difference in effects between women from low- and high-SES families is small and statistically insignificant, suggesting that household financial resources are not the binding constraint. The daughter's own income, reflected in her work status, matters more than household SES—the effect among young women who do not work is nearly twice the magnitude of the effect among those who do—though the difference is not significant at conventional levels.

An alternative mechanism is that the subsidy relaxes a constraint with both financial and social dimensions. Ineligible young women—who would have to pay up to \$980 for an abortion—may need to involve their parents to cover the cost. For women coming from families with strict attitudes toward

abortion, this may preclude having the abortion altogether. The subsidy therefore allows young women to obtain an abortion without involving their parents. To explore this hypothesis, we estimate Equation (1) separately for women from strict and lenient backgrounds. The results, shown in Columns (6) and (7) of Table 1, support this interpretation. Consistent with the subsidy allowing pregnant young women to bypass social frictions, the effects differ sharply by abortion attitudes. Among women from lenient backgrounds, the subsidy increased the abortion ratio by 5.6 percentage points (7.2% of the baseline). Among women from strict backgrounds, the effect was much larger—14.8 percentage points, or 29.1% of the baseline. The difference between the two groups is statistically significant ($p < 0.01$).¹²

A concern is that heterogeneity by abortion attitudes may instead reflect socioeconomic differences.¹³ To address this concern, Table A.2 reports estimates of Equation (1) separately by the interaction of abortion attitudes with household SES and with the daughter’s employment status. At both SES levels, the effects are large for strict families and small for lenient families. By employment status, the effect is large only among nonworking daughters from strict families. This pattern suggests that the subsidy matters only when both conditions hold: the daughter cannot finance the procedure herself, and her parents are unwilling to help. These splits yield small samples and imprecise estimates, so the results should be interpreted cautiously. Nonetheless, the patterns support an important role for social frictions.

We next turn to the identifying assumptions underlying our difference-in-differences design. Equation (1) relies on the standard parallel-trends assumption that, in the absence of the policy change, treatment and control groups would have evolved similarly. We assess this assumption by estimating the following regression:

$$abort_{it} = \sum_{\substack{k=2009 \\ k \neq 2013}}^{2016} \delta_k \mathbb{1}\{y_t = k\} \cdot T_i + \gamma_{a_i} + \gamma_{y_t} + \gamma_{m_t} + \beta X_{it} + \epsilon_{it}, \quad (2)$$

where the coefficients of interest, δ_k , trace the dynamic effects of the policy. For $k > 2013$, δ_k captures the policy’s impact in year k , while for $k < 2013$ it reflects potential pre-trends. The baseline year (2013) is normalized to zero, $\delta_{2013} = 0$.

We estimate Equation (2) separately for women from lenient and strict backgrounds, with results shown in Figure 2. Estimates for the full sample are shown in Figure A.2. Among women from lenient backgrounds, we observe modest differential trends between treatment and control groups both before


¹²Table A.3 shows that each strict subgroup (religious Jews, Ultra-Orthodox Jews, and Arabs) exhibits substantially larger effects than the lenient group. All differences are statistically significant except for the Ultra-Orthodox, where small sample size limits precision. Effects do not differ significantly across strict subgroups, supporting our decision to pool them to increase statistical power.

¹³In our analytic sample, 18% of low-SES women come from strict communities, compared with 8% of high-SES women (Table A.1).

and after the reform, making it difficult to determine whether the policy had no effect or only a small one. In contrast, among women from strict backgrounds, there is no evidence of differential pre-trends prior to the reform, followed by a large and sustained increase in abortion ratios for the treated group thereafter.

These results suggest that the large post-reform response among women from strict backgrounds is not driven by pre-trends, but rather reflects a causal effect of the subsidy.

Figure 2: Dynamic Difference-in-Differences Estimates of the Subsidy Effect



Notes: The figure plots the dynamic difference-in-differences coefficients from Equation 2, which measure changes in abortion ratios between treated women (ages 20–23) and control women (ages 16–19) from 2009–2016. Estimates are shown separately for communities with lenient and strict attitudes toward abortion. Each point corresponds to a coefficient δ_k ; 2013 is the omitted year. The vertical dashed line marks the 2014 reform. The sample includes all unmarried women aged 16–23 who conceived for the first time between 2009–2016. All regressions include month fixed effects and the interaction of ethno-religious-group and education fixed effects. Standard errors are clustered by age-by-year at conception, and shaded regions depict 95% confidence intervals.

The event-study estimates for the strict population suggest a gradual post-reform increase rather than an immediate jump. Such a pattern is consistent with gradual information diffusion, as more women learn about the subsidy. Information diffusion, however, is unlikely to explain why the subsidy effect is larger in strict communities. Doing so would require awareness of the subsidy to spread faster in strict than in lenient communities, whereas the setting suggests the opposite is more likely. Abortion is widely regarded as taboo in strict communities (Appendix B.1), so information about the subsidy would, if anything, be expected to diffuse more slowly there.¹⁴

Next, in Section 3.3, we examine the robustness of our findings.

¹⁴As an extreme example, “kosher phones” commonly used in Ultra-Orthodox communities restrict internet access and even text messaging, limiting information flows (Rosenberg et al., 2019).

3.3 Robustness

We address the main threats to identification here and leave the details, along with a further set of robustness exercises, to Appendix D. Because our main analysis conditions on being pregnant, a subsidy-induced change in contraceptive behavior—which we do not observe directly—would bias our estimates. Re-estimating Equation (1) on all unmarried women aged 16–23, with conception as the outcome, yields negative and statistically insignificant coefficients (Table D.1); the same null rules out substitution from unrecorded illegal abortions toward legal ones, which would have shown up as additional pregnancies (Appendix D.3). Treatment and comparison groups are balanced on most observables, with education the main exception (Table D.2), which is why our preferred specification includes ethno-religious-group-by-education fixed effects; results hold with and without them (Appendix D.1). Finally, the bounds in brackets in Table 1, which allow for differential trends estimated from pre-treatment data (Rambachan and Roth, 2023), exclude zero for the strict population, include zero for the lenient one, and do not overlap (Appendix D.2).

The findings are also robust to alternative specification and sample choices: narrower age bandwidths (Appendix D.4), samples not restricted to unmarried women (Appendix D.5), the exclusion of former-USSR migrants, whose historical relationship with abortion differs (Appendix D.6), alternative sets of controls (Tables D.3, D.4, and D.5), and the reason recorded for approval (Appendix D.7). Across all of these exercises, the subsidy has a large and statistically significant positive effect among women from strict backgrounds; for women from lenient backgrounds the evidence is less consistent, and we cannot distinguish no effect from a modest positive one.

3.4 External Validity

To assess the generalizability of our findings, we extend two existing U.S. studies of parental involvement in abortion decisions. These exercises show that the relationship between community attitudes toward abortion and the effects of policies that increase abortion autonomy is not unique to Israel.¹⁵

First, we extend Myers (2017), who studies policies granting minors confidential access to abortion across U.S. states. Using the original replication data and specification, we split states by abortion attitudes and re-estimate the effect separately for strict and lenient states.¹⁶ As shown in the first two columns of Table 2, confidential access reduces teenage births substantially more in strict states: 58.9% relative to the baseline rate, compared with 30.3% in lenient states.

¹⁵The U.S. studies use location-level data. For comparability, we also reproduce our main Israeli finding using location-level aggregates: Table A.4 shows that the subsidy effect is smaller in more lenient locations.

¹⁶We measure strictness using the share of the population adhering to Christian denominations with the most conservative views on abortion, based on the 1971 Survey of Churches and Church Membership (Glenmary Research Center, 1974). These denominations are the Church of Jesus Christ of Latter-Day Saints, the Church of the Nazarene, the Southern Baptist Convention, the Catholic Church, the Church of God, and the United Methodist Church. States above the U.S. median are classified as “strict.”

Table 2: U.S. Evidence on Heterogeneous Autonomy Effects by Abortion Attitudes

	Myers (2017)		Joyce et al. (2020)	
	Lenient	Strict	Lenient	Strict
Treatment Effect	-0.046*** (0.010)	-0.089*** (0.011)	-0.673 (0.841)	-3.526*** (0.547)
Baseline Mean	0.152	0.151	14.100	10.100
P-value	0.004		0.004	
N	152,279	124,233	514	600

Notes: Columns (1)–(2) extend Myers (2017), estimating the effect of confidential abortion access for minors on giving birth before age 19. Columns (3)–(4) extend Joyce et al. (2020), estimating the effect of parental involvement laws on abortion incidence among 15–17 year-olds. States are classified as strict or lenient based on pre-policy religious composition, as described in the text. All specifications follow the original studies. Robust standard errors are in parentheses. P-values test equality of the estimates in Columns (1) and (2), and Columns (3) and (4). * $p < 0.10$, ** $p < 0.05$, and *** $p < 0.01$.

Second, we similarly extend Joyce et al. (2020), who study the introduction of parental involvement laws and abortion incidence among minors.¹⁷ The last two columns of Table 2 show little effect of parental involvement laws in lenient states, but a large reduction in abortion incidence in strict states, equivalent to 34.9% of the baseline rate.

Together, these extensions show that policies changing young women’s autonomy over abortion decisions have larger effects in communities with stricter abortion attitudes, consistent with our findings for Israel.

4 Model: Abortion Decisions with Conflicting Attitudes

This section studies the abortion decision within a single family: a pregnant young woman and her parents, who may disagree about whether she should carry the pregnancy to term. Our goal is to understand how an abortion subsidy affects this decision, whose behavior it changes, and how these effects relate to autonomy. We build the family’s problem in three layers, isolating a price channel, a liquidity channel, and a family-transfer channel. The family-level objects we characterize here provide the inputs to the population analysis in Section 5, which aggregates across families to derive the subsidy’s community-level effects.

Appendix E presents the same model in more formal detail, and Appendix F collects the proofs.

¹⁷We use the authors’ replication data and replicate Model C from their Table 2 using CDC abortion-occurrence data for 15–17 year-olds. We classify states as strict or lenient using the same procedure as above, but based on the 1980 Survey of Churches and Church Membership (Glenmary Research Center, 1982).

4.1 Setup

Agents and timing. We consider a family consisting of a young woman—the *daughter*, indexed by D —and her *parents*, indexed by P . The daughter lives for two periods, $t = 1, 2$, and earns exogenous incomes Y_1^D and Y_2^D , with $Y_1^D < Y_2^D$. Her parents act only in period 1 and earn income Y^P . Within the family, abortion attitudes (ξ^D, ξ^P) are fixed and commonly known. With probability $\pi \in (0, 1)$, the daughter becomes pregnant in period 1; let $E \in \{0, 1\}$ indicate pregnancy. She privately observes E and chooses whether to obtain an abortion, $A \in \{0, 1\}$, where $A = 1$ denotes an abortion. The model studies the family’s problem conditional on this pregnancy realization.¹⁸

The abortion procedure has a monetary cost $p > 0$. Let $S \in \{0, 1\}$ indicate whether the government subsidizes the procedure. We focus on the comparative static from no subsidy ($S = 0$) to a full subsidy ($S = 1$), which reduces the out-of-pocket cost from p to zero.

The decision unfolds in four stages. First, before the daughter’s pregnancy status is known, parents choose an allowance T . We assume that they choose this allowance as if the daughter were not pregnant. Second, a pregnancy is realized; the daughter privately observes whether she is pregnant. Third, if pregnant, the daughter decides whether to reveal her pregnancy to her parents, generating the HIDE and TELL branches of the game; let $R \in \{0, 1\}$ indicate revelation. Fourth, in the TELL branch parents decide whether to cover the abortion cost, denoted $k \in \{0, 1\}$, and the daughter then chooses consumption and whether to obtain the procedure.

Preferences. The daughter values consumption, (C_1^D, C_2^D) , and bears a cost ξ^D when she obtains an abortion. This term is the *net* cost of abortion relative to carrying the pregnancy to term. It captures non-monetary considerations—moral, religious, or physical—as well as the monetary costs of child-rearing, including the daughter’s anticipation of parental support absent an abortion.¹⁹ Accordingly, ξ^D may be positive or negative. Formally, the daughter’s utility is

$$V^D = u(C_1^D) + u(C_2^D) - \xi^D \cdot A, \quad (3)$$

where the flow utility u is strictly increasing, strictly concave, and twice continuously differentiable, with $u'(0^+) = \infty$.²⁰

¹⁸Period 1 should be read as the short window—about two months—over which a pregnant woman must decide on and finance an abortion. Period 2 is the rest of her economic life.

¹⁹Disclosure carries no direct cost, so a daughter whose parents would not cover the procedure is indifferent between hiding and telling. We break this tie by the sign of parental attitudes: she tells her parents when they weakly approve of the abortion ($\xi^P \leq 0$) and hides when they oppose it (Appendix E.4). Appendix G shows this is the strict best response once involving one’s parents carries a cost proportional to their abortion distaste, and that the results in Section 4.3 are unchanged.

²⁰A non-pregnant daughter trivially has $A = 0$ and, taking Equation (3) literally, would have the same utility as a daughter who carries the pregnancy to term, even though she bears no child-related costs. Because pregnancy is exogenous in the model, adding a constant to the non-pregnant payoff to capture the absence of these costs would not affect any decision. We omit this constant for notational brevity.

Parental preferences are analogous, with parental consumption C^P and a *net* abortion cost ξ^P . This term bundles both non-monetary considerations (moral, religious, or normative) and the financial side of the no-abortion alternative; in particular, it nets out the parents' anticipated post-birth contribution to child-rearing. We do not give parents a separate child-rearing transfer instrument: because such transfers cannot be credibly committed *ex ante*, the daughter's anticipation of them is folded into ξ^D and the parents' anticipated cost is folded into ξ^P . Parents bear this cost only if the daughter obtains an abortion *and* discloses it; if she hides and self-finances, parents remain uninformed and do not directly experience it.²¹

Parents are altruistic toward their daughter. We adopt a *paternalistic-altruism* formulation (Pollak, 1988; Doepke and Zilibotti, 2017): parents value their daughter's consumption but not the abortion cost ξ^D that she would bear, reflecting the possibility that parents disagree with the daughter about what an abortion costs her. Any such cost that parents believe she would bear is instead folded into their own abortion distaste ξ^P .²² Parents' utility is

$$V^P = u(C^P) + \alpha^P \cdot [u(C_1^D) + u(C_2^D)] - \xi^P \cdot A \cdot R, \quad (4)$$

where $\alpha^P > 0$ measures parental altruism and the product $A \cdot R$ encodes that abortion distaste is payoff-relevant to parents only on disclosed paths.

Parental transfers. Parents can transfer resources to their daughter in two distinct ways: an *allowance* and *abortion coverage*.

The allowance represents ordinary intra-family consumption sharing that would be in place even absent the pregnancy. We denote it by $T \geq 0$: parents choose T freely but cannot extract resources from the daughter. We take T to be chosen before pregnancy is realized and as if the daughter were not pregnant. This timing rules out strategic reductions in ordinary transfers intended to limit hidden self-funding. Because pregnancy affects only 4.7% of the relevant population, however, allowing parents to account for pregnancy risk would likely yield an allowance close to the no-pregnancy benchmark.

A fixed share $\phi \in [0, 1]$ of the allowance is provided in kind. The daughter must consume this share, ϕT , in period 1. In practice, this in-kind component may arise through earmarked payments, such as tuition, rent, or meal plans, or through direct consumption sharing. The remaining share, $1 - \phi$, is fungible cash. We treat ϕ as a fixed parameter.

²¹We could add a separate child-rearing support cost to the parental payoff when the daughter does not abort. We omit it because it is constant across the parental decisions studied here and therefore does not affect behavior. Thus, Equation (4) normalizes the no-abortion payoff as the baseline, and ξ^P measures the parent's net cost of a disclosed abortion relative to that baseline.

²²Belley and Lochner (2007, Appendix B) adopt the same structure in the schooling context: parents value their own consumption and their child's consumption utility, hold their own psychic valuation of the child's schooling choice, and place some weight in $[0, 1]$ on the child's own taste for it. Equation (4) sets that weight to zero.

In the TELL branch, parents may additionally choose whether to cover the abortion cost, $k \in \{0, 1\}$. When $k = 1$, parents pay the out-of-pocket cost, $(1 - S)p$, and the daughter bears no direct cost. When $k = 0$, the daughter bears the full out-of-pocket cost. We refer to k as *parental coverage*. Unlike the allowance, coverage is chosen *after* parents observe the pregnancy through the daughter's disclosure.

Importantly, the daughter can use the fungible share of the transfer to finance an abortion, even if parents choose not to provide coverage ($k = 0$). The in-kind share, by contrast, cannot be redirected in this way. The composition of parental transfers therefore influences whether a parental refusal to cover the procedure can, in practice, prevent the abortion from occurring.

4.2 Three Layers of the Subsidy's Effect

We now present the daughter's problem in three layers to isolate the distinct forces through which the subsidy affects abortion decisions. The layers are separate models, not sequential stages of a single model: they share the same primitives but differ in the budget set available to the daughter. Layer 1 gives the daughter a lifetime budget; Layer 2 adds a short-run liquidity constraint; and Layer 3 adds family transfers. In all three layers, the daughter chooses consumption and abortion to maximize her utility V^D , defined in Equation (3). Note that the four-stage timing described in Section 4.1 is fully realized only in Layer 3, since Layers 1 and 2 have no parental transfers and therefore no allowance, disclosure, or coverage decisions.

In each layer, we sort daughters into three groups, defined by their behavior with and without the subsidy. A daughter is a *never-taker* if she does not abort even when the procedure is free ($S = 1$); an *always-taker* if she aborts even at the full price ($S = 0$); and a *complier* if she aborts only when subsidized. The subsidy changes the abortion decision of—and only of—compliers. The sorting is governed by the daughter's abortion cost ξ^D ; in Layer 3 the parents' attitude ξ^P enters as well.

The three layers differ in who belongs to the complier group. Layer 1 isolates a pure price effect: compliers are daughters who decline the procedure at the full price but obtain it when it is free—they are deterred by the price itself. Layer 2 widens this group by adding daughters who are not deterred by the price per se: they would pay the full price if they could, but cannot because they lack liquidity in period 1. Layer 3 then narrows the complier group by removing daughters who can obtain an abortion through parental support.

Layer 1: lifetime budget. The daughter freely smooths consumption across periods, facing only a single lifetime budget:

$$C_1^D + C_2^D + p(1 - S)A \leq Y_1^D + Y_2^D.$$

This frictionless benchmark isolates a pure price effect: by lowering the effective price of the procedure, the subsidy raises the number of women who abort, a standard demand response to a lower price. The abortion decision is governed by two thresholds, illustrated in Figure 3, Panel (a). Because a free abortion ($S = 1$) carries no monetary cost, the daughter takes it whenever her net distaste is weakly negative; the never-taker boundary is therefore $\bar{\xi}_1^D = 0$. She self-funds at the full price only if her distaste is sufficiently negative to outweigh the utility cost of paying p , which defines an always-taker boundary $\underline{\xi}_1^D < 0$. Compliers are the daughters in between.

This price effect is likely to be small. The procedure costs only NIS 2,100–3,500, a very small fraction of a woman’s lifetime income—at most about 0.6%, on conservative assumptions.²³ For a daughter who can borrow against her future earnings, the utility cost of paying for it is correspondingly small. Her decision instead turns on ξ^D , which bundles together forces far larger than this price: the monetary cost of raising a child and the non-monetary costs and benefits of motherhood and abortion. The price is decisive only when these large forces nearly cancel each other, leaving ξ^D close to zero—so the complier band $(\underline{\xi}_1^D, 0]$ is narrow.

Layer 2: lifetime budget + liquidity constraint. The daughter is now barred from borrowing against future income. She faces both a lifetime budget and a first-period liquidity constraint:


$$\begin{aligned} \text{lifetime budget:} \quad & C_1^D + C_2^D + p(1 - S)A \leq Y_1^D + Y_2^D, \\ \text{liquidity constraint:} \quad & p(1 - S)A + C_1^D \leq Y_1^D. \end{aligned}$$

Figure 3, Panel (b), illustrates how decision thresholds move relative to Layer 1. Preferences are unchanged, so the never-taker boundary is unchanged, $\bar{\xi}_2^D = \bar{\xi}_1^D = 0$: a daughter who would decline a free abortion continues to do so regardless of liquidity. The always-taker boundary, $\underline{\xi}_2^D$, changes in two ways. First, if first-period income does not exceed the price, $Y_1^D \leq p$, self-funding is infeasible and there are no always-takers: abortion becomes possible only under the subsidy. In that case, $\underline{\xi}_2^D \rightarrow -\infty$. Second, even when self-funding is feasible, paying p comes entirely out of period-1 consumption, whose marginal utility is high because $Y_1^D < Y_2^D$. The utility cost of paying the full price is therefore higher than under intertemporal smoothing, so $\underline{\xi}_2^D < \underline{\xi}_1^D$ strictly and the complier region expands.

The liquidity channel is potentially large. The relevant resource is no longer lifetime income but the income on hand during the roughly two-month window in which the decision must be made. In our sample of young unmarried pregnant women, 72.7% did not have stable employment in the year

²³We approximate the daughter’s lifetime income as the present discounted value of her earnings over a 35-year career, discounted at a 3% real rate. Starting from the low level of below-median-SES young women in our sample (NIS 15,652 per year) and assuming 3% real annual wage growth, lifetime income is about NIS 547,000. The path is conservative—after 35 years her earnings reach only about NIS 43,000, still short of the previous generation’s level (low-SES fathers in our sample earn NIS 50,283 per year)—yet the procedure costs at most 0.6% of lifetime income.

Figure 3: The subsidy's effect across the three layers



Notes: This figure illustrates how abortion decisions and the effect of an abortion subsidy depend on the abortion attitudes of both daughters and parents. Blue regions are *always-takers*, who abort with or without the subsidy; orange regions are *compliers*, who abort only when subsidized; and gray regions are *never-takers*, who do not abort. The ordering of the thresholds across layers is formalized in Lemma 1 in Appendix E.

before conception, making an out-of-pocket abortion very difficult.²⁴ Even among daughters with stable employment, the procedure costs, on average, 38–63% of their earnings over the two-month window—hardly a negligible expense.²⁵

Layer 3: lifetime budget + liquidity constraint + family transfers. The daughter is liquidity-constrained as in Layer 2 but now interacts with her parents through the game described in Section 4.1. Her budget constraints become

$$\begin{aligned} \text{lifetime budget:} \quad & C_1^D + C_2^D + p(1-S)(1-k)A \leq Y_1^D + Y_2^D + T, \\ \text{liquidity constraint:} \quad & p(1-S)(1-k)A + C_1^D \leq Y_1^D + T, \\ \text{in-kind floor:} \quad & C_1^D \geq \phi T. \end{aligned}$$

The in-kind floor states that the daughter must consume the in-kind portion ϕT of her allowance; only the fungible portion $(1-\phi)T$ can be redirected to paying for an abortion procedure. Combining the floor with the liquidity constraint yields the *funding constraint* for an out-of-pocket abortion,

$$p(1-S)(1-k)A \leq Y_1^D + (1-\phi)T,$$

which makes transparent why ϕ matters: a daughter with a generous but largely in-kind allowance may still be unable to pay for a procedure her parents oppose.

Figure 4 summarizes the four-stage timing of the Layer-3 game. Parents choose T ex ante—before the daughter’s pregnancy status is realized and as if she will not become pregnant—to maximize their utility V^P , subject to $C^P + T \leq Y^P$.²⁶ After the pregnancy is realized and the daughter has decided whether to disclose, parents choose $k \in \{0, 1\}$ in the TELL branch, anticipating the daughter’s subsequent abortion choice, to maximize V^P subject to $C^P + T + pk(1-S)A \leq Y^P$. Appendix E states the game formally and solves it by backward induction.


Figure 3, Panel (c), shows how family transfers change the decision thresholds relative to Layer 2. These transfers do not affect preferences and therefore leave the never-taker boundary unchanged, $\bar{\xi}_3^D = \bar{\xi}_2^D = 0$. They do, however, remove some daughters from the complier group. Some can now self-finance using the fungible part of their allowance, which raises the always-taker boundary, $\underline{\xi}_3^D \geq \underline{\xi}_2^D$. Others can obtain parental coverage when parents are sufficiently supportive, as summarized by the

²⁴A daughter has *stable employment*—equivalently, “works” in Table 1 and below—if her earnings in the year before conception exceed six months of the minimum wage.

²⁵Liquidity constraints are likely relevant for daughters from both high- and low-SES families. Table A.1 shows that, although total household income is more than four times higher in high-SES families, daughters’ own income is only 15% higher.

²⁶We assume parental altruism α^P is large enough that the non-pregnancy optimum is interior, $T > 0$.

Figure 4: Extensive form of the Layer-3 game



Notes: Extensive form of the four-stage game presented in Section 4; details are provided in Appendix E.4. Rounded boxes are decision nodes, the circle is Nature, and shaded rectangles are terminal subgames. Symbols: T (parental allowance), E (pregnancy realization), R (daughter's disclosure), k (parental coverage), and A (abortion).

parental-coverage threshold $\bar{\xi}^P$.

The same panel also illustrates how the subsidy affects the daughter's autonomy: we say she has autonomy when the abortion outcome depends only on her own attitude, ξ^D , and not on her parents' attitude, ξ^P . With the subsidy, only the vertical boundary at $\xi^D = 0$ remains relevant—she aborts if and only if $\xi^D \leq 0$, regardless of her parents' attitude, and therefore has autonomy. Without the subsidy, the outcome may also depend on ξ^P , implying a lack of autonomy.

The resulting partition of the abortion-attitude space, (ξ^D, ξ^P) , has three regions. *Never-takers* are families with $\xi^D > 0$: the daughter does not abort under either subsidy regime. Daughters in the *always-taker* region abort regardless of the subsidy. This region includes all families in which the daughter wants an abortion ($\xi^D \leq 0$) and at least one of two conditions holds: the daughter is willing to pay the full price herself ($\xi^D \leq \underline{\xi}_3^D$), or the parents are willing to cover the procedure ($\xi^P \leq \bar{\xi}^P$). The always-taker region could be further subdivided according to who pays for the abortion. However, because our focus is on the subsidy's effect on abortion take-up, what matters here is whether the abortion occurs, not who pays for it. We therefore treat the always-taker region as a single region.²⁷

Greater autonomy changes abortion outcomes only for the remaining families: daughters who want an abortion ($\xi^D \leq 0$), do not self-finance at the full out-of-pocket price ($\xi^D > \underline{\xi}_3^D$), and have parents who are unwilling to cover the procedure ($\xi^P > \bar{\xi}^P$).²⁸ These are the *compliers*, or mismatch families. The resulting complier region is the rectangle²⁹

$$\text{complier region} = (\underline{\xi}_3^D, 0] \times (\bar{\xi}^P, \infty). \quad (5)$$

4.3 Abortion Attitudes and the Subsidy Effect

We now formalize the model's implications for how the effect of the subsidy varies with abortion attitudes. Let $A_l(\xi^D, \xi^P; S) \in \{0, 1\}$ denote the daughter's abortion outcome in Layer l under subsidy regime $S \in \{0, 1\}$, and define

$$\Delta A_l(\xi^D, \xi^P) = A_l(\xi^D, \xi^P; 1) - A_l(\xi^D, \xi^P; 0).$$

Thus, $\Delta A_l(\xi^D, \xi^P) = 1$ indicates that the subsidy changes the daughter's abortion decision in Layer l , i.e., that she is a complier.

Proposition 1 (Abortion attitudes and the subsidy effect). *The complier indicator ΔA_l satisfies the following properties:*

²⁷Who pays for the procedure matters for welfare; we return to this issue in Section 6.

²⁸Here, “do not self-finance” may reflect either literal infeasibility or the fact that paying p out of pocket is too costly relative to the daughter's abortion benefit.

²⁹Appendix F.6 derives Equation (5).

(a) Weakly decreasing in the daughter's abortion distaste when self-funding is infeasible. In any layer $l \in \{1, 2, 3\}$, if self-funding is infeasible ($\xi_l^D = -\infty$), then ΔA_l is weakly decreasing in ξ^D .

(b) Weakly increasing in the parents' abortion distaste. In Layer 3, ΔA_3 is weakly increasing in ξ^P .

Proof. See Appendix F.7.

Proposition 1 shows that liquidity constraints alone cannot explain why the subsidy has larger effects in more abortion-averse communities. When self-funding is infeasible, the subsidy effect is weakly *decreasing* in the daughter's own abortion distaste ξ^D (Proposition 1(a)). Layer 3 provides a potential explanation through parental transfers: the subsidy effect is weakly *increasing* in the parents' abortion distaste ξ^P (Proposition 1(b)). This is the autonomy mechanism at work: the subsidy is most likely to change the daughter's decision when parental opposition would otherwise prevent an abortion she wants. However, if more abortion-averse communities have both more abortion-averse parents and more abortion-averse daughters, these two forces work in opposite directions. Whether the subsidy has larger effects in such communities therefore depends on how parental and daughter attitudes vary across communities, a question we take to the data next in Section 5.

Proposition 1 is silent on two important questions: the welfare consequences of the subsidy and the daughter's decision to disclose her pregnancy. We return to both in Section 6.

Finally, these results do not rely on disclosure being costless. Appendix G extends the model by adding a direct cost of involving one's parents, proportional to their abortion distaste, and shows that Proposition 1 continues to hold (Proposition 5). A direct cost of parental involvement therefore reinforces the same mechanism highlighted by the baseline model.

5 Taking the Model to the Data

Section 4 studied one family with fixed attitudes (ξ^D, ξ^P) and asked when the subsidy changes the daughter's abortion decision. To take the model to the data, we now aggregate it to the level of a community. Our object of interest is the *abortion ratio*—the share of pregnancies that end in abortion—and our focus is how the subsidy moves it. A *community* is indexed by $h \in \{0, 1\}$, with $h = 1$ indicating a strict community and $h = 0$ a lenient one. Within a community, there is a continuum of families, each of them characterized by abortion attitudes (ξ^D, ξ^P) , which are distributed following a joint distribution f_h . All families in both communities share the same economic environment—incomes, the price p , the in-kind share ϕ , and parental altruism α^P .³⁰

³⁰We hold the economic environment fixed across communities so as to isolate the role of attitudes; we discuss the consequences of *income* differences across families in Section 5.3.

The abortion ratio in community h under subsidy regime S is

$$A(S | h) \equiv \Pr(A_3(\xi^D, \xi^P; S) = 1 | h), \quad (6)$$

where the probability is taken with respect to the joint distribution f_h of (ξ^D, ξ^P) in community h . The subsidy's effect is $\Delta A(h) \equiv A(1 | h) - A(0 | h)$. The effect equals the mass of compliers—daughters whose abortion decision is changed by the subsidy—which corresponds to the mismatch region in Figure 3, Panel (c).

When the procedure is free, every daughter below the never-taker threshold, $\bar{\xi}_3^D$, obtains an abortion. Thus, the abortion ratio under the full subsidy is

$$A(1 | h) = \Pr(\xi^D \leq \bar{\xi}_3^D | h).$$

Because the thresholds depend only on the (common) economic environment and not on realized attitudes,³¹ the three thresholds in Panel (c) of Figure 3 ($\bar{\xi}_3^D$, $\bar{\xi}^P$, and $\bar{\xi}_3^D$) are common across communities. What differs across communities is the distribution of abortion attitudes f_h , which determines how the population mass is spread across the regions of the figure.

Section 4 left us with a question. The subsidy effect for a single family is weakly increasing in the parents' abortion distaste, ξ^P , but not in the daughter's own distaste, ξ^D (Proposition 1). Whether the effect is larger in stricter communities therefore cannot be read off the single-family problem alone: it depends on how daughter and parent attitudes are distributed.

We answer that question in three steps. Section 5.1 derives the model's predictions for how subsidy effects at the community level vary with strictness and relates these predictions to the empirical estimates in Section 3. The results point to intergenerational mismatch in abortion attitudes as the central explanation for the observed subsidy effects. Section 5.2 validates this mechanism using independent survey data on abortion attitudes. Finally, Section 5.3 turns to the role of income. We find suggestive evidence that the intergenerational-mismatch channel bites only when the daughter lacks the resources to pay for the abortion herself.

5.1 Community Strictness and Subsidy Effects

When comparing the *total* subsidy effect, $\Delta A(h)$, across communities, two opposing forces are at work. The first is the *abortion pool*: the share of pregnancies that would end in abortion if the procedure were free. Fewer daughters in stricter communities want an abortion in the first place, pushing the effect down. The second is the *subsidy-induced share*: the share of the abortion pool made up of compliers

³¹Lemma 1(d) in Appendix E.

rather than always-takers. Fewer parents in stricter communities are willing to cover their daughters' abortions, increasing the mass of compliers and pushing the effect up. The following decomposition makes this precise:

$$\overbrace{\Delta A(h)}^{\text{total effect}} = \overbrace{A(1 | h)}^{\text{abortion pool}} \cdot \overbrace{\frac{\Delta A(h)}{A(1 | h)}}^{\text{subsidy-induced share}}. \quad (7)$$

Because the two forces move in opposite directions as communities become stricter, the total effect is theoretically ambiguous, even though each factor is monotone on its own. We now formalize this intuition, beginning with what it means for one community to be stricter than another.

Assumption 1 (Bivariate MLRP across communities). *The strict and lenient communities' joint densities, f_1 and f_0 , have a common support that is a sublattice of $\mathbb{R}^2$ (so that $x \vee y$ and $x \wedge y$ remain in the support whenever x, y do), and are ranked in the bivariate likelihood-ratio order: for all $x = (\xi^P, \xi^D)$, $y = ((\xi^P)', (\xi^D)')$ in the common support,*

$$f_1(x) f_0(y) \leq f_1(x \vee y) f_0(x \wedge y),$$

where $\vee$ and $\wedge$ are the coordinatewise maximum and minimum in (ξ^P, ξ^D) . We further assume that $A(1 | h) > 0$ for both $h \in \{0, 1\}$.

Assumption 1 is a standard stochastic-ordering condition—the bivariate likelihood-ratio order (Karlin and Rinott, 1980; Shaked and Shanthikumar, 2007). It formalizes a simple idea: in a stricter community, families in which *both* daughters and parents are abortion-averse are more common.³² The assumption imposes neither independence between ξ^D and ξ^P nor any specific parametric form. The next proposition characterizes how the subsidy effect varies across communities under this condition.

Proposition 2 (Community strictness and subsidy effects). *Maintaining the model of Section 4, under Assumption 1:*

- (i) *The abortion pool, $A(1 | h)$, is weakly smaller in the strict community ($h = 1$) than in the lenient community ($h = 0$), while the subsidy-induced share, $\Delta A(h)/A(1 | h)$, is weakly larger.*
- (ii) *Without further assumptions, the total effect, $\Delta A(h)$, can be larger or smaller in the strict community than in the lenient community.*

³²Assumption 1 restricts both the marginal and joint distributions of attitudes. Marginally, it implies that ξ^D and ξ^P are each stochastically higher in the strict community; jointly, it rules out distributions in which ξ^D and ξ^P are so negatively correlated in the strict community that both-averse families are more common in the lenient community despite the marginal shift.

Proof. See Appendix H.

Proposition 2 delivers two distinct sets of implications. First, it gives two sharp predictions: stricter communities should have a smaller abortion pool but a larger subsidy-induced share. Because these predictions rest on weak assumptions, they provide useful tests of the model. Second, the proposition leaves the total effect unsigned: without further restrictions, the total effect $\Delta A(h)$ can be larger or smaller in stricter communities, depending on which force dominates.

We now revisit the estimates in Table 1 through this lens. Empirically, the abortion pool, $A(1 | h)$, corresponds to the abortion ratio when the subsidy is in place. The total subsidy effect, $\Delta A(h)$, corresponds to the estimated treatment effect reported in the table. The subsidy-induced share, $\Delta A(h)/A(1 | h)$, is the ratio of these two objects.

The data line up with both sharp predictions. First, the abortion pool is smaller in strict communities, as predicted: when the procedure is free, 65.7% of pregnancies end in abortion in strict communities, compared with 82.9% in lenient communities. Second, the subsidy-induced share moves in the opposite direction. Among women in the abortion pool, 22.5% are compliers in strict communities, compared with 6.8% in lenient communities. This agreement supports the model's central mechanism: stricter communities have a smaller pool of women who would choose abortion when it is free, but a larger share of that pool is constrained by parental opposition absent the subsidy.

The remaining question is which of these two forces dominates. Proposition 2 leaves this theoretically ambiguous, but Table 1 gives a clear empirical answer: the subsidy effect is much larger in strict communities. Equation (7) shows why. The abortion pool is only about 20% smaller in strict communities, while the subsidy-induced share is more than three times larger. The smaller abortion pool is therefore more than offset by the larger induced share, generating a larger total effect in strict communities.

These patterns are informative about how abortion attitudes vary across communities: the larger effect in strict communities points to greater intergenerational mismatch, with the difference in parental opposition between strict and lenient communities exceeding the corresponding difference among daughters. We examine this implication directly in Section 5.2, using survey data from an independent source on abortion attitudes.

5.2 Testing the Intergenerational-Mismatch Mechanism

Equation (7) decomposes the total effect of the subsidy into two components. These components differ across communities on highly asymmetric scales: the abortion pool is only about 20% smaller in strict communities, while the subsidy-induced share is more than three times as large. Why is the cross-community difference so much larger for one object than for the other?

One potential explanation is greater *intergenerational mismatch* in abortion attitudes in strict communities. In lenient communities, parents and daughters hold similarly permissive attitudes: many daughters want an abortion, generating a large abortion pool, and many parents are willing to cover the procedure, leaving few compliers. In strict communities, by contrast, daughters are only somewhat more abortion-averse than lenient daughters, so the abortion pool shrinks only modestly, while parents are much more abortion-averse and much more likely to refuse coverage. The subsidy therefore bites harder in strict communities because it closes a larger wedge between daughters' demand for abortion and parents' willingness to finance it. In the language of Section 4, the subsidy grants autonomy in both communities, but autonomy changes behavior only in mismatch families, which are far more prevalent in strict communities.

Survey evidence. We now test this mechanism directly using public opinion data. We use data from the Israeli World Values Survey (WVS) Wave 4, collected in 2001 (Inglehart et al., 2014), and define *Young* women as those aged 18–23.³³ We contrast this group with older generations, defined as men and women aged 45–75. We classify respondents as *Religious* if they report that religious faith is an important quality that children should be encouraged to learn at home.³⁴ Attitudes toward abortion are based on responses to whether abortion is justifiable on a scale from 1 (never justifiable) to 10 (always justifiable), with *Lenient* attitudes defined as responses of 6 or higher.

To quantify intergenerational mismatch in abortion attitudes, we estimate the following regression:

$$\text{Lenient}_i = \beta_0 + \beta_Y \text{Young}_i + \beta_R \text{Religious}_i + \beta_{YR} (\text{Young}_i \times \text{Religious}_i) + \epsilon_i, \quad (8)$$

where β_R captures the difference in abortion attitudes between religious and secular respondents, β_Y measures the difference between younger and older respondents in secular communities—our measure of intergenerational mismatch—and β_{YR} measures how much larger this mismatch is in religious communities than in secular ones. The error term ϵ_i captures idiosyncratic variation.

OLS estimates of Equation (8) using Israeli data are reported in Column (1) of Table 3. The coefficient on *Young* is small and statistically insignificant, indicating little intergenerational mismatch in abortion attitudes among secular communities. In contrast, the coefficient on *Religious* is negative, large, and statistically significant, confirming that religious respondents hold stricter views on abortion. This validates our classification of these groups as “strict” in Section 3.2. Finally, the interaction term *Young* $\times$ *Religious* is large, positive, and statistically significant, consistent with the prediction that intergenerational mismatch is greater in religious communities.³⁵

³³ Although our main empirical analysis includes women aged 16–23, 18 was the minimum age surveyed in the WVS.

³⁴ Respondents were asked to rank up to five qualities that were important to impart to children.

³⁵ Column (2) of Table 3 reports analogous results for the United States, consistent with the U.S. evidence in Section 3.4 that parental involvement laws have larger effects on abortion outcomes in more religious states.

Table 3: Survey Evidence on Intergenerational Mismatch in Abortion Attitudes

	Israel (WVS)	US (Pew)
Young	-0.072 (0.064)	0.018 (0.016)
Religious	-0.415*** (0.043)	-0.363*** (0.006)
Young $\times$ Religious	0.218** (0.104)	0.085*** (0.029)
N	517	19,527
Dep. Var. Mean	0.429	0.665

Notes: This table reports OLS estimates of Equation (8) using public opinion data from Israel (World Values Survey Wave 4, 1999–2004) and the United States (Pew Research Center Religious Landscape Study, 2023–2024). Each column presents a regression of an indicator for holding a lenient attitude toward abortion on indicators for being a young woman (ages 18–23), for belonging to a more religious or conservative group, and their interaction. The coefficient on *Young* captures intergenerational differences among secular respondents; the coefficient on *Religious* captures average differences between religious and secular individuals; and the interaction term *Young $\times$ Religious* measures how the intergenerational gap varies with religiosity. Standard errors are reported in parentheses. * $p < 0.10$, ** $p < 0.05$, and *** $p < 0.01$.

Model calibration. The results in Table 3 are consistent with the interpretation that the subsidy has larger effects in strict communities because these communities exhibit a larger intergenerational mismatch in abortion attitudes. We now ask whether this mismatch is quantitatively large enough to explain the estimated difference in subsidy effects across communities. To do so, we calibrate the model using the estimated effects for the strict group, and then use the calibrated model, together with the observed distribution of abortion attitudes, to predict the subsidy effect for the lenient group. This section summarizes the calibration strategy and main results; Appendix I provides the details.

We interpret the observed distribution of abortion attitudes in the WVS as the empirical counterpart of the distributions of the abortion-distaste parameters in the model, ξ^D and ξ^P . We then choose the model thresholds $(\underline{\xi}_3^D, \bar{\xi}_3^D, \bar{\xi}^P)$ so that the predicted subsidy effects match the estimated effects for the strict group in Table 1. Holding these thresholds fixed, we combine them with the observed distribution of abortion attitudes in the lenient group to predict the subsidy effect in that group. We then compare this prediction to the estimated lenient-group effect.³⁶

The key identifying restriction in this exercise is that the thresholds are common across strict and lenient communities. By holding the thresholds fixed, the exercise asks whether differences in abortion attitudes alone—with all other primitives held fixed—can account for the observed difference in subsidy effects.³⁷ In empirical terms, the restriction amounts to assuming that religious and secular respondents interpret the WVS abortion-attitudes scale in the same way, even though their responses differ.

³⁶The calibration recovers the corresponding cutoffs on the WVS scale, which we assume maps into utility up to an unknown location shift. Thus, $\bar{\xi}_3^D = 0$ in utility units need not correspond to zero on the WVS scale, and its WVS-scale counterpart must also be calibrated. Appendix I provides details.

³⁷Lemma 1(d) in Appendix E supports this restriction by showing that the thresholds do not depend on the distributions of ξ^D or ξ^P .

We implement the calibration under progressively more flexible approaches. First, under two simplifying assumptions—that daughters cannot self-finance the procedure and that abortion attitudes are independent within families—the model generates an exact prediction for the lenient-group subsidy effect. Second, we relax the within-family independence assumption and study how the prediction varies with the correlation in abortion attitudes. Third, we relax the no-self-financing assumption and report bounds over all parameter values that rationalize the strict-group estimates.

The no-self-financing assumption is motivated by the limited financial resources of daughters in our sample: 72.7% did not have stable employment in the year before conception, and, among those who did, the procedure would cost 38–63% of two months’ earnings. The within-family independence assumption is imposed for tractability: it allows the model to generate a closed-form prediction. We relax this assumption below.

Under these assumptions, the model predicts a subsidy effect for the lenient group of $\Delta A = 0.050$. This prediction is very close to the estimated lenient-group effect, 0.056, and much smaller than the estimated strict-group effect, 0.148. Thus, the observed difference in the marginal distributions of abortion attitudes across generations is quantitatively large enough to explain the difference in subsidy effects across communities.

We next examine the role of a second form of mismatch: within-family mismatch. The previous exercise captures intergenerational mismatch through differences in the marginal distributions of attitudes between parents and daughters. However, the subsidy effect may also depend on the within-family correlation in attitudes, denoted by ρ . An alternative explanation for our findings is that attitudes are less strongly correlated between parents and daughters in strict communities, making the subsidy more likely to relax a binding parental constraint.

To evaluate this possibility, we recalibrate the model using the same procedure described above, but allow ρ to vary between 0 and 0.99, from independence to near-perfect correlation. We first impose the same value of ρ in both groups. We then allow the correlation to differ across groups by fixing $\rho = 0.5$ in the strict group and varying ρ between 0 and 0.99 in the lenient group. Figure I.1 reports the results. Across the full range of values considered, the predicted lenient-group effect is at most about half of the strict-group effect. Moreover, the results do not support the alternative explanation that the strict community has a weaker parent-daughter correlation in abortion attitudes; if anything, the calibration implies a stronger correlation in the strict group.³⁸

Finally, we relax the assumption that daughters cannot self-finance an abortion by allowing the self-financing threshold ξ_3^D to be finite rather than fixing it at $-\infty$. We instead impose an alternative restriction: a share of daughters—calibrated to be at least as large as the share not employed in our

³⁸In Panel (b) of Figure I.1, the predicted lenient-group effect matches its estimated counterpart when ρ_{lenient} is approximately 0.35, while ρ_{strict} is fixed at 0.5. Thus, the model fits the data best when the within-group correlation in abortion attitudes is stronger in the strict group.

sample—cannot self-finance the procedure, regardless of their abortion attitudes. We then consider all combinations of $(\xi_3^D, \bar{\xi}_3^D, \bar{\xi}^P, \rho)$ that match the estimated strict-group subsidy effect and, for each combination, compute the implied lenient-group effect. This exercise yields the bounds $\Delta A_{\text{lenient}} \in [0.002, 0.131]$. The estimated lenient-group effect lies within these bounds, whereas the estimated strict-group effect lies above them.

In sum, the calibration shows that the observed intergenerational mismatch in abortion attitudes is large enough to account for the heterogeneous subsidy effects. When the model is disciplined by the strict-group estimates and the observed distribution of abortion attitudes, it predicts a much smaller effect in lenient communities, close to the effect observed in the data. This conclusion is robust to allowing for different degrees of within-family correlation in attitudes and to relaxing the assumption that daughters cannot self-finance the procedure.

5.3 Parental Income versus the Daughter’s Liquidity

We now examine how income shapes the subsidy effect, drawing on two patterns in the data. First, the effect is essentially flat in parental resources: Table 1 reports treatment effects of 0.071 for low-SES families and 0.064 for high-SES families—a small and statistically insignificant difference—even though high-SES households earn about 4.7 times as much as low-SES households. Second, the effect varies with the daughter’s own resources: the point estimate among daughters who do not work is nearly twice that among working daughters, although the difference is not statistically significant (Columns (4) and (5) of the same table). Taken together, these patterns are puzzling for a simple liquidity story: if the binding constraint were the family’s ability to pay for the procedure, the subsidy should matter most where resources are scarcest—among low-SES families—and that is not what we find.

The next proposition reconciles these patterns in three results. First, once parental income is sufficiently large relative to the procedure cost, the subsidy effect becomes insensitive to further income increases. Second, in this saturated regime, the subsidy has no effect on daughters who can self-finance the procedure. Third, still in the saturated regime, the effect vanishes if parental transfers include any fungible cash. Together, these results identify when the subsidy matters and through which channel.

Assumption 2 (Vanishing marginal utility). $u'(c) \rightarrow 0$ as $c \rightarrow \infty$.

Assumption 2 is satisfied by CRRA utility and by most standard concave specifications.

Proposition 3 (Income, liquidity, and the subsidy effect). *Let $\Delta A(h, Y^P, Y_2^D, Y_1^D, \phi)$ denote the population subsidy effect $\Delta A(h) \equiv A(1 | h) - A(0 | h)$ from Equation (6), written to make explicit its dependence on incomes and the in-kind share; the price p and parental altruism α^P are held fixed. Under the model of Section 4 and Assumption 2, fix (h, Y_1^D, ϕ) . Then $\Delta A(h, Y^P, Y_2^D, Y_1^D, \phi)$ converges to a limit $\Delta A^\infty(h, Y_1^D, \phi)$ as $\min(Y^P, Y_2^D) \rightarrow \infty$, and:*

(a) Saturation in adult incomes. For any $b > 0$, the difference

$$\Delta A(h, b \cdot Y^P, b \cdot Y_2^D, Y_1^D, \phi) - \Delta A(h, Y^P, Y_2^D, Y_1^D, \phi)$$

tends to 0 as $\min(Y^P, Y_2^D) \rightarrow \infty$.

(b) Daughter's hidden option. If $Y_1^D \geq p$, then $\Delta A^\infty(h, Y_1^D, \phi) = 0$ for any h and any ϕ .

(c) Liquid transfers. If $\phi < 1$, then $\Delta A^\infty(h, Y_1^D, \phi) = 0$ for any h and any Y_1^D .

Proof. See Appendix J.

Proposition 3(a) explains the first pattern. Once adult incomes—parental income Y^P and the daughter's future income Y_2^D —are large relative to the procedure cost, further increases no longer affect parents' willingness to support their daughter's abortion, and the subsidy effect becomes flat in income. The data suggest that even low-SES families are in or near this saturated regime: over the roughly two-month decision window, low-SES households earn 2.9 times the upper bound of the procedure cost. Viewed through the proposition, the similar effects among low- and high-SES families are exactly what the model predicts, despite the large income differences between the two groups.

Proposition 3(b) turns to the second pattern: the daughter's own income at the time of pregnancy. Working daughters in our sample earn, on average, NIS 5,563 over the two-month decision window, against a procedure cost of NIS 2,100–3,500, a large share of their earnings. One might expect the subsidy to still matter for these women. The proposition says otherwise, and the reason is the saturated regime itself. Because parental transfers already leave the daughter with high period-1 consumption, her marginal utility of cash is low: giving up p out of her own earnings is cheap in utility terms. Her income therefore matters only through liquidity—whether she *can* pay—and not through affordability. Once $Y_1^D \geq p$, the subsidy effect is zero; it makes no difference whether her earnings exceed the cost by little or by much.

Proposition 3(b) also predicts how the work-status gradient should vary across communities. Because the proposition holds for any h , working daughters should show no effect in strict and lenient communities alike. Among non-working daughters, the effect is large only where many daughters who want an abortion face parents unwilling to pay—and Sections 5.1–5.2 showed that such mismatch families are concentrated in strict communities. The model therefore predicts a large effect for a single group: non-working daughters from strict backgrounds.

This is the pattern we see in Table A.2: in strict communities the effect falls from 0.173 among non-working daughters to a statistical zero (-0.016) among working daughters, while effects for lenient families are small regardless of work status. As noted in Section 3.2, the samples underlying these splits are small, so we read these results as supportive but not definitive.

Finally, Proposition 3(c) concerns the composition of parental transfers: in-kind versus liquid support. If transfers had any meaningful fungible component ($\phi < 1$), the subsidy effect should be small even for non-working daughters: as parental income grows, the fungible share $(1 - \phi)T$ also grows without bound, eventually giving daughters who want an abortion enough cash to pay for the procedure themselves. The large effect among non-working daughters from strict communities is therefore inconsistent with the saturated limit under $\phi < 1$ and points to parental support being predominantly in kind—reaching daughters in forms they cannot redirect against their parents’ wishes. As a result, daughters without earnings of their own may be unable to fund the procedure even when their parents have substantial resources.³⁹

In sum, parental income matters little because families appear to be past the point where resources constrain parents’ choices. What matters is the conjunction that defines the complier region: the subsidy changes outcomes only when the daughter cannot finance the procedure herself *and* her parents are unwilling to do so. The first condition determines whether the subsidy grants her autonomy; the second determines whether that autonomy changes behavior. This refines the reading of the income splits in Section 3.2: the subsidy does relax a financial constraint, but the binding margin is the daughter’s own liquidity in the face of parental opposition, not household resources.

6 Welfare

The preceding sections established that the subsidy operates through an abortion attitude mismatch: liquidity-constrained daughters who want an abortion but cannot afford one, and whose parents oppose abortion and refuse to finance the procedure. This section asks what that result implies for welfare. Throughout, welfare means *private family welfare*—the payoffs of the daughter and her parents.⁴⁰

First, we sign the subsidy’s welfare effect for each group of families: it is zero for never-takers, small and positive for always-takers, and, for compliers, it depends crucially on whether, when evaluating an abortion that the daughter hides from her parents, parental attitudes enter the welfare calculation. Second, we show that, if parental attitudes do enter, the data do not allow us to compare the utility gain of daughters who obtain an abortion because of the subsidy with the utility loss of parents who oppose it. Finally, we frame these results in terms of autonomy—who holds decision-making authority within the family—and connect them to parental consent laws for abortion.

³⁹Propositions 3(b) and 3(c) are limiting results: the subsidy effect vanishes only as adult incomes grow without bound. At finite incomes, the model allows small effects for working daughters and for daughters receiving some liquid transfers. The practical interpretation is not that transfers are literally fully in kind, but that the in-kind share is high enough to leave daughters unable to finance an abortion even when parental income is large relative to the procedure cost.

⁴⁰We evaluate welfare conditional on the daughter’s pregnancy. Since the subsidy affects only families in which the daughter is pregnant, these are the only families whose welfare it changes.

6.1 The welfare effect of the subsidy

We evaluate welfare using the daughter's utility V^D from Equation (3), and the parents' utility net of the altruism term, $w^P \equiv u(C^P) - \xi^P \cdot A \cdot R$. Family welfare is the sum $W \equiv V^D + w^P$.⁴¹

Computing family welfare requires one crucial normative choice: should parental attitudes enter the welfare calculation? We consider two alternatives. The *knowledge benchmark* includes parental attitudes, ξ^P , only when parents know about the abortion. The *action benchmark* includes them whenever an abortion occurs, including abortions the daughter hides. Parents make an active decision only when the pregnancy is revealed, and on that branch the two benchmarks coincide. Equilibrium behavior is therefore identical under either benchmark.

We write ΔW for the change in family welfare as the subsidy is introduced. The following proposition signs ΔW for each of the three groups of families.

Proposition 4 (The welfare effect of the subsidy). *By the partition of families into always-takers, never-takers, and compliers, presented in Figure 3, Panel (c):*

- (a) Never-takers. $\Delta W = 0$ under either benchmark.
- (b) Always-takers. $\Delta W > 0$ under either benchmark.
- (c) Compliers.

$$\Delta W = \begin{cases} -\xi^D - \min\{\xi^P, 0\} \geq 0, & \text{under the knowledge benchmark,} \\ -\xi^D - \xi^P, & \text{under the action benchmark.} \end{cases}$$

Proof. See Appendix K.

Figure 5 illustrates Proposition 4 by displaying the sign of ΔW across the attitude space. Part 4(a) is immediate: the subsidy changes nothing for never-takers. Part 4(b) establishes that the subsidy unambiguously benefits always-takers.⁴² This gain is somewhat mechanical, since we focus on the family receiving the subsidy and abstract from how the subsidy is financed. Appendix L presents corresponding results that account for the fiscal cost of the subsidy. It also provides a market-completion interpretation of the subsidy effect. Financing a period-1 procedure by taxing the daughter's own period-2 income functions as a loan against future earnings, while pooling the tax cost across all daughters—most of whom never become pregnant—insures her against the pregnancy shock. The subsidy therefore completes two markets she cannot access on her own: credit and insurance.

⁴¹Stripping the altruism term avoids double-counting the daughter's consumption, which enters V^P through α^P ; none of our conclusions depends on this choice. Results are also unchanged if we use any weighted sum of the two felicities with strictly positive generation weights (Remark K.1).

⁴²For families in which parents pay for the procedure even though their daughter could self-finance, ΔW remains strictly positive, but its magnitude differs across benchmarks; see Remark K.2.

Part 4(c) concerns compliers: the families whose abortion decision the subsidy changes. For them the subsidy changes nothing in the budget—without it the daughter forgoes an abortion she cannot fund, with it she obtains one she does not pay for—so ΔW is the family’s valuation of the abortion itself. The daughter gains $-\xi^D \geq 0$ by revealed preference. Whether her parents’ valuation also enters depends on whether they learn of the abortion: a complier discloses to approving parents who merely would not pay ($\bar{\xi}^P < \xi^P \leq 0$) and hides from opposed ones ($\xi^P > 0$), as Section 6.3 explains. Under the knowledge benchmark only informed parents count, so the approving parents add their gain $-\xi^P \geq 0$ while the opposed parents add nothing: $\Delta W = -\xi^D - \min\{\xi^P, 0\} \geq 0$. Under the action benchmark parental distaste counts whether or not they learn, $\Delta W = -\xi^D - \xi^P$, whose sign is ambiguous. The benchmarks thus agree wherever the abortion is disclosed and differ only where it is concealed. We take up the ambiguous case next in Section 6.2.

6.2 The welfare sign under the action benchmark

Under the action benchmark, the welfare effect of the subsidy in the complier region nets the daughter’s gain, $-\xi^D \geq 0$, against her parents’ abortion distaste, ξ^P . As Figure 5 shows, the sign of the net effect flips at the boundary $\xi^P = -\xi^D$: families below this boundary, for whom the daughter’s gain exceeds the parents’ distaste, gain on net, while families above it lose.

This split raises an empirical question: how much of the complier mass lies on each side of the boundary? The following corollary establishes that data on abortion decisions alone cannot answer this question.


Corollary 1 (The welfare sign is not identified by behavior). *On the complier region, every family has the same abortion outcome—no abortion under $S = 0$, an abortion under $S = 1$. Data on abortion behavior therefore reveal that a family is a complier, but not where it lies within the complier region: they bound attitudes only by $\xi^D \in (\xi_3^D, 0]$ and $\xi^P > \bar{\xi}^P$, while placing no upper bound on ξ^P . Consequently, under the action benchmark, data on abortion behavior identify neither the sign of ΔW for any complier family nor the complier mass on either side of the boundary $\xi^P = -\xi^D$.*

Proof. See Appendix K.

The non-identification is structural, not a limitation of our dataset. A complier is defined by three threshold crossings: the daughter would take up a free abortion, would not pay for it herself at the full price, and her parents would not pay for it on her behalf. In particular, parental refusal reveals $\xi^P > \bar{\xi}^P$ and nothing more: a parent who mildly disapproves and a parent for whom abortion is unthinkable refuse support in the same way, and ξ^P is unbounded above.⁴³

⁴³The survey data in Section 5.2 cannot resolve this question either. Signing ΔW on the complier region requires comparing the *magnitudes* of ξ^D and ξ^P in common utility units, whereas the WVS measures abortion attitudes on an ordinal scale.

Figure 5: The welfare effect of the subsidy



Notes: Sign of the change in family welfare ΔW from $S : 0 \rightarrow 1$ (Proposition 4), over the attitude partition of Figure 3, Panel (c). The figure is drawn under the *action* benchmark: never-takers (white) are unaffected, always-takers (light green) gain, and in the complier rectangle $\Delta W = -\xi^D - \xi^P$ is positive below the dashed locus $\xi^P = -\xi^D$ (dark green) and negative above it (red). Under the *knowledge* benchmark the red region vanishes. Axes are schematic and the dashed locus is not drawn to scale; its position within the rectangle is not identified (Corollary 1).

Given the non-identification result in Corollary 1, the central question for evaluating the welfare effects of the subsidy is therefore whether welfare should be assessed under the action benchmark or the knowledge benchmark. We develop this distinction further in the next section.

6.3 Welfare as a matter of decision authority

The subsidy leaves never-takers unaffected, gives always-takers a small transfer of order p , and, for compliers—the only families whose behavior it changes—has welfare consequences that depend on whether parents' distaste, ξ^P , should count for an abortion they do not know about.

Whether parents in complier families learn about the abortion depends on their attitudes.⁴⁴ For a complier whose parents approve of abortion but would not pay for it, $\bar{\xi}^P < \xi^P \leq 0$, the subsidy induces a *disclosed* abortion: her parents learn of the abortion and approve, so the two welfare benchmarks coincide. For a complier whose parents oppose abortion, $\xi^P > 0$, the subsidy induces a *concealed* abortion. In these cases, the subsidy shifts decision-making authority within the family: it grants the daughter autonomy to act without involving a parent who would have blocked her. These concealed abortions are the cases in which the welfare implications of greater autonomy are ambiguous.

To assess whether parents' attitudes, ξ^P , should enter welfare when they never learn of the abortion, it is useful to clarify what those attitudes represent in the model. In the limited-altruism structure of Equation (4), parents value their daughter's consumption but do not adopt her valuation of the abortion

⁴⁴Proposition 6 in Appendix K.

itself. Any view they hold about what the abortion means—for her or for the family—is their own and is captured by ξ^P . For example, any burden that an unwanted pregnancy, if carried to term, would impose on the daughter and, through altruism, on the parents is also reflected in ξ^P . Asking whether ξ^P should count for a concealed abortion is therefore asking *whose valuation of the costs and benefits of the daughter’s abortion is authoritative*: hers alone, as in the knowledge benchmark, or the family’s jointly, as in the action benchmark? The two benchmarks thus take opposite stances on whether to defer to the daughter’s autonomy: under the knowledge benchmark, only her valuation matters for the welfare consequences of her abortion; under the action benchmark, her parents’ valuation matters as well.

This framing connects the subsidy to the policy that allocates decision authority explicitly: parental-consent requirements.⁴⁵ Both are autonomy policies: consent laws allocate the decision by statute, while subsidies reallocate it de facto, by removing the financial dependence that gives opposed parents veto power. The two instruments operate on related margins because both affect families in which parents oppose abortion. They differ, however, in which daughters are constrained. Removing parental-consent requirements enables abortions for daughters who can self-finance; the subsidy enables abortions for those who cannot. Despite this difference, the central welfare issue is the same. In both cases, welfare cannot be signed using data on observed behavior alone; the key question is whether welfare should be assessed under the action benchmark or the knowledge benchmark.

This perspective aligns closely with the literature on parental-consent laws for abortion. These laws, which require minors to involve a parent before accessing abortion services, have been extensively studied. Some argue that they protect minors from hasty decisions and promote family involvement (MacBride, 2001; New, 2011). In contrast, a substantial empirical literature finds that such requirements delay access to care (Colman and Joyce, 2009), increase teen birth and early marriage rates (Myers and Ladd, 2020), and reduce young women’s educational attainment and labor-market outcomes (U.S. Government Accountability Office, 2025).⁴⁶

Our paper does not take a position on the normative question of who should have decision-making authority over abortion decisions. Rather, our contribution is to clarify that this question lies at the center of the welfare analysis. In this light, abortion subsidies should be understood not merely as financial transfers, but as part of a broader class of policies that shape the allocation of decision-making authority between parents and children.⁴⁷ This framing connects our analysis to a broader literature in

⁴⁵Appendix M provides a formal discussion of parental-consent laws and compares them with subsidies through the lens of the model.

⁴⁶See U.S. Government Accountability Office (2025) for a comprehensive review of the broader impacts of parental-involvement laws on reproductive and socioeconomic trajectories.

⁴⁷This insight extends to other domains in which public policy loosens intergenerational constraints, highlighting that the effects of economic instruments such as subsidies often operate through social rather than purely financial channels. Student-loan programs, for example, enable young adults to pursue higher education even when their parents are unwilling or unable to contribute (Keane and Wolpin, 2001; Belley and Lochner, 2007; Lochner and Monge-Naranjo, 2011), thereby

economics, philosophy, and law on the trade-off between shielding young individuals from potentially harmful decisions and respecting their developing autonomy (e.g., Lundberg et al., 2009; Doepke and Zilibotti, 2017; Tunick, 2023).

7 Conclusion

This paper examines how government subsidies shape abortion decisions using administrative data covering all legal abortions and births in Israel, linked to rich demographic and socioeconomic information. We find that a subsidy covering the full cost of the procedure substantially increases abortion take-up, with the response concentrated among young women from communities with strict abortion attitudes. Among women from lenient backgrounds, we cannot distinguish a small effect from no effect at all. To explain this pattern, we develop a two-period model of the family in which a liquidity-constrained daughter may need to rely on her parents to finance an abortion, while the subsidy allows her to obtain one without their financial support. The subsidy therefore increases the daughter’s autonomy: the abortion outcome comes to depend only on her own attitude rather than her parents’. This changes behavior when there is an intergenerational mismatch, so that the daughter wants an abortion her parents would refuse to finance, a mismatch that is more common in strict communities. Economic barriers to abortion can thus operate through a social channel. Survey evidence on abortion attitudes across generations supports this mechanism, and the pattern extends beyond Israel. In the United States, policies that grant or restrict young women’s autonomy over abortion decisions, including confidential access and parental involvement laws, have much larger effects in more religious communities.

Our analysis reframes debates over the public funding of abortion for young women. The central issue is not simply a standard welfare calculation that weighs gains to recipients against the fiscal cost, but a question of decision-making authority: who should decide whether a pregnancy continues, the young woman or her parents? Our framework does not resolve this normative question, but it clarifies how the policy should be evaluated by placing abortion subsidies in the same policy space as parental involvement laws, which explicitly allocate decision-making authority.

Several avenues for future research remain. Extending the data to more recent years would allow researchers to study longer-run effects of the subsidy on marriage, education, and labor-market outcomes. Richer data could also help distinguish parental involvement from stigma and information barriers by directly measuring within-family attitudes and pregnancy disclosure. Our framework also abstracts from male partners, whom we do not observe for pregnancies ending in abortion. More broadly, future work could examine how partners, families, and other social forces shape the effects of reproductive-health subsidies in other settings.

shifting control over educational investment from the family to the student.

References


- Abadie, A., Athey, S., Imbens, G. W., and Wooldridge, J. M. (2022). When Should You Adjust Standard Errors for Clustering? *The Quarterly Journal of Economics*, 138(1):1–35.
- Akerlof, G. A., Yellen, J. L., and Katz, M. L. (1996). An Analysis of Out-of-Wedlock Childbearing in the United States. *The Quarterly Journal of Economics*, 111(2):277–317.
- Almond, D., Li, H., and Zhang, S. (2019). Land Reform and Sex Selection in China. *Journal of Political Economy*, 127(2):560–585.
- Altındağ, O. and Joyce, T. (2022). Another day, another visit: Impact of Arkansas’ mandatory waiting period for women seeking an abortion by demographic groups. *Journal of Public Economics*, 213:104715.
- Amir, D. (2015). *The silencing of abortion in Israel: A feminist and international perspective, institutional and personal dilemmas*. Sociology and Anthropology series. Israel.
- Amsterdamski, S., Kyzer, L., and Agberia, R. (2021). Radio conversation.
- Ananat, E. O., Gruber, J., Levine, P. B., and Staiger, D. (2009). Abortion and Selection. *The Review of Economics and Statistics*, 91(1):124–136.
- Bagnoli, M. and Bergstrom, T. (2005). Log-concave probability and its applications. *Economic Theory*, 26(2):445–469.
- Belley, P. and Lochner, L. (2007). The Changing Role of Family Income and Ability in Determining Educational Achievement. *Journal of Human Capital*, 1(1):37–89.
- Bhalotra, S., Clots-Figueras, I., and Iyer, L. (2021). Religion and abortion: The role of politician identity. *Journal of Development Economics*, 153:102746.
- Bitler, M. and Zavodny, M. (2001). The effect of abortion restrictions on the timing of abortions. *Journal of Health Economics*, 20(6):1011–1032.
- Blank, R. M., George, C. C., and London, R. A. (1996). State abortion rates: The impact of policies, providers, politics, demographics, and economic environment. *Journal of Health Economics*, 15(5):513–553.
- Central Bureau of Statistics (Israel) (2018). Religion and Self-Definition of Extent of Religiosity: Selected Data from the Society in Israel. Technical Report 10.
- Colman, S. and Joyce, T. (2009). Minors’ Behavioral Responses to Parental Involvement Laws: Delaying Abortion Until Age 18. *Perspectives on Sexual and Reproductive Health*, 41(2):119–126.
- Colman, S. and Joyce, T. (2011). Regulating Abortion: Impact on Patients and Providers in Texas. *Journal of Policy Analysis and Management*, 30(4):775–797.
- Cook, P. J., Parnell, A. M., Moore, M. J., and Pagnini, D. (1999). The effects of short-term variation in abortion funding on pregnancy outcomes. *Journal of Health Economics*, 18(2):241–257.
- Dickman, S. L., White, K., Sierra, G., and Grossman, D. (2022). Financial Hardships Caused by Out-of-Pocket Abortion Costs in Texas, 2018. *American Journal of Public Health*, 112(5):758–761.
- Doepke, M. and Zilibotti, F. (2017). Parenting With Style: Altruism and Paternalism in Intergenerational Preference Transmission. *Econometrica*, 85(5):1331–1371.
- Doran, F. and Nancarrow, S. (2015). Barriers and facilitators of access to first-trimester abortion services for women in the developed world: A systematic review. *Journal of Family Planning and Reproductive Health Care*, 41(3):170–180.
- Eckstein, Z., Keane, M., and Lifshitz, O. (2019). Career and Family Decisions: Cohorts Born 1935-1975. *Econometrica*, 87(1):217–253.
- Einav, T., Korlinsky, D., and Yifrah, A. (2017). Israel National Health Interview Survey 2013-2015. Technical report, ICDC.
- Fischer, S., Royer, H., and White, C. (2018). The impacts of reduced access to abortion and family planning services on abortions, births, and contraceptive purchases. *Journal of Public Economics*, 167(C):43–68.

- Forsstrom, M. P. (2021). Abortion Costs and Single Parenthood: A Life-Cycle Model of Fertility and Partnership Behavior. *Labour Economics*, 69:101977.
- Foster, A. M., Daoud, F., Abed, S., Sa'da, K. A., and Al-Ayasah, H. (2007). Illicit sex, abortion, and so-called "honor-killings": Attitudes and opinions of female university students in Palestine. In *National Abortion Federation Annual Meeting*.
- Fried, M. G. (2000). Abortion in the United States: Barriers to Access. *Health and Human Rights*, 4(2):174.
- Glenmary Research Center (1974). Churches and Church Membership in the United States, 1971.
- Glenmary Research Center (1982). Churches and Church Membership in the United States, 1980.
- Grossman, D., Grindlay, K., and Burns, B. (2016). Public funding for abortion where broadly legal. *Contraception*, 94(5):453–460.
- Guttmacher (2018). Induced Abortion Worldwide. Technical report, Guttmacher Institute.
- Haas-Wilson, D. (1996). The Impact of State Abortion Restrictions on Minors' Demand for Abortions. *The Journal of Human Resources*, 31(1):140.
- Hoehn-Velasco, L., Dhingra, N., and Pineda-Torres, M. (2025). The Consequences of Abortion Funding Bans. Working Paper 34548, National Bureau of Economic Research, Cambridge, MA.
- Inglehart, R., Haerpfer, C., Moreno, A., Welzel, C., Kizilova, K., Diez-Medrano, J., Lagos, M., Norris, P., Ponarin, E., Puranen, B., et al. (2014). World Values Survey: Round Four - Country-Pooled Datafile.
- Joyce, T. and Kaestner, R. (2001). The Impact of Mandatory Waiting Periods and Parental Consent Laws on the Timing of Abortion and State of Occurrence among Adolescents in Mississippi and South Carolina. *Journal of Policy Analysis and Management*, 20(2):263–282.
- Joyce, T. J., Kaestner, R., and Ward, J. (2020). The Impact of Parental Involvement Laws on the Abortion Rate of Minors. *Demography*, 57(1):323–346.
- Kane, T. J. and Staiger, D. (1996). Teen motherhood and abortion access. *The Quarterly Journal of Economics*, 111(2):467–506.
- Karlin, S. and Rinott, Y. (1980). Classes of orderings of measures and related correlation inequalities. I. Multivariate totally positive distributions. *Journal of Multivariate Analysis*, 10(4):467–498.
- Keane, M. P. and Wolpin, K. I. (2001). The Effect of Parental Transfers and Borrowing Constraints on Educational Attainment. *International Economic Review*, 42(4):1051–1103.
- Kearney, M. S. and Levine, P. B. (2012). Why Is the Teen Birth Rate in the United States So High and Why Does It Matter? *Journal of Economic Perspectives*, 26(2):141–163.
- Kelner, Y. (2013). Basket of medicine 2014: free abortions without medical reasoning. *Ynet*.
- Levine, P. and Staiger, D. (2002). Abortion as Insurance. Technical Report w8813, National Bureau of Economic Research, Cambridge, MA.
- Levine, P. B. (2003). Parental involvement laws and fertility behavior. *Journal of Health Economics*, 22(5):861–878.
- Levine, P. B. and Staiger, D. (2004). Abortion Policy and Fertility Outcomes: The Eastern European Experience. *The Journal of Law and Economics*, 47(1):223–243.
- Levine, P. B., Trainor, A. B., and Zimmerman, D. J. (1996). The effect of Medicaid abortion funding restrictions on abortions, pregnancies and births. *Journal of Health Economics*, 15:555–578.
- Lindo, J., Myers, C., Schlosser, A., and Cunningham, S. (2020). How Far Is Too Far? New Evidence on Abortion Clinic Closures, Access, and Abortions. *Journal of Human Resources*, 55(4):1137–1160.
- Lindo, J. M. and Pineda-Torres, M. (2019). New Evidence on the Effects of Mandatory Waiting Periods for Abortion. Working Paper 26228, National Bureau of Economic Research.
- Lochner, L. J. and Monge-Naranjo, A. (2011). The Nature of Credit Constraints and Human Capital. *American Economic Review*, 101(6):2487–2529.
- Londoño-Vélez, J. and Saravia, E. (2025). The Impact of Being Denied a Wanted Abortion on Women and Their Children. *The Quarterly Journal of Economics*, 140(2):1061–1110.

- Lundberg, S., Romich, J. L., and Tsang, K. P. (2009). Decision-making by children. *Review of Economics of the Household*, 7(1):1–30.
- MacBride, J. A. (2001). Parental Consent Legislation Benefits Minors Seeking Abortions. *Buffalo Women's Law Journal*, 10(1):75–92.
- Myers, C. and Ladd, D. (2020). Did parental involvement laws grow teeth? The effects of state restrictions on minors' access to abortion. *Journal of Health Economics*, 71:102302.
- Myers, C. K. (2017). The Power of Abortion Policy: Reexamining the Effects of Young Women's Access to Reproductive Control. *Journal of Political Economy*, 125(6):2178–2224.
- Myers, C. K., Dench, D. L., and Pineda-Torres, M. (2025). The road not taken: How driving distance and appointment availability shape the effects of abortion bans. Technical Report 33548, NBER.
- New, M. J. (2011). Analyzing the effectiveness of parental involvement laws in the United States. *Technical Report*.
- Newman, M. (2017). 15,000 illegal abortions performed in Israel each year, activists claim. *The Times of Israel*.
- Oberman, M. (2021). How Abortion Laws Do and Don't Work. *Wisconsin Journal of Law, Gender & Society*, 36(2):163.
- Pollak, R. A. (1988). Tied Transfers and Paternalistic Preferences. *American Economic Review*, 78(2):240–244.
- Rambachan, A. and Roth, J. (2023). A More Credible Approach to Parallel Trends. *Review of Economic Studies*, 90(5):2555–2591.
- Rimalt, N. (2017). When Rights Don't Talk: Abortion Law and the Politics of Compromise. *Yale Journal of Law and Feminism*, 28:55.
- Rosenberg, H., Blondheim, M., and Katz, E. (2019). It's the text, stupid! Mobile phones, religious communities, and the silent threat of text messages. *New Media & Society*, 21(11-12):2325–2346.
- Sade, Y. (2023). The top deciles bear most of the burden of serving in the army. *Calcalist*.
- Shaked, M. and Shanthikumar, J. G. (2007). *Stochastic Orders*. Springer, New York.
- Shapiro, G. K. (2014). Abortion law in Muslim-majority countries: An overview of the Islamic discourse with policy implications. *Health Policy and Planning*, 29(4):483–494.
- Singh, S., Remez, L., Sedgh, G., Kwok, L., and Onda, T. (2018). Abortion Worldwide 2017: Uneven Progress and Unequal Access. Technical report, Guttmacher Institute.
- Tunick, M. (2023). State Authority, Parental Authority, and the Rights of Mature Minors. *The Journal of Ethics*, 27(1):7–29.
- Upadhyay, U. D. (2022). Barriers Push People into Seeking Abortion Care Later in Pregnancy. *American Journal of Public Health*, 112(9):1280–1281.
- U.S. Government Accountability Office (2025). Abortion access: State restrictions and the economic impact on women and families. *GAO Technical Report*.
- Venator, J. and Fletcher, J. (2021). Undue Burden Beyond Texas: An Analysis of Abortion Clinic Closures, Births, and Abortions in Wisconsin. *Journal of Policy Analysis and Management*, 40(3):774–813.
- Zionov, A. (2021). Living Arrangements of Young Adults in Israel in 2018. Technical Report 070, Israel Central Bureau of Statistics.


A Additional Figures and Tables

Figure A.1: Abortion Ratios in Israel by Ethno-Religious Group



Notes: The figure reports abortion ratios (the percentage of pregnancies that end in abortion) for two populations: all women aged 16–45 who conceived between 2009–2013 (left panel), and the analytic sample—defined in Section 3.1—of unmarried women aged 16–23 who conceived for the first time during the same period (right panel). Within each population, observations are grouped by ethno-religious category: all women, secular Jews, religious Jews (including Ultra-Orthodox Jews), and Israeli Arabs.

Figure A.2: Full Sample Dynamic Difference-in-Differences Estimates



Notes: The figure plots the dynamic difference-in-differences coefficients from Equation (2), which measure changes in abortion ratios between treated women (ages 20–23) and control women (ages 16–19) over 2009–2016. Estimates are pooled across all communities, without separating locations by abortion attitudes. Each point corresponds to a coefficient δ_k ; 2013 is the omitted year. The vertical dashed line marks the 2014 subsidy reform. The sample includes all unmarried women aged 16–23 who conceived for the first time between 2009–2016. All regressions include month fixed effects and the interaction of ethno-religious-group and education fixed effects. Standard errors are clustered by age-by-year at conception, and shaded regions depict 95% confidence intervals.

Table A.1: Descriptive Statistics by Socioeconomic Status

Variable	Low SES	High SES
Daughter's Income (Shekels)	15,652	18,121
HH Income (Shekels)	60,224	280,643
Father's Income (Shekels)	50,283	190,467
Father's Months Worked	9	11
Post-secondary educ. (>12 yrs)	0.08	0.14
Religious Jew (share)	0.08	0.06
Arab (share)	0.10	0.02

Notes: This table reports baseline characteristics for women from low- and high-SES families. Socioeconomic status is defined based on whether household earnings fall below or above the median in our sample. Reported variables include the daughter's income, combined parental income, the father's income, the father's months worked in the past year, the share with post-secondary education (more than 12 years of schooling), and indicators for religious Jewish and Arab background. All earnings measures correspond to the year of conception.

Table A.2: Effect of the 2014 Reform by Socioeconomic Status and by Daughter's Own Earnings

<i>Panel A: By Socioeconomic Status</i>				
	Lenient		Strict	
	Low SES	High SES	Low SES	High SES
Treatment Effect	0.058*** (0.014) [-0.07, 0.03]	0.052*** (0.014) [-0.06, 0.03]	0.141*** (0.033) [0.09, 0.35]	0.142*** (0.047) [-0.17, 0.27]
P-value (SES)	0.770		0.984	
P-value (Attitudes)			0.020	0.067
Baseline Mean	0.704	0.844	0.442	0.681
N	17,783	15,543	3,775	1,487
<i>Panel B: By Daughter's Own Earnings</i>				
	Lenient		Strict	
	Does not work	Works	Does not work	Works
Treatment Effect	0.054*** (0.012) [-0.04, 0.04]	0.041** (0.017) [-0.12, 0.08]	0.173*** (0.029) [0.01, 0.26]	-0.016 (0.056) [-0.4, 0.32]
P-value (Works)	0.551		0.003	
P-value (Attitudes)			0.000	0.326
Baseline Mean	0.786	0.736	0.479	0.597
N	25,531	9,700	3,908	1,356

Notes: This table reports difference-in-differences estimates of the effect of the 2014 reform on abortion ratios within communities with lenient and strict attitudes toward abortion. Panel A further splits each attitude group by the family's socioeconomic status (low and high SES); Panel B splits it by the daughter's employment status in the year before conception, defined as in Table 1: "Works" columns include daughters whose own earnings exceed six months of the minimum wage, and "Does not work" columns include all others. The sample consists of all unmarried women aged 16–23 who became pregnant for the first time between 2009–2016. "Lenient" and "Strict" communities are defined as in Table 1. In Panel A, "P-value (SES)" tests the equality of treatment effects between low- and high-SES families within each strictness group, and "P-value (Attitudes)" tests the equality between lenient and strict communities within each SES group. In Panel B, "P-value (Works)" and "P-value (Attitudes)" report the analogous tests by work status. All regressions include month fixed effects and the interaction of ethno-religious-group and education fixed effects. Standard errors, clustered by age-by-year at conception, are in parentheses. Upper and lower bounds that account for differential pre-trends appear in brackets. * $p < 0.10$, ** $p < 0.05$, and *** $p < 0.01$.

Table A.3: Subsidy Effect on Abortion Utilization By Ethno-Religious Group

	Lenient	All Strict	All Strict Jews	Religious Jews	Ultra-Orthodox	Arabs
Treatment Effect	0.056*** (0.010)	0.148*** (0.027)	0.150*** (0.038)	0.131*** (0.041)	0.164* (0.090)	0.129*** (0.036)
P-value (vs. lenient)			0.018	0.078	0.232	0.052
P-value (vs. all strict)			0.961	0.729	0.863	0.671
Baseline Mean	0.773	0.509	0.630	0.705	0.453	0.372
N	35,231	5,264	2,802	1,966	836	2,462

Notes: This table reports OLS estimates of Equation 1, which identifies the effect of the 2014 subsidy reform on abortion ratios for women aged 16–23 who became pregnant for the first time between 2009–2016. Column (1) reports estimates for women from lenient backgrounds. Column (2) reports estimates for all women from strict communities (religious Jews, Ultra-Orthodox Jews, and Israeli Arabs). Columns (3)–(6) further disaggregate strict communities into: religious and Ultra-Orthodox Jews (Column 3), religious Jews only (Column 4), Ultra-Orthodox Jews only (Column 5), and Israeli Arabs (Column 6). All regressions include month fixed effects and the interaction of ethno-religious-group and education fixed effects. Standard errors, clustered by age-by-year at conception, are reported in parentheses. Reported p-values test differences in treatment effects between each subgroup and either Column (1) (lenient) or Column (2) (all strict), as indicated. * $p < 0.10$, ** $p < 0.05$, and *** $p < 0.01$.

Table A.4: Subsidy Effect Heterogeneity by Location Leniency Share

	Level Effect		Percent Effect	
	(1)	(2)	(3)	(4)
Lenient Share	-0.260*** (0.064)	-0.302*** (0.079)	-82.500*** (15.899)	-84.607*** (20.213)
HH Income	X	X	X	X
Daughter's Income		X		X
Daughter Working Share		X		X
N	23	23	23	23

Notes: Each column reports results from a regression of the location-specific treatment effect from Equation (1) on the share of lenient individuals in the location, estimated on the 23 locations with at least 100 pre-reform pregnancies. “Level Effect” (Columns (1)–(2)) refers to the estimated impact in percentage points; “Percent Effect” (Columns (3)–(4)) scales this impact by the location’s pre-reform abortion rate. Columns (2) and (4) add daughters’ own average income and daughters’ average employment rate (employment defined as in Table 1) on top of household income. Both income variables are measured in 10,000 NIS. X indicates that the control is included. All regressions weight locations by the number of pre-reform pregnancies. Standard errors in parentheses. * $p < 0.10$, ** $p < 0.05$, and *** $p < 0.01$.

B Israeli Context

B.1 Abortion and Contraceptive Use: Norms and Prevalence

Prevalence of Abortion. Abortion is common in Israel despite the committee approval requirement. In the pre-policy period (2009–2013) abortions accounted for about 10% of pregnancies. Although this may appear high, Israel’s legal abortion ratio is relatively low compared to international benchmarks (Figure B.1). Globally, 25% of pregnancies end in abortion; in Europe the rate is 26%, and in North America 16% (Guttmacher, 2018).

Public Discourse. Public engagement with abortion policy in Israel is limited compared to the United States. A Google Trends comparison (Figure B.2) shows that search activity for the term “abortion” (*hapala*) in Hebrew is much less reactive to events than in the U.S. The modest rise in 2014 coincides with the implementation of the cost-elimination reform but remains muted relative to U.S. peaks, such as during President Trump’s election in 2016 or the Supreme Court confirmation hearings for Judge Kavanaugh. This pattern underscores the low salience of abortion in Israeli public debate.

Political Salience. Abortion is not as polarizing a political issue in Israel and is often described as a “silenced phenomenon” (Amir, 2015). Liberal parties occasionally criticize the committee system but avoid pursuing reform, fearing that renewed debate could lead to more restrictive laws (Oberman, 2021; Rimalt, 2017). Two legislative attempts to expand or abolish the committee process—in 2004 and 2006—did not pass. Religious parties have likewise failed to tighten restrictions. Bills introduced in 2008 and 2017 sought to prohibit late-term abortions or add a religious representative to the committee but did not gain sufficient support.


Cultural and Religious Attitudes. Israel’s population is religiously diverse: 75% Jewish, 18.6% Muslim, 2% Christian, and 4.4% other or unaffiliated. As shown in Figure A.1, abortion ratios vary widely across groups, reflecting heterogeneity in social norms and perceived moral costs. These differences are central to understanding variation in young women’s reproductive autonomy.

Judaism generally holds more permissive views on abortion than Christianity or Islam, prioritizing the mother’s life and health (Amir, 2015). Abortions are permitted in two broad cases: (i) risk to the woman’s life, or (ii) if the child would be born into an “unstable life,” which may include cases such as unmarried parents, very young or old mothers, or severe congenital disorders. Despite this relative permissiveness, Israel’s “demographic project” aims to preserve high Jewish fertility, creating a tension between religious interpretation and policy practice.

Religiosity within the Jewish population is strongly correlated with marital patterns, fertility, contraceptive use, and abortion attitudes. Secular Jews (about 45% of the Jewish population) generally support abortion, have high contraceptive use rates, and low fertility. In contrast, religious and Ultra-Orthodox Jews (about 30% combined) strongly oppose abortion, use contraception rarely (Figure B.3), and have high fertility. Marriage and childbearing also occur earlier among religious women: about half of them are married by the age of 23, and four out of five of those who conceived were married.

Muslim and Arab Populations. The Israeli-Arab population, predominantly Muslim, views abortion as highly taboo. Among Muslims, 11% identify as secular, 57% as traditional, and 31% as religious (Central Bureau of Statistics (Israel), 2018). Islam generally prohibits abortion except to save the mother's life, which may lead some women to seek abortions outside the legal system. Reliable data on illegal or self-induced abortions in Israel are scarce, but evidence from neighboring populations offers context. A 2006 survey of Palestinian women found that 10% had self-induced an abortion, and that 25% of respondents believed self-induced abortions are necessary for unmarried women to avoid honor-based violence (Foster et al., 2007).⁴⁸ These findings highlight the strong social stigma surrounding abortion in Arab communities.


Figure B.1: Abortion Ratios Worldwide



Notes: The figure reports abortion ratios (abortions per pregnancy) among women aged 15–49 by region. Global data: Guttmacher Institute (<https://data.guttmacher.org/>); Israeli data: Central Bureau of Statistics.


⁴⁸“Honor killings” refer to the practice of family members harming women accused of bringing dishonor to the family through premarital or extramarital sex.

Figure B.2: Google Searches for “Abortion” (Israel vs. U.S.)



Notes: The figure compares search intensity for “abortion” in the U.S. and its Hebrew equivalent (*hapala*) in Israel, 2009–2019. Search levels are normalized to January 2009.

Figure B.3: Self-Reported Contraceptive Use by Religiosity



Source: Israel National Health Interview Survey (2013–2015)

Notes: The figure reports self-reported contraceptive use by religiosity, using the Israel National Health Interview Survey (2013–2015) (Einav et al., 2017). The data include Jewish and Muslim women by self-identified religiosity. No Ultra-Orthodox women in the survey reported using contraception.

B.2 Abortion Committee

Israel's abortion committee system is exceptional in the global context and reflects the country's explicit demographic agenda. Officially, the committee process was introduced due to medical concerns that abortion might harm women's future fertility (Amir, 2015). This concern aligned with Israel's broader *demographic project*—an effort to reverse the decline of the global Jewish population following the Holocaust. In pursuit of this project, Israel has adopted extensive pro-natalist policies, including subsidized daycare, monthly child allowances, paid parental leave, and public coverage of infertility treatments. In contrast, contraception is not covered by national health insurance, and abortions are illegal without prior approval from an abortion committee. Consistent with this policy framework, Israel maintains the highest fertility rate among OECD countries.

When a woman seeks an abortion, her doctor refers her to one of 42 abortion committees operating across hospitals and clinics nationwide (See the full list of committees [here](#)). Each committee comprises two medical professionals and a social worker, at least one of whom must be a woman. All abortions, including those performed by private doctors, require committee authorization.

The committee grants approval if *at least one* of the following criteria is met (see rates in Table [B.1a](#)):

1. The woman is under 18 or over 40 years old;
2. The pregnancy occurred outside marriage;
3. The pregnancy resulted from rape or incest;
4. The pregnancy endangers the woman's life or health (including mental health); or
5. The fetus has a congenital disorder.

These criteria are grounded in Jewish law, which prioritizes the mother's life and health and considers the fetus part of her body until birth. Jewish law also views children born out of wedlock as socially disadvantaged, providing religious justification for abortion in such cases (Amir, 2015).

At the committee, women complete paperwork and pay a 400 NIS (\$112 in 2014 rates) fee, which was also eliminated by the 2014 policy. They first meet with the social worker, who evaluates eligibility and provides counseling. In practice, social workers act as the effective gatekeepers, while committees generally serve as a formality (Oberman, 2021). Only married women aged 18–40 with healthy pregnancies are formally ineligible. In these cases, social workers often help women secure psychiatric documentation stating emotional distress or adjustment difficulties, which qualifies under the “risk to woman's health” criterion (Oberman, 2021).⁴⁹

⁴⁹Alternatively, some women report medication use that could endanger the fetus.

Although the system appears restrictive, nearly all applications are approved. In our data, 99% of applications were approved and 97% acted upon. The process is confidential, and neither partner nor parental consent is required. Late-term cases (beyond 24 weeks, about 1% of abortions) are reviewed by a special committee with stricter standards, which we exclude from analysis. High approval rates likely reflect the social workers' gatekeeping role, rather than leniency in policy. Alternative explanations—such as cross-border “abortion tourism”—are implausible given Israel's geopolitical isolation and the more restrictive abortion laws in neighboring countries.

To reduce financial barriers, the government progressively expanded abortion subsidies: first in 2001 to women under 18, then in 2008 to those aged 19 or younger, and in 2014 to women aged 20–32 under any approval criterion (see the 2014 change in Table [B.1b](#)).

B.3 Childrearing Cost in Israel

The cost of childrearing in Israel is substantially lower than in most other high-income countries, particularly the United States. Education and healthcare are publicly funded: healthcare coverage for children costs approximately \$3 per month, and pre-kindergarten is free for children aged three and older. Even below age three, public or subsidized daycare options are widely available and typically cost less than a full-time minimum wage salary.⁵⁰

As part of its pro-natalist demographic agenda, the government provides several direct and indirect financial supports. These include a one-time birth grant (USD 145–484 depending on birth order), a monthly child allowance through social security (USD 41–52 per child), and a tax credit for working parents (USD 118 per month per child). Together, these transfers substantially reduce the monetary cost of raising children, especially for low- and middle-income households.

Parental leave policies further mitigate the opportunity cost of childbearing. By law, mothers are entitled to 14 weeks of paid maternity leave, fully financed by the government, with the option to extend for an additional 12 weeks unpaid. During this period, employers may not terminate the parent's employment. Fathers are formally eligible to share the leave, though uptake remains rare. Some employers voluntarily grant extended unpaid leave beyond the statutory period, but they are not required to do so.

Israel's small geographic size and strong intergenerational family networks also facilitate mothers' reentry into the labor market. Many parents rely on grandparents or extended family for childcare during the child's first year, before entry into the public education system. These patterns are particularly pronounced among religious families, where women typically marry and have children at younger ages and combine employment with childrearing (Figure [B.4](#)).

⁵⁰Israel's minimum full-time monthly wage in 2014 was NIS 4,300 (USD 1,200).

Table B.1: Eligibility Criteria for Abortions and Subsidies

(a) Abortion Eligibility Criteria and Pre-2014 Subsidies

Eligibility Criteria for Abortion	Share of Approvals by Criteria	Free Pre-2014
Out of Marriage or Illegal Act	50.3%	X ✓
Risk for Woman or Fetus	40.6%	✓ ✓
Age < 18 or Age ≥ 40	9%	✓ X


Notes: This table shows the eligibility criteria (column 1) for obtaining a legal abortion in Israel and the proportion of applications that are approved by the committee for each criterion (column 2). In the third column, we show the eligibility criterion for a subsidized abortion pre-2014. While “out-of-marriage” and illegal act are both under the same eligibility criteria, only abortions approved due to an illegal act were subsidized prior to 2014.

(b) 2014 Change in Abortion Subsidy (Identification Strategy)

Age	Free?	
	Pre-2014	Post-2014
Age ≤ 19	✓	✓
19 < Age < 33	X	✓
Age ≥ 33	X	X

Notes: This table highlights the change in eligibility for a fully subsidized abortion following the 2014 policy, which serves as a natural experiment for this paper. Women aged 19 and under were already fully subsidized by the government and therefore unaffected by the change and women age 33 and older were not included in the subsidy expansion and thus never treated. This change in funding applies to women aged 20-32 regardless of what criteria their abortion was approved under, but as can be seen in Table B.1a, of the potential criteria that apply to women aged 20-32, the out-of-marriage criterion is the only one not eligible for a subsidy prior to 2014.

Figure B.4: Life-Cycle Decisions of Women in Strict Communities



Notes: The figure plots four life-cycle outcomes—marriage, parenthood, employment, and higher education—by age, for women in strict communities (religious Jewish and Arab women, pooled and population-weighted), ages 16–30. Each point is a *cross-sectional share*: the fraction of women at a given age who, in that year, are currently married, are a parent, are working, or are studying, conditional on being unmarried at the time of conception. The marriage and parenthood series are not cumulative distribution functions—a woman’s status can change from year to year, so the shares need not be monotonic in age (for example, the share currently married declines slightly at older ages as the unmarried-at-baseline cohort thins out). Panel (a) shows family formation (marriage and parenthood); panel (b) shows employment and higher education; the series labeled “Part-time Emp. (Uncond.)” is the share of all women at that age (not only of those working) who are employed part-time.

C Sample Restrictions

Unmarried Women. We restrict the analysis to unmarried women for both conceptual and empirical reasons. First, the 2014 policy change affected only pregnancies outside marriage, since funding coverage for married women did not change; unmarried women are therefore the group whose financing was actually altered by the reform. Second, pregnancies outside marriage are automatically approved for legal abortion (Table B.1a), whereas married women obtain approval only under specific circumstances (e.g., health risk or infidelity) and are far less likely to terminate a pregnancy (3.9%). Estimated effects are similar when married women are included in the sample (Table D.7). Marital status is measured at the month of conception.

First Pregnancies. To avoid endogeneity in sample construction, we restrict the analysis to each woman's first observed pregnancy. This approach ensures that treatment assignment is not affected by previous fertility or abortion decisions, which could themselves depend on earlier exposure to the policy.

Age Range. Our primary sample includes unmarried women aged 16–23 at the time of conception. This range captures those most likely to be directly affected by the 2014 policy, which expanded full funding for abortions from age 19 to 32. In robustness analyses, we focus on a narrower 18–21 age band to align more closely with the lower age cutoff and maximize comparability around the policy threshold. Younger women are also the most relevant group for testing our model's predictions, as they face higher financial and social constraints on abortion access.

A potential concern is that women aged 18–19 are often serving in the military, 20-year-olds may still be serving, and those aged 21–22 have typically completed service (Sade, 2023). Because military service might influence fertility decisions, we conduct robustness checks excluding women aged 20 and re-defining the treated group as 21–22; results are unchanged.

Time Period. We restrict the sample to conceptions between January 2009 and March 2016. Starting in 2009 avoids contamination from an earlier policy change that expanded abortion funding to women aged 19, since before then coverage varied by military status.

D Subsidies' Effect on Abortion Decisions: Robustness

D.1 Changes in Sample Composition

Treatment and comparison groups may have undergone differential compositional changes around the time of the reform. To assess this possibility, we estimate the following simple difference-in-differences specification:

$$outcome_{it} = \beta (Post_t \times T_i) + \gamma_{a_i} + \gamma_{y_t} + \epsilon_{it} , \quad (9)$$

where the outcome is a baseline characteristic of woman i : her parents' household income in the year of conception, her own income and months worked in the years before conception, an indicator for post-secondary education, or an ethno-religious-group indicator. Equation (9) is analogous to our main specification (Equation 1), but excludes controls to focus on compositional changes.

Results are presented in Table D.2. Panel A reports estimates for all unmarried pregnant women aged 16–23, while Panels B and C show results separately for women from strict and lenient backgrounds. The estimates indicate that treatment and control groups are balanced across most observable characteristics, with the exception of education.

To assess whether this imbalance affects our main estimates, we re-estimate Equation (1)—our primary specification for identifying the effects of the subsidy—under alternative sets of controls and examine the sensitivity of the results. Table D.3 reports results for the full analytic sample, while Tables D.4 and D.5 present results for the lenient and strict subsamples, respectively. Column (1) presents estimates without controls. Columns (2)–(4) sequentially add each component of our baseline control set, while Column (5) includes all baseline controls jointly. Columns (6)–(8) introduce additional controls, including location fixed effects, interactions between ethno-religious group and age or year, an indicator for high SES, and an indicator for employment status (defined as in Table 1). Column (9) reports an additional robustness check using a triple-differences design that exploits married women—who were unaffected by the reform—as a control group.

Across all specifications, the estimated effects remain stable and statistically indistinguishable from the baseline (see p -values in Tables D.3–D.5). These findings strengthen our confidence that the results are not driven by compositional changes across cohorts. To be conservative, we adopt the specification with education-ethno-religious-group fixed effects as our preferred specification throughout the paper.

D.2 Accounting for Pre-Trends

The estimates reported in Table 1 rely on the standard parallel-trends assumption between treatment and control groups. However, Figure 2 suggests a modest differential pre-trend for the lenient population. For the strict population, although pre-trends are not apparent, the relatively large standard

errors prevent us from fully ruling out their presence.

To account for possible deviations from parallel trends, we apply the methodology of Rambachan and Roth (2023), which provides bounds on treatment effects that account for differential trends estimated from pre-treatment data. Their approach incorporates statistical uncertainty in the estimation of those trends. The resulting bounds are reported in brackets in Table 1.

For the lenient population, the bounds include zero, reinforcing that we cannot statistically distinguish between no effect and a small positive effect. In contrast, for the strict population, the bounds confirm a robust positive effect even when allowing for potential pre-trends. Moreover, the two sets of bounds do not overlap, providing further evidence that the subsidy's impact was significantly larger for the strict population.

As an additional robustness check, we relax the assumption that pre- and post-treatment trends evolve at the same rate by allowing the slope of the differential trend to vary by up to 4% before and after the reform, following the range proposed by Rambachan and Roth (2023). The results, shown in Figure D.1, again confirm our main conclusion: among women from strict backgrounds, we can reject the null of no effect across all specifications within the considered range.

D.3 Effect of Subsidy on Pregnancy

By reducing the cost of abortion, the subsidy also reduced the cost of becoming pregnant. This could, in principle, lower incentives for contraceptive use and thereby increase pregnancy rates. Such an effect would bias our estimates, since our main analysis conditions on being pregnant.

To examine this possibility, we re-estimate Equation (1) using all unmarried women aged 16–23 between 2009 and 2016, without conditioning on pregnancy. We consider two outcomes: (i) an indicator for whether a woman conceived in that year, and (ii) an indicator for whether she both conceived and obtained an abortion.

Results are reported in Table D.1. The estimates for pregnancy show no evidence that the subsidy increased conception rates—the coefficients are negative and statistically insignificant. The abortion estimates, while noisier due to the inclusion of non-pregnant women (who were unaffected by the policy), still reveal a statistically significant increase in abortions among women from strict backgrounds, consistent with our main findings. For women from lenient backgrounds, the effects are statistically indistinguishable from zero.

The null effect on pregnancy also alleviates concerns that the subsidy induced women who would otherwise obtain an illegal abortion to obtain a legal one. Legal abortions are observed in the administrative data and therefore appear as pregnancies, whereas illegal abortions are not recorded. If substantial substitution from illegal to legal abortions had occurred, we would expect to see an increase in observed pregnancies. We find no such increase.

D.4 Age bandwidth

In our main specification, women in the treatment group (ages 20–23) are older than those in the control group (ages 16–19). Although we include age fixed effects, differential age-specific trends could still bias our estimates. To assess this concern, Table D.6 presents estimates of Equation (1) using alternative age bandwidths. First, we restrict the sample to women aged 18–21 (instead of 16–23), ensuring that treated and control groups are more comparable in age. While standard errors increase slightly due to the smaller sample, point estimates remain stable.

Another potential concern is that 20-year-old secular Jewish women are likely to still be in military service, which fully covers all healthcare including abortion. As a robustness check, we estimate a specification using 18–19-year-olds as the control group and 21–22-year-olds as the treatment group, excluding 20-year-olds. The results remain virtually unchanged. Overall, the estimates do not differ significantly across specifications, as shown by the p -values in Table D.6. Importantly, across all bandwidths, the conclusion that the policy’s effect is substantially larger for women from strict backgrounds remains robust.

D.5 Selection into Marriage

Our main analysis focuses on unmarried women who conceived between 2009 and 2016. If the policy influenced marriage decisions, this restriction could introduce selection bias. To assess this possibility, Table D.7 reports estimates from Equation (1) without excluding married women.

In levels, coefficients appear larger for lenient-background women when married women are included. However, this difference primarily reflects disparities in baseline abortion ratios: among all strict-background women, only 3.2% of pregnancies end in abortion, compared to 31.3% among all lenient-background women. When expressing the treatment effects as percentages of baseline means, the results are consistent across specifications. For strict-background women, the effect equals 23.3% of the baseline when including all pregnancies and 29.1% when focusing on unmarried women; for lenient-background women, the corresponding figures are 4.4% and 7.2%. Moreover, the policy effect for the strict population ($p < 0.01$) and the difference in percent effects between strict and lenient populations ($p = 0.03$) are both statistically significant.

Overall, the results are consistent with those obtained from the unmarried-only sample (Table 1), indicating that selection into marriage is unlikely to drive our findings.

D.6 Former-USSR Migration

The pre-trends observed among women from lenient backgrounds may partly reflect demographic shifts associated with the large influx of migrants from the former USSR in the early 1990s. To assess this

possibility, we exclude women whose families originated from the former USSR and re-estimate Equations (2) and (1). Figure D.2 compares estimated trends for all unmarried pregnant women from lenient backgrounds with those obtained after excluding this group. Removing the former-USSR-origin population attenuates the pre-trends but does not eliminate them entirely. Table D.8 shows that the estimated effect of the subsidy remains robust. Moreover, when accounting for pre-trends using the approach of Rambachan and Roth (2023) (bounds reported in brackets in Table D.8), our conclusion remains unchanged: the subsidy had little to no effect on abortion rates among women from lenient backgrounds.

D.7 Effects by Approval Reason

In Israel, abortions are legally approved if at least one of the following conditions is met: (1) the woman is under 18 or over 40; (2) the pregnancy occurred outside marriage; (3) the pregnancy resulted from rape or incest; (4) the pregnancy endangers the woman's life or health; or (5) the fetus has a congenital disorder (see Table B.1a). Before 2014, only women aged 19 or younger and those meeting medical or legal criteria were eligible for fully subsidized abortions. Therefore, the 2014 reform essentially extended eligibility to out-of-marriage abortions for women between the ages of 20 and 32. Because health-related abortions require a more stringent approval process, the reform likely induced some substitution from health-related to out-of-marriage approvals.


Figure D.3 confirms this pattern: a large rise in out-of-marriage approvals coincides with a decline in health-related ones. This evidence supports the interpretation that the overall increase in abortions after the reform was driven by the policy itself rather than by confounding trends.

Note that our main results do not distinguish between approval reasons. Thus, they capture the net increase in the overall abortion ratio, accounting for this reallocation across approval categories.

D.8 Taking Stock


This appendix has presented a series of robustness analyses of our main findings. Taken together, the results reveal a consistent pattern: the subsidy had a large and statistically significant positive effect on abortion rates among women from strict backgrounds. For women from lenient backgrounds, the evidence is less consistent across specifications, and we cannot empirically distinguish between no effect and a small positive effect.

Figure D.1: Bounds on the Effect of the Subsidy on Abortion Utilization




Notes: This figure displays bounds on the treatment effect of the 2014 subsidy reform constructed using the method of Rambachan and Roth (2023), which allows for potential differential pre-trends between treated and control groups. Bounds are shown separately for women from lenient and strict backgrounds. The x-axis reports the maximum allowed deviation of the post-treatment trend from the pre-treatment trend, up to 0.04 following Rambachan and Roth (2023), and the y-axis reports the corresponding treatment-effect bounds.

Figure D.2: Dynamic Difference-in-Differences Excluding Former-Soviet-Union Origin



Notes: The figure plots the dynamic difference-in-differences coefficients from Equation 2, which measure changes in abortion ratios between treated women (ages 20–23) and control women (ages 16–19) over 2009–2016. The sample consists of unmarried women aged 16–23 who conceived for the first time during this period. Estimates are shown for all women from lenient backgrounds and for the subsample excluding women from families originating in the Former Soviet Union. Each point corresponds to a coefficient δ_k ; 2013 is the omitted year. The vertical dashed line marks the 2014 subsidy reform. All regressions include month fixed effects and the interaction of ethno-religious-group and education fixed effects. Standard errors are clustered by age-by-year at conception, and shaded regions depict 95% confidence intervals.

Figure D.3: Impact of the 2014 Subsidy Expansion on Abortion Approval by Reason



Notes: This figure reports OLS estimates of Equation 1 separately for each reason under which an abortion may be approved by the committee. The sample consists of unmarried women aged 16–23 who became pregnant for the first time between 2009–2016. Estimates are shown for three groups: all women, women from lenient-attitude families, and women from strict-attitude families. All regressions include month fixed effects and the interaction of ethno-religious-group and education fixed effects. Standard errors are clustered by age-by-year at conception.

Table D.1: Impact of the 2014 Subsidy Expansion on Abortion and Pregnancy

	Abortion			Pregnancy		
	Full Sample	Lenient	Strict	Full Sample	Lenient	Strict
Treatment Effect	-0.00064 (0.00183)	-0.00890 (0.00693)	0.00091* (0.00055)	-0.00335 (0.00329)	-0.00479 (0.00735)	-0.00459 (0.00297)
P-value (levels)		0.158			0.98	
P-value (%)		0.059			0.384	
Baseline Mean	0.02908	0.05417	0.00861	0.04673	0.07656	0.02238
N	1,268,430	570,014	698,416	1,268,430	570,014	698,416

Notes: This table reports OLS estimates of Equation 1 using all unmarried women aged 16–23 observed between 2009–2016. Outcomes are indicators for obtaining an abortion (left panel) and for becoming pregnant (right panel). Unlike the main analysis, which conditions on pregnancy, this table uses the population of all unmarried women; abortion and pregnancy outcomes therefore represent unconditional probabilities. For each outcome, Column (1) reports results for the full sample, and Columns (2) and (3) report estimates separately for women from lenient and strict backgrounds. P-values test whether the treatment effects differ between lenient and strict groups, with the first row comparing the level effects (percentage points) and the second row comparing percent effects (scaling by each group’s baseline mean). All regressions include month fixed effects and the interaction of ethno-religious-group and education fixed effects. Standard errors, clustered by age-by-year at conception, are reported in parentheses. * $p < 0.10$, ** $p < 0.05$, and *** $p < 0.01$.

Table D.2: Testing for Composition Changes Around the 2014 Reform

Panel A: Full Sample (Unmarried, 16-23 year-olds)

	HH Income (log)	HH Income	Own Income	Months Worked	Post-secondary educ. (>12 yrs)	Religious Jew	Arab
Treatment Effect	0.004 (0.026)	4871.871 (3229.873)	319.649 (310.634)	0.187 (0.204)	-0.052*** (0.010)	-0.009** (0.005)	-0.003 (0.006)
Baseline Mean	11.592	175525.230	10798.972	11.193	0.135	0.068	0.065
N	32,619	32,619	37,148	37,148	40,495	41,371	41,371

Panel B: Strict Abortion Attitudes (Unmarried, 16-23 year-olds)

	HH Income (log)	HH Income	Own Income	Months Worked	Post-secondary educ. (>12 yrs)
Treatment Effect	-0.136 (0.089)	-3502.203 (7595.393)	154.942 (503.960)	-0.524 (0.426)	-0.066*** (0.012)
Baseline Mean	11.237	130453.537	10365.622	10.389	0.116
N	3,886	3,886	5,433	5,433	5,264

Panel C: Lenient Abortion Attitudes (Unmarried, 16-23 year-olds)

	HH Income (log)	HH Income	Own Income	Months Worked	Post-secondary educ. (>12 yrs)
Treatment Effect	0.030 (0.025)	6991.934** (3335.005)	330.914 (322.496)	0.259 (0.219)	-0.052*** (0.011)
Baseline Mean	11.640	181620.959	10873.207	11.331	0.138
N	28,733	28,733	31,715	31,715	35,231

Notes: This table reports OLS estimates of Equation 9, which tests for changes in the composition of baseline characteristics around the 2014 subsidy reform. Outcomes include household income in the year of conception (in logs and in levels), the woman's own income and months worked in the years before conception, an indicator for post-secondary education (more than 12 years of schooling), and indicators for being a religious Jew or an Israeli Arab. Panel A reports results for all unmarried women aged 16–23 who conceived for the first time between 2009–2016, while Panels B and C report results separately for women from strict and lenient backgrounds, respectively. All regressions include age and year fixed effects. Standard errors are clustered at the age-by-year at conception level and reported in parentheses. * $p < 0.10$, ** $p < 0.05$, and *** $p < 0.01$.

Table D.3: Robustness of the Estimated Subsidy Effect to Sequentially Adding Controls (Full Sample)

	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)
Treatment Effect	0.068*** (0.014)	0.069*** (0.014)	0.065*** (0.012)	0.071*** (0.011)	0.068*** (0.011)	0.062*** (0.010)	0.076*** (0.011)	0.076*** (0.012)	0.092*** (0.031)
Baseline Mean	0.723	0.723	0.723	0.738	0.738	0.738	0.738	0.734	0.132
P-value		0.969	0.887	0.844	0.998	0.73	0.633	0.639	0.487
N	41,371	41,371	41,371	40,495	40,495	40,427	40,495	38,588	342,822
Month FE		X			X	X	X	X	X
Ethno-Relig FE			X		X	X	X	X	X
Education FE				X	X	X	X	X	X
Ethno-Relig $\times$ Education FE					X	X	X	X	X
Location FE						X			
Ethno-Relig $\times$ Age FE							X	X	
Ethno-Relig $\times$ Year FE							X	X	
SES FE								X	
Employed FE								X	

Notes: This table reports OLS estimates of Equation 1 for all unmarried women aged 16–23 who conceived for the first time between 2009–2016. All specifications include age and year fixed effects, and columns sequentially add additional sets of controls, as indicated in the table. Column (9) reports estimates from a triple-differences specification that uses married women aged 16–23 as an unaffected comparison group. Standard errors are clustered at the age-by-year at conception level and reported in parentheses. P-values correspond to tests of equality between each column's estimate and that in Column (1). * $p < 0.10$, ** $p < 0.05$, and *** $p < 0.01$.

Table D.4: Robustness of the Estimated Subsidy Effect to Sequentially Adding Controls (Lenient-Background)

	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)
Treatment Effect	0.057*** (0.013)	0.057*** (0.013)	0.057*** (0.013)	0.057*** (0.010)	0.056*** (0.010)	0.052*** (0.010)	0.066*** (0.011)	0.067*** (0.012)	0.083*** (0.032)
Baseline Mean	0.760	0.760	0.760	0.773	0.773	0.772	0.773	0.770	0.186
P-value		0.969	0.992	0.988	0.971	0.777	0.568	0.582	0.45
N	35,862	35,862	35,862	35,231	35,231	35,176	35,231	33,326	223,833
Month FE		X			X	X	X	X	X
Ethno-Relig FE			X		X	X	X	X	X
Education FE				X	X	X	X	X	X
Ethno-Relig $\times$ Education FE					X	X	X	X	X
Location FE						X			
Ethno-Relig $\times$ Age FE							X	X	
Ethno-Relig $\times$ Year FE							X	X	
SES FE								X	
Employed FE								X	

Notes: This table reports OLS estimates of Equation 1 for lenient-background unmarried women aged 16–23 who conceived for the first time between 2009–2016. All specifications include age and year fixed effects, and columns sequentially add additional sets of controls, as indicated in the table. Column (9) reports estimates from a triple-differences specification that uses married women aged 16–23 as an unaffected comparison group. Standard errors are clustered at the age-by-year at conception level and reported in parentheses. P-values correspond to tests of equality between each column's estimate and that in Column (1). * $p < 0.10$, ** $p < 0.05$, and *** $p < 0.01$.

Table D.5: Robustness of the Estimated Subsidy Effect to Sequentially Adding Controls (Strict-Background)

	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)
Treatment Effect	0.134*** (0.030)	0.134*** (0.030)	0.135*** (0.027)	0.150*** (0.030)	0.148*** (0.027)	0.123*** (0.026)	0.142*** (0.027)	0.142*** (0.027)	0.191*** (0.046)
Baseline Mean	0.487	0.487	0.487	0.509	0.509	0.509	0.509	0.509	0.030
P-value		0.995	0.975	0.69	0.721	0.787	0.832	0.836	0.293
N	5,509	5,509	5,509	5,264	5,264	5,251	5,264	5,262	118,989
Month FE		X			X	X	X	X	X
Ethno-Relig FE			X		X	X	X	X	X
Education FE				X	X	X	X	X	X
Ethno-Relig $\times$ Education FE					X	X	X	X	X
Location FE						X			
Ethno-Relig $\times$ Age FE							X	X	
Ethno-Relig $\times$ Year FE							X	X	
SES FE								X	
Employed FE								X	

Notes: This table reports OLS estimates of Equation 1 for strict-background unmarried women aged 16–23 who conceived for the first time between 2009–2016. All specifications include age and year fixed effects, and columns sequentially add additional sets of controls, as indicated in the table. Column (9) reports estimates from a triple-differences specification that uses married women aged 16–23 as an unaffected comparison group. Standard errors are clustered at the age-by-year at conception level and reported in parentheses. P-values correspond to tests of equality between each column's estimate and that in Column (1). * $p < 0.10$, ** $p < 0.05$, and *** $p < 0.01$.

Table D.6: Sensitivity of Estimated Subsidy Effects to Alternative Age Bandwidths

	All Women			Lenient			Strict		
	16-23	18-21	18-22	16-23	18-21	18-22	16-23	18-21	18-22
Treatment Effect	0.068*** (0.011)	0.053*** (0.013)	0.053*** (0.012)	0.056*** (0.010)	0.044*** (0.012)	0.044*** (0.011)	0.148*** (0.027)	0.110*** (0.032)	0.112*** (0.031)
P-value		0.377	0.355		0.437	0.444		0.363	0.387
Baseline Mean	0.738	0.740	0.732	0.773	0.776	0.766	0.509	0.510	0.513
N	40,495	23,182	22,582	35,231	20,067	19,599	5,264	3,115	2,983

Notes: This table reports OLS estimates of Equation 1 using alternative age bandwidths for defining treated and control groups. Columns (1)–(3) report results for all unmarried women who became pregnant for the first time; Columns (4)–(6) and (7)–(9) report results separately for women from lenient and strict backgrounds, respectively. For each group, we estimate the model using three age ranges: ages 16–23 (our primary analytic bandwidth), ages 18–21, and ages 18–22 excluding age 20. P-values compare estimates from the primary bandwidth (16–23) with those obtained using the alternative bandwidths. All regressions include age-at-conception fixed effects, month fixed effects, and the interaction of ethno-religious-group and education fixed effects. Standard errors, clustered by age-by-year at conception, are reported in parentheses. * $p < 0.10$, ** $p < 0.05$, and *** $p < 0.01$.

Table D.7: Effect of the Subsidy on Abortion Utilization (Including Married Women)

	Full Sample	Lenient	Strict
Treatment Effect	0.0057 (0.0050)	0.0139 (0.0174)	0.0074*** (0.0022)
P-value (levels)			0.709
P-value (%)			0.031
Baseline Mean	0.1618	0.3126	0.0318
N	193,959	89,796	104,163

Notes: This table reports OLS estimates of Equation 1, which identifies the effect of the 2014 subsidy reform on abortion ratios for all women aged 16–23 who became pregnant for the first time between 2009–2016, including both married and unmarried women. Columns (2) and (3) report estimates separately for women from communities with lenient and strict attitudes toward abortion. The reported p-values test whether the treatment effects differ across attitude groups, with the first row comparing level effects (percentage points) and the second row comparing percent effects (scaling by each group’s baseline abortion rate). All regressions include month fixed effects and the interaction of ethno-religious-group and education fixed effects. Standard errors, clustered by age-by-year at conception, are reported in parentheses. * $p < 0.10$, ** $p < 0.05$, and *** $p < 0.01$.

Table D.8: Effect of the Subsidy Excluding Former-Soviet-Union Origin

	All Lenient	Lenient (Excluding FSU)
Treatment Effect	0.056*** (0.01) [-0.04, 0.03]	0.04*** (0.01) [-0.04, 0.03]
P-value	0.261	
Baseline Mean	0.773	0.789
N	35,231	30,106

Notes: This table reports OLS estimates of Equation 1 for unmarried women aged 16–23 from lenient-attitude communities who became pregnant for the first time between 2009–2016. Column (1) reports estimates for the full lenient sample; Column (2) reports estimates for the same sample excluding women from families originating in the Former Soviet Union. The reported p-value tests whether the treatment effects differ between Columns (1) and (2). All regressions include month fixed effects and the interaction of ethno-religious-group and education fixed effects. Standard errors, clustered by age-by-year at conception, are reported in parentheses. Bounds that account for potential differential pre-trends, constructed using the method of Rambachan and Roth (2023), appear in brackets. $*p < 0.10$, $**p < 0.05$, and $***p < 0.01$.

E Formal statement of the family model

This appendix restates the family model from Sections 4.1–4.2 in self-contained, formal terms. Appendix E.1 defines the primitives and preferences. Appendices E.2–E.3 formulate Layers 1 and 2 as constrained-optimization problems for the daughter. Appendix E.4 formulates Layer 3 as a four-stage sequential game with a parental transfer and solves it by backward induction.

The three layers are independent models built on the same primitives, defined in Appendix E.1. They are not successive stages of a single model. Instead, they differ in the budget set available to the daughter. Layer 3 additionally introduces a parental-transfer game.

E.1 Primitives

A family consists of two players, the daughter (D) and her parents (P), and a two-period horizon $t \in \{1, 2\}$ in which the parents consume only in period 1. The primitives are:

- **Endowments.** Daughter incomes $Y_1^D, Y_2^D \in \mathbb{R}_{++}$ with $Y_1^D < Y_2^D$, and parental income $Y^P \in \mathbb{R}_{++}$.
- **Policy.** A procedure price $p \in \mathbb{R}_{++}$ and a subsidy $S \in \{0, 1\}$ that scales the out-of-pocket cost to $(1 - S)p$.
- **Family environment.** An in-kind share $\phi \in [0, 1]$ of the parental allowance and an altruism weight $\alpha^P \in \mathbb{R}_{++}$.
- **Attitudes.** A pair $(\xi^D, \xi^P) \in \mathbb{R}^2$, fixed for a given family.
- **Pregnancy.** An indicator $E \in \{0, 1\}$ with $\Pr(E = 1) = \pi \in (0, 1)$, observed privately by the daughter.
- **Flow utility.** A function $u : \mathbb{R}_+ \rightarrow \mathbb{R} \cup \{-\infty\}$ that is twice continuously differentiable on $\mathbb{R}_{++}$, strictly increasing, strictly concave, and satisfies $u'(0^+) = +\infty$.

All primitives, including the attitude vector (ξ^D, ξ^P) , are common knowledge; only the pregnancy indicator E is privately observed by the daughter. Daughter utility is V^D as in equation (3) and parental utility is V^P as in equation (4), where $A \in \{0, 1\}$ is the abortion indicator and $R \in \{0, 1\}$ the disclosure indicator. The altruism component of V^P is over the daughter's *realized* consumption, regardless of whether parents observe the pregnancy; the disclosure indicator R enters V^P only through the direct abortion term $-\xi^P \cdot A \cdot R$. There is no discounting between periods.

Assumption 3 (Atomless attitudes). *Within each community, the attitudes (ξ^D, ξ^P) have an atomless joint density f_h on $\mathbb{R}^2$; we write f when the community is left implicit. Atomlessness ensures that finite-threshold loci have f -measure zero in the population aggregates (Lemma 4).*

Tie-breaking conventions. At attitude-threshold ties—points at which the daughter’s abortion decision or the parents’ coverage decision is exactly indifferent—the daughter chooses $A = 1$ and parents choose $k = 1$. Because the joint attitude density f is atomless, these loci have f -measure zero and the convention does not affect population predictions.⁵¹

Feasibility conventions. We impose strictly positive consumption: any action that would imply $C_t^D = 0$ for some period t or $C^P = 0$ is infeasible. This convention is automatic when $u(0) = -\infty$ and a normalization when $u(0)$ is finite.

Regularity assumption: Interior parental transfer. We assume that parental altruism α^P is sufficiently strong for parents to optimally make a strictly positive transfer to their daughter. Formally, the optimal transfer is interior if and only if $\alpha^P > u'(Y^P)/u'(Y_1^D)$.

E.2 Layer 1: lifetime budget

In Layer 1 the daughter solves

$$V_1^D(\xi^D, S) \equiv \max_{\substack{C_1^D, C_2^D \geq 0 \\ A \in \{0,1\}}} u(C_1^D) + u(C_2^D) - \xi^D A, \quad (10)$$

subject to $C_1^D + C_2^D + (1 - S)pA \leq Y_1^D + Y_2^D$. Let $W(y) \equiv 2u(y/2)$ denote the value of consuming $y > 0$ smoothly across two periods.⁵² Eliminating (C_1^D, C_2^D) in (10) reduces the problem to a binary choice $A \in \{0, 1\}$ with payoff $W(Y_1^D + Y_2^D - (1 - S)pA) - \xi^D A$. Two thresholds govern the decision:

$$\bar{\xi}_1^D \equiv 0, \quad \underline{\xi}_1^D \equiv \begin{cases} W(Y_1^D + Y_2^D - p) - W(Y_1^D + Y_2^D), & \text{if } Y_1^D + Y_2^D > p, \\ -\infty, & \text{otherwise.} \end{cases} \quad (11)$$

The daughter aborts iff $\xi^D \leq \bar{\xi}_1^D = 0$ when $S = 1$ and iff $\xi^D \leq \underline{\xi}_1^D$ when $S = 0$. The second case in the definition of $\underline{\xi}_1^D$ is the feasibility convention: when the lifetime budget cannot accommodate the full-price abortion, no daughter is an always-taker.

⁵¹In contrast with abortion and coverage, disclosure indifference can arise on a positive-measure region of (ξ^D, ξ^P) -space (Lemma 3). We therefore state the disclosure tie-breaking rule separately in Appendix E.4. This rule does not affect any abortion outcomes.

⁵²Under the strict-positive-consumption convention, $W(y) = -\infty$ at $y = 0$.

E.3 Layer 2: liquidity constraint

In Layer 2 the daughter solves

$$V_2^D(\xi^D, S) \equiv \max_{\substack{C_1^D, C_2^D \geq 0 \\ A \in \{0,1\}}} u(C_1^D) + u(C_2^D) - \xi^D A, \quad (12)$$

subject to

$$\begin{aligned} C_1^D + C_2^D + (1 - S)pA &\leq Y_1^D + Y_2^D, \\ C_1^D + (1 - S)pA &\leq Y_1^D. \end{aligned}$$

Since $Y_1^D < Y_2^D$ and there is no discounting, the liquidity constraint binds at the optimum: the daughter consumes $C_1^D = Y_1^D - (1 - S)pA$ and $C_2^D = Y_2^D$. Write $U_2(A, S) \equiv u(Y_1^D - (1 - S)pA) + u(Y_2^D)$ for the resulting consumption payoff. The Layer-2 thresholds are

$$\bar{\xi}_2^D \equiv U_2(1, 1) - U_2(0, 1) = 0, \quad \underline{\xi}_2^D \equiv \begin{cases} U_2(1, 0) - U_2(0, 0), & \text{if } Y_1^D > p, \\ -\infty, & \text{otherwise.} \end{cases} \quad (13)$$

The daughter aborts iff $\xi^D \leq \bar{\xi}_2^D = 0$ when $S = 1$ and iff $\xi^D \leq \underline{\xi}_2^D$ when $S = 0$.

E.4 Layer 3: the family game

Layer 3 augments Layer 2 with the family transfer mechanism of Section 4.1, played as a four-stage sequential game. Fix a subsidy regime $S \in \{0, 1\}$ at the start; the model compares the $S = 0$ and $S = 1$ games and we suppress the dependence of strategies on S where no confusion arises.

Extensive form. Figure 4 displays the extensive form of the game. The game proceeds in four stages:

- (i) **Allowance.** Parents choose an allowance $T \geq 0$ to maximize V^P on the non-pregnancy branch, subject to $C^P + T \leq Y^P$. The optimum is denoted T_{NP}^* .
- (ii) **Pregnancy realization.** Nature reveals E privately to the daughter.
- (iii) **Disclosure.** If $E = 1$, the daughter chooses $R \in \{0, 1\}$ given $(\xi^D, \xi^P, T_{\text{NP}}^*)$; if $E = 0$, $R = 0$ by convention.
- (iv) **Coverage and abortion.** In the TELL branch ($R = 1$), parents observe disclosure and choose $k \in \{0, 1\}$ to maximize V^P , defined in (4). In the HIDE branch, $k = 0$ by convention. The

daughter then chooses $(A, C_1^D, C_2^D) \in \{0, 1\} \times \mathbb{R}_+^2$ to maximize V^D , defined in (3), subject to

$$C_1^D + (1 - S)(1 - k)pA \leq Y_1^D + T_{\text{NP}}^*, \quad (14)$$

$$C_1^D + C_2^D + (1 - S)(1 - k)pA \leq Y_1^D + Y_2^D + T_{\text{NP}}^*, \quad (15)$$

$$C_1^D \geq \phi T_{\text{NP}}^*. \quad (16)$$

Parental consumption is the residual $C^P = Y^P - T_{\text{NP}}^* - (1 - S)pkAR \geq 0$. Constraints (14) and (16) jointly imply the *funding constraint*

$$(1 - S)(1 - k)pA \leq Y_1^D + (1 - \phi)T_{\text{NP}}^*, \quad (17)$$

so that only the cash share $(1 - \phi)T_{\text{NP}}^*$ of the allowance is available to finance an out-of-pocket abortion.

Remark E.1 (Non-pregnancy Allowance Assumption). Stage (i) treats T as chosen on the non-pregnancy branch, rather than optimized against the pregnancy probability π . Because $\Pr(E = 1) = \pi$ is small in the relevant age bracket (4.7%), the optimal allowance that would account for pregnancy is likely close to T_{NP}^* . The equilibrium definition below incorporates this assumption.

Strategies. The non-trivial decision nodes are: parents' stage-(i) transfer; the pregnant daughter's stage-(iii) disclosure; parents' stage-(iv) coverage in the TELL branch; and the daughter's stage-(iv) abortion-and-consumption choice. Formally, the strategies are:

- the parents' stage-(i) allowance $T \in \mathbb{R}_+$;
- the pregnant daughter's stage-(iii) disclosure choice $R(\xi^D, \xi^P) \in \{0, 1\}$;
- the parents' stage-(iv) coverage choice in the TELL branch $k_T(\xi^D, \xi^P) \in \{0, 1\}$;
- the pregnant daughter's stage-(iv) action $(A^{D-P}, C_1^{D-P}, C_2^{D-P})(\xi^D, \xi^P; k) \in \{0, 1\} \times \mathbb{R}_+^2$, given attitudes and the coverage k realized in stage (iv) ($k = 0$ on the HIDE branch);
- the non-pregnant daughter's stage-(iv) consumption $(C_1^{D-\text{NP}}, C_2^{D-\text{NP}}) \in \mathbb{R}_+^2$.

The goal of the family-level analysis is to characterize abortion behavior and the subsidy's effect as functions of family attitudes (ξ^D, ξ^P) , holding the economic environment $(Y_1^D, Y_2^D, Y^P, p, \phi, \alpha^P, u)$ fixed. We therefore write strategies as functions of attitudes and suppress dependence on the environment from the notation. The allowance T_{NP}^* is itself a function of the environment alone—not of attitudes—so its dependence is absorbed into the same suppression. Strategies are also regime-specific: each is a function of $S \in \{0, 1\}$ in addition to its other arguments. We make the regime S explicit when it matters and suppress it otherwise.

Branches. The game has three branches: the *non-pregnant* branch ($E = 0$), the *HIDE* branch ($E = 1, R = 0$), and the *TELL* branch ($E = 1, R = 1$). We characterize the game by backward induction. Because the stage-(i) allowance is chosen as if the daughter were not pregnant (Remark E.1), it depends only on the non-pregnant daughter's consumption value. We therefore start with the non-pregnant daughter's consumption, then work back to the stage-(i) allowance, and finally forward to the pregnant daughter's stage-(iv) action and the parental coverage choice in the *TELL* branch.

Non-pregnant daughter's consumption. The non-pregnant daughter's consumption problem at allowance T is

$$(C_1^{D-NP^*}(T), C_2^{D-NP^*}(T)) \in \arg \max_{C_1, C_2 \geq 0} u(C_1) + u(C_2) \quad (18)$$

subject to $C_1 + C_2 \leq Y_1^D + Y_2^D + T$, $C_1 \leq Y_1^D + T$, and $C_1 \geq \phi T$.

Stage (i): Allowance. Treating the daughter as non-pregnant (Remark E.1), parents choose

$$T_{NP}^* \equiv \arg \max_{0 \leq T \leq Y^P} u(Y^P - T) + \alpha^P [u(C_1^{D-NP^*}(T)) + u(C_2^{D-NP^*}(T))]. \quad (19)$$

Stage (iv): Pregnant daughter. For any $k \in \{0, 1\}$, the pregnant daughter chooses

$$(A^{D-P^*}(k), C_1^{D-P^*}(k), C_2^{D-P^*}(k)) \in \arg \max_{\substack{A \in \{0,1\} \\ C_1, C_2 \geq 0}} u(C_1) + u(C_2) - \xi^D A \quad (20)$$

subject to constraints (14)–(16) evaluated at $T = T_{NP}^*$; the dependence on S and on the daughter's own ξ^D is suppressed in the notation.

Stage (iv): The TELL branch. Parents observe disclosure and choose k to maximize

$$\begin{aligned} k_T^*(\xi^D, \xi^P) \in \arg \max_{k \in \{0,1\}} & u(Y^P - T_{NP}^* - (1 - S)pk A^{D-P^*}(k)) \\ & + \alpha^P [u(C_1^{D-P^*}(k)) + u(C_2^{D-P^*}(k))] \\ & - \xi^P A^{D-P^*}(k). \end{aligned} \quad (21)$$

Equilibrium definition. An *equilibrium* of the Layer-3 game is a strategy profile—an allowance T , a disclosure rule $R(\cdot)$, a coverage rule $k_T(\cdot)$, and the pregnant and non-pregnant daughters' stage-(iv) actions—such that the allowance solves the non-pregnancy problem (19), that is, $T = T_{NP}^*$. At every other decision node, the acting player maximizes their utility given the rest of the profile, with ties broken according to the conventions in Appendix E.1.

Remark E.2 (Interpretation of the Equilibrium Concept). Parents choose the stage-(i) transfer as if the daughter were not pregnant, rather than to maximize expected utility over the pregnancy distribution.

Thus, the equilibrium defined above is not a perfect Bayesian equilibrium. Conditional on this stage-(i) transfer, all remaining decision nodes form subgames of perfect information.

Thresholds. Define the three Layer-3 thresholds as follows:

- *Never-taker*, $\bar{\xi}_3^D$: the largest ξ^D at which the daughter aborts under $S = 1$.
- *Sponsorship*, $\bar{\xi}^P$: the largest ξ^P at which parents prefer $k = 1$ over $k = 0$ in stage (iv) of the TELL branch under $S = 0$, evaluated on the complier-relevant region $\xi^D \in (\underline{\xi}_3^D, 0]$. When parental coverage is infeasible ($Y^P - T_{NP}^* \leq p$), set $\bar{\xi}^P \equiv -\infty$.
- *Self-fund*, $\underline{\xi}_3^D$: the largest ξ^D at which the daughter aborts in the HIDE branch under $S = 0$, given the pre-committed allowance T_{NP}^* . When self-funding is infeasible (see the Remark on the operational form of the feasibility convention below), set $\underline{\xi}_3^D \equiv -\infty$.

These take values $\bar{\xi}_3^D = 0$ (Lemma 1(a)), $\bar{\xi}^P = u(Y^P - T_{NP}^* - p) - u(Y^P - T_{NP}^*)$ when parental coverage is feasible (Lemma 2), and $\underline{\xi}_3^D = v_p(T_{NP}^*) - v(T_{NP}^*)$ when self-funding is feasible (derived in Appendix F, where v and v_p are defined in equations (23)–(24)). The resulting partition of (ξ^D, ξ^P) -space appears in Figure 3, Panel (c).

Remark E.3 (Operational form of the feasibility convention). The strict-positive-consumption convention in Appendix E.1 is implemented through feasibility conditions at three margins, with the corresponding thresholds set to $-\infty$ when feasibility fails. In Layer 2, self-funding requires $Y_1^D > p$ (strict); hence $\underline{\xi}_2^D = -\infty$ when $Y_1^D \leq p$, since paying for the procedure would otherwise force $C_1^D = 0$. In Layer 3, hidden self-funding requires

$$Y_1^D + (1 - \phi)T_{NP}^* \geq p,$$

with the boundary $Y_1^D + (1 - \phi)T_{NP}^* = p$ admissible only when $\phi > 0$: at the boundary the post-payment feasible interval for C_1^D collapses to the single point $\{\phi T_{NP}^*\}$, so the in-kind constraint binds, and $\phi T_{NP}^* > 0$ keeps C_1^D strictly positive even when the daughter spends her entire cash on the procedure. Hence $\underline{\xi}_3^D = -\infty$ when $Y_1^D + (1 - \phi)T_{NP}^* < p$, and also when $\phi = 0$ and $Y_1^D + T_{NP}^* = p$. Parental coverage requires

$$Y^P - T_{NP}^* > p$$

(strict, since parental consumption has no in-kind cushion); so $\bar{\xi}^P = -\infty$ when $Y^P - T_{NP}^* \leq p$. These conditions, and the associated $-\infty$ thresholds, are used throughout the proofs.

Decision thresholds across layers. The following lemma collects the ordering of the thresholds defined in equations (11) and (13) and in the Layer-3 thresholds above. It is the formal counterpart of

Figure 3.

Lemma 1 (Decision thresholds). *In each layer $l \in \{1, 2, 3\}$, the daughter's decision is governed by an always-taker threshold $\underline{\xi}_l^D$ and a never-taker threshold $\bar{\xi}_l^D$. These thresholds satisfy:*

(a) Never-taker threshold. *The never-taker threshold is common across layers:*

$$\bar{\xi}_1^D = \bar{\xi}_2^D = \bar{\xi}_3^D = 0.$$

(b) Liquidity. *Introducing the liquidity constraint in Layer 2 weakly lowers the always-taker threshold:*

$$\underline{\xi}_2^D \leq \underline{\xi}_1^D,$$

with strict inequality when $Y_1^D < Y_2^D$ and self-funding is feasible. If first-period income does not exceed the price, $Y_1^D \leq p$, self-funding is infeasible and $\underline{\xi}_2^D = -\infty$.

(c) Family transfers. *Introducing family transfers in Layer 3 weakly raises the always-taker threshold:*

$$\underline{\xi}_3^D \geq \underline{\xi}_2^D.$$

Layer 3 also introduces a parental-coverage threshold, $\bar{\xi}^P$, as shown in Figure 3, Panel (c).

(d) Determinants. *All thresholds depend on the family's economic environment, but not on the realized attitudes (ξ^D, ξ^P) .*

Proof. See Appendix F.5.

Threshold-structure lemmas. The three thresholds and the rectangular partition rest on two properties recorded in the lemmas below: parents' coverage decision in the complier-relevant region is determined by the scalar threshold $\bar{\xi}^P$ (Lemma 2), and the daughter discloses iff $\xi^P \leq 0$, a boundary distinct from $\bar{\xi}^P$ (Lemma 3).

Lemma 2 (Parental-coverage single-crossing). *Restrict attention to the complier-relevant region $\xi^D \in (\underline{\xi}_3^D, 0]$, where the daughter does not self-fund in the hide branch under $S = 0$. Assume parental coverage is feasible, $Y^P - T_{NP}^* > p$. Under $S = 0$ and given disclosure, the parental value difference between $k = 1$ and $k = 0$ in stage (iv) is*

$$D_P(\xi^P) \equiv [u(Y^P - T_{NP}^* - p) - u(Y^P - T_{NP}^*)] - \xi^P,$$

strictly decreasing and affine in ξ^P , and independent of ξ^D . Hence parents prefer $k = 1$ iff $\xi^P \leq \bar{\xi}^P$, where

$$\bar{\xi}^P \equiv u(Y^P - T_{NP}^* - p) - u(Y^P - T_{NP}^*) < 0$$

depends only on environmental primitives. If parental coverage is infeasible ($Y^P - T_{NP}^* \leq p$), set $\bar{\xi}^P = -\infty$ and the willing-parent region is empty.

Proof. See Appendix F.2.

Disclosure tie-breaking. We state the disclosure tie-breaking rule (cf. the tie-breaking conventions in Appendix E.1): at any disclosure indifference, the daughter chooses $R = 1$ if $\xi^P \leq 0$ and $R = 0$ if $\xi^P > 0$ —she tells parents who weakly approve of the abortion and hides from those who oppose it. This selection is used only to resolve payoff-equivalent disclosure choices and has no effect on abortion outcomes (Lemma 3 below).⁵³

Remark E.4 (Microfoundation for the tie-break). The disclosure tie-break is not arbitrary. Appendix G shows that, once involving one's parents carries a cost proportional to their abortion distaste, telling approving parents and hiding from opposed ones is the strict best response. The theoretical results in Section 4.3 are unchanged.

Lemma 3 (Disclosure characterization). *Under $S = 0$, on the complier-relevant region $\xi^D \in (\xi_3^D, 0]$, the pregnant daughter obtains a covered abortion iff parents would willingly cover ($\xi^P \leq \bar{\xi}^P$) and no abortion otherwise; her disclosure follows $R^* = 1$ iff $\xi^P \leq 0$ —strictly preferred where it secures coverage ($\xi^P \leq \bar{\xi}^P$) and selected by the tie-breaking rule where it does not ($\xi^P > \bar{\xi}^P$). The disclosure boundary ($\xi^P = 0$) and the coverage boundary ($\xi^P = \bar{\xi}^P$) are distinct, and no disclosure threshold in ξ^D arises. On the never-taker region $\xi^D > 0$, the daughter's disclosure choice does not affect the abortion outcome ($A^* = 0$ in every branch). On the self-fund region $\xi^D \leq \xi_3^D$, the parental coverage choice in stage (iv) is governed by the primitive condition*

$$\alpha^P[v(T_{NP}^*) - v_p(T_{NP}^*)] \geq u(Y^P - T_{NP}^*) - u(Y^P - T_{NP}^* - p), \quad (22)$$

which depends only on the environmental primitives (through α^P , Y^P , p , u , T_{NP}^* , $v(T_{NP}^*)$ and $v_p(T_{NP}^*)$) and is independent of ξ^P : when (22) holds, parents cover and the daughter strictly prefers to disclose, so disclosure is uniquely pinned by strict incentives; when it fails, hide and tell are payoff-equivalent for the daughter and the tie-breaking rule selects $R^* = 1$ iff $\xi^P \leq 0$. Either way the abortion outcome is $A^* = 1$.

⁵³The tie-breaking is needed because, in the parsimonious model, disclosure has no direct payoff cost. When parents would refuse coverage, hide and tell yield the same daughter payoff ($V^{D,T} = V^{D,H} = v(T_{NP}^*)$) on the complier-relevant region; see the proof of Lemma 3), so the selection is not pinned down by strict incentives. Adding a disclosure cost proportional to parental distaste, $g\xi^P$ with any $g > 0$, makes this selection the strict best response—telling approving parents is a benefit, telling opposed ones a cost—without changing abortion outcomes or any of the population-level results (Appendix G).

Under $S = 1$ the abortion outcome is independent of disclosure (the procedure is free and consumption does not depend on A); the maintained tie-breaking rule selects $R^* = 1$ iff $\xi^P \leq 0$, but this selection has no effect on abortion outcomes or population aggregates.

Proof. See Appendix F.3.

Existence and uniqueness.

Lemma 4 (Existence and uniqueness of equilibrium outcomes). (a) *An equilibrium of the Layer-3 game exists.* (b) *The daughter's abortion choice A^* is uniquely pinned down outside the null set $\mathcal{N}_0$ of finite threshold loci (defined in the proof). On the willing-parent complier-relevant region $\{\xi^D \in (\underline{\xi}_3^D, 0], \xi^P \leq \bar{\xi}^P\} \setminus \mathcal{N}_0$, the disclosure choice, parental coverage choice, and consumption profile are uniquely pinned down by strict incentives. On the unwilling-parent complier-relevant region $\{\xi^D \in (\underline{\xi}_3^D, 0], \xi^P > \bar{\xi}^P\}$, hiding and disclosing are payoff-equivalent (Lemma 3), and $R^* = \mathbf{1}[\xi^P \leq 0]$ is uniquely selected by the disclosure tie-breaking convention; the abortion outcome ($A^* = 0$ under $S = 0$, $A^* = 1$ under $S = 1$) is invariant to this selection. On the never-taker region $\xi^D > 0$, disclosure and coverage are payoff-irrelevant ($A^* = 0$ in every branch). On the self-fund region $\xi^D \leq \underline{\xi}_3^D$, parental coverage is uniquely pinned by the primitive inequality of Lemma 3; disclosure is then uniquely pinned by strict incentives when parents would cover and payoff-equivalent otherwise. The abortion outcome $A^* = 1$ on this region is invariant to any selection. Since f is atomless, $\mathcal{N}_0$ has zero f -measure, and all population aggregates—the complier mass and the abortion ratios $A(S | h)$ defined in Section 5—are unique.*

Proof. See Appendix F.1 for existence and Appendix F.4 for uniqueness of equilibrium outcomes.

F Proofs for the family model

This appendix collects the proofs of the structural results that underlie the Layer-3 game in Appendix E.4 and the summary Proposition 1 of Section 4.3 and the threshold Lemma 1 of Appendix E. Throughout, all primitives are those of Appendix E.1; in particular, u is strictly increasing, strictly concave, twice continuously differentiable on $\mathbb{R}_{++}$, with $u'(0^+) = +\infty$, the attitude density f is atomless on $\mathbb{R}^2$ (Assumption 3), and ties are broken by the conventions in Appendix E.1. The proofs are organized in dependency order. We first establish existence of an equilibrium selection (Lemma 4, part a). We then prove the two structural properties of the Layer-3 game: parental-coverage single-crossing (Lemma 2, a one-line closed-form argument under pre-commitment) and disclosure characterization (Lemma 3, an immediate consequence of limited altruism and the constancy of T_{NP}^* across branches). We then return to uniqueness of equilibrium outcomes (Lemma 4, part b), which uses the structural results. Finally, we prove Lemma 1 and then Proposition 1.

Threshold notation. The proofs below use four threshold objects, with consistent typography across layers. We collect them here for reference.

Symbol	Meaning	Location
$\bar{\xi}_l^D$	Never-taker threshold in Layer $l \in \{1, 2, 3\}$ (equals 0)	(11), (13)
$\underline{\xi}_l^D$	Always-taker / self-funding threshold in Layer $l \in \{1, 2\}$	(11), (13)
$\underline{\xi}_3^D$	Layer-3 self-fund threshold: largest ξ^D at which the daughter aborts in hide under $S = 0$; $-\infty$ when infeasible	Appendix E.4
$\bar{\xi}^P$	Parental-coverage threshold on the complier-relevant region	Lemma 2

Overbar denotes a never-taker / coverage-style threshold (above which the relevant action is *not* taken); underbar denotes an always-taker / self-funding threshold (below which the action is strictly preferred). The Layer-3 parental-coverage threshold $\bar{\xi}^P$ is region-specific: it characterizes parental willingness to cover when the daughter would not self-fund under $k = 0$ (the complier-relevant region $\xi^D \in (\underline{\xi}_3^D, 0]$), and does *not* apply in the self-funder region $\xi^D \leq \underline{\xi}_3^D$ where the daughter aborts regardless of coverage. Under pre-commitment, no separate disclosure threshold in ξ^D arises (Lemma 3); the disclosure boundary is $\xi^P = 0$, distinct from the parental-coverage threshold $\bar{\xi}^P$.

The complier-relevant region. Throughout the proofs of the Layer-3 lemmas we refer to the *complier-relevant region* as the rectangle $\xi^D \in (\underline{\xi}_3^D, 0]$: the daughters who would abort if the procedure were free ($\xi^D \leq 0$) but cannot fund a hidden abortion at the full price ($\xi^D > \underline{\xi}_3^D$). The Layer-3 lemmas

characterize equilibrium objects on this region; outside it, abortion outcomes are determined by the basic threshold structure and not by parental coverage or disclosure dynamics.

Notation and shorthand. Given a parental transfer (T, k) , write $\mathcal{B}^D(T, k; S)$ for the set of $(A, C_1^D, C_2^D) \in \{0, 1\} \times \mathbb{R}_+^2$ satisfying the daughter's three Layer-3 constraints (14)–(16), and write

$$\mathcal{F}(T, k; S) \equiv \{A \in \{0, 1\} : (1 - S)(1 - k)pA \leq Y_1^D + (1 - \phi)T\}$$

for the set of *funding-feasible* abortion choices given (T, k, S) (the joint implication of liquidity and in-kind, which is the funding constraint (17)). We use $U_2(A, S) \equiv u(Y_1^D - (1 - S)pA) + u(Y_2^D)$ for the Layer-2 consumption payoff (equation (13)).

For Layer 3 we define two consumption-only value functions of the parental allowance $T \geq 0$ that recur throughout. Let

$$v(T) \equiv \max_{C_1, C_2 \geq 0} u(C_1) + u(C_2) \quad \text{s.t.} \quad C_1 \leq Y_1^D + T, \quad C_1 \geq \phi T, \quad C_1 + C_2 \leq Y_1^D + Y_2^D + T, \quad (23)$$

$$v_p(T) \equiv \max_{C_1, C_2 \geq 0} u(C_1) + u(C_2) \quad \text{s.t.} \quad C_1 \leq Y_1^D + T - p, \quad C_1 \geq \phi T, \quad C_1 + C_2 \leq Y_1^D + Y_2^D + T - p, \quad (24)$$

where v_p is defined only when self-funding is feasible at T in the sense of the operational feasibility convention (Appendix E.4): $Y_1^D + (1 - \phi)T \geq p$, with strict inequality required when $\phi = 0$ (no in-kind cushion). At the boundary $Y_1^D + (1 - \phi)T = p$ with $\phi > 0$, the post-payment feasible interval for C_1^D collapses to $\{\phi T\}$: the in-kind constraint binds and the daughter consumes exactly her in-kind allowance in period 1. Here $v(T)$ is the daughter's *no-cost* consumption value at allowance T : it applies whenever no out-of-pocket abortion expenditure enters her budget — both in the non-pregnant hide baseline (no abortion) and in the tell-cover branch with $A = 1$ (parents pay p). $v_p(T)$ is her *hide-self-fund* value at allowance T : she receives the same allowance but pays p out of cash, so the liquidity and lifetime constraints are tightened by p , while the in-kind floor stays at ϕT .

Both v and v_p are continuous in T on their domains by the maximum theorem (the feasible set in (23) and (24) is a continuous, compact-valued correspondence in T , and the objective is continuous). Both are strictly increasing in T . This follows from the regime formulas below: in the liquidity-bound, smoothing, and in-kind-bound regimes, the relevant consumption argument(s) entering the value increase with T (with regime-specific slopes 1, 1/2, and the pair $(\phi, 1 - \phi)$ respectively), and strict monotonicity of u gives a strict increase in value. When $\phi < 1$, the fungible component $(1 - \phi)T$ rises strictly; when $\phi = 1$, the in-kind floor and total resources rise one-for-one with T .

Strict concavity of $v(T)$ in T is the load-bearing property used to pin down a unique parental

allowance in Lemma 4, Step 5. It does not follow automatically from strict concavity of u plus a convex parametrization (a value function can be affine over an interval if the optimizer travels along a linear path). We verify it directly by case analysis. The daughter's optimum at allowance T falls into one of three regimes, defined by which constraint binds:

- *Smoothing* (neither liquidity nor in-kind binds): $C_1^D = C_2^D = (Y_1^D + Y_2^D + T)/2$, so $v(T) = 2u((Y_1^D + Y_2^D + T)/2)$.
- *Liquidity-bound* (liquidity binds, in-kind slack): $C_1^D = Y_1^D + T$, $C_2^D = Y_2^D$, so $v(T) = u(Y_1^D + T) + u(Y_2^D)$.
- *In-kind-bound* (in-kind binds, liquidity slack): $C_1^D = \phi T$, $C_2^D = Y_1^D + Y_2^D + (1 - \phi)T$, so $v(T) = u(\phi T) + u(Y_1^D + Y_2^D + (1 - \phi)T)$.

Equivalently, since flow utility is strictly increasing the lifetime budget binds at the optimum, and the optimal allocation has the closed form

$$C_1^*(T) = \min \left\{ \max \left\{ \frac{Y_1^D + Y_2^D + T}{2}, \phi T \right\}, Y_1^D + T \right\}, \quad C_2^*(T) = Y_1^D + Y_2^D + T - C_1^*(T),$$

which traces out the three regimes above as T varies. Each expression is strictly concave in T as a sum/composition of strictly concave u with linear arguments (the only borderline case is in-kind-bound with $\phi = 0$, where the first term reduces to the constant $u(0)$ — but $\phi = 0$ rules out the in-kind regime entirely, since $\phi T = 0$ never binds when liquidity allows $C_1 > 0$). The right derivative $v'(T)$ takes the corresponding regime expression (liquidity-bound: $u'(Y_1^D + T)$; smoothing: $u'((Y_1^D + Y_2^D + T)/2)$; in-kind-bound: $\phi u'(\phi T) + (1 - \phi)u'(Y_1^D + Y_2^D + (1 - \phi)T)$) and is continuous across regime boundaries by the matching of one-sided derivatives at the projection kinks. The possible regimes are ordered along T as liquidity-bound (small T , when Y_2^D is large) $\rightarrow$ smoothing $\rightarrow$ in-kind-bound (large T , when $\phi > 1/2$; the in-kind regime is absent when $\phi \leq 1/2$), and within each regime v' is strictly decreasing in T (strict concavity of u applied to the relevant linear argument); continuity across boundaries preserves this strict monotonicity. Hence v' is continuous and strictly decreasing on the domain, so v is strictly concave.

The same argument applies to v_p on its (smaller) domain. The liquidity cap becomes $C_1 \leq Y_1^D + T - p$ and the lifetime cap becomes $C_1 + C_2 \leq Y_1^D + Y_2^D + T - p$, while the in-kind floor stays at ϕT (not $\phi(T - p)$). The projection formula $C_1^*(T) = \min \{ \max \{ (Y_1^D + Y_2^D + T - p)/2, \phi T \}, Y_1^D + T - p \}$ has the same regime structure, and the strict-monotonicity argument for v'_p goes through identically on the self-funding domain where v_p is defined.

For Layer 3 we also use $\tilde{U}(A; T, k, S)$ for the daughter's stage-(iv) maximized consumption value:

$$\tilde{U}(A; T, k, S) \equiv \max_{C_1, C_2 \geq 0} u(C_1) + u(C_2) \quad \text{s.t. } (A, C_1, C_2) \in \mathcal{B}^D(T, k; S), \quad (25)$$

with the convention $\tilde{U}(A; T, k, S) = -\infty$ when $A \notin \mathcal{F}(T, k; S)$ or the constraint set in (25) is empty. Note $\tilde{U}(0; T, k, S) = v(T)$ (no abortion) and $\tilde{U}(1; T, k = 1, S = 0) = v(T)$ (parents cover); the hide-self-fund case gives $\tilde{U}(1; T, k = 0, S = 0) = v_p(T)$ when funding-feasible.

F.1 Existence and selection (proof of Lemma 4, part a)

We construct an equilibrium by setting the stage-(i) parental allowance to T_{NP}^* per the pre-commitment assumption and solving stages (ii)–(iv) by backward induction. The selection is unique on all dimensions outside a measure-zero null set of indifference loci; uniqueness is established in Section F.4 after the structural lemmas.

Technical preliminary: handling $u(0) = -\infty$. Appendix E.1 allows $u : \mathbb{R}_+ \rightarrow \mathbb{R} \cup \{-\infty\}$, which can take $u(0) = -\infty$, and the strict-positive-consumption convention declares any action with $C_t^D = 0$ or $C^P = 0$ infeasible. The naive feasible set (after removing zero-consumption boundary allocations) is then not closed, so we cannot invoke Weierstrass on it directly. We argue via maximizing sequences.

We first record a boundary argument for the continuous consumption subproblems used below: the non-pregnant daughter's consumption problem and the pregnant daughter's stage-(iv) problem conditional on a feasible abortion action A and a fixed coverage choice k . The parental stage-(iv) coverage choice is binary and is handled directly in Step 3; the stage-(i) parental transfer problem is handled separately in Step 5. Each continuous subproblem has the form $\max u(C_1) + u(C_2)$ subject to a closed bounded constraint set $\Omega \subseteq \mathbb{R}_+^2$. By $Y_1^D, Y_2^D > 0$ and the interior-transfer regularity ($T_{\text{NP}}^* > 0$), each such subproblem admits at least one strictly feasible allocation $x^0 \in \Omega \cap \mathbb{R}_{++}^2$ with $u(x^0) > -\infty$ and finite objective value V^0 .

Let $\{x_n\} \subset \Omega$ be a maximizing sequence (objective value $\rightarrow \sup$). We show every maximizing sequence is eventually bounded away from each coordinate-zero face, so a subsequence converges in a compact subset of $\mathbb{R}_{++}^2$.

Suppose, for contradiction, that some consumption coordinate i satisfies $(x_n)_i \rightarrow 0$ along a subsequence. The case $u(0) = -\infty$ is immediate: the objective along the subsequence tends to $-\infty$, contradicting the maximizing property (the benchmark x^0 has finite value). Assume henceforth that $u(0)$ is finite, and split by the vanishing coordinate.

Case $i = C_1^D$. The in-kind floor $C_1^D \geq \phi T$ together with $C_1^D \rightarrow 0$ forces $\phi T = 0$ (otherwise the floor would already bound C_1^D away from zero), so the floor is slack. We construct a feasible perturbation that shifts mass from C_2^D to C_1^D : this preserves the lifetime cap (sum unchanged) and tightens the liquidity cap $C_1^D \leq Y_1^D + T$ from below, which is feasible since C_1^D was near zero. The donor coordinate C_2^D must stay bounded away from zero along the subsequence: if both coordinates converged to zero, the objective would converge to $2u(0)$, strictly below the value at the strictly positive benchmark x^0 ,

contradicting the maximizing property. Let $\delta > 0$ be a lower bound on $(x_n)_{C_2^D}$ along the subsequence and set $B \equiv u'(\delta/2) < \infty$. By the Inada condition $u'(0^+) = +\infty$, choose $\varepsilon \in (0, \delta/2)$ small enough that $u'(2\varepsilon) > B$. By concavity of u ,

$$\begin{aligned} u((x_n)_{C_1^D} + \varepsilon) - u((x_n)_{C_1^D}) &\geq u'((x_n)_{C_1^D} + \varepsilon) \varepsilon, \\ u((x_n)_{C_2^D}) - u((x_n)_{C_2^D} - \varepsilon) &\leq u'((x_n)_{C_2^D} - \varepsilon) \varepsilon \leq B \varepsilon. \end{aligned}$$

For n large enough that $(x_n)_{C_1^D} \leq \varepsilon$, the gain bound satisfies $u'((x_n)_{C_1^D} + \varepsilon) \geq u'(2\varepsilon) > B$, so the perturbed allocation's objective exceeds the original by at least $(u'(2\varepsilon) - B) \varepsilon > 0$, uniformly in n large. This contradicts $\{x_n\}$ being a maximizing sequence.

Case $i = C_2^D$. No trade-off is needed: the in-kind floor $C_1^D \geq \phi T$ does not constrain C_2^D , and the liquidity constraint $C_1^D + q \leq Y_1^D + T$ (with $q \geq 0$ the daughter's out-of-pocket payment) leaves the lifetime constraint $C_1^D + C_2^D + q \leq Y_1^D + Y_2^D + T$ slack by at least $Y_2^D - C_2^D$. Along the subsequence with $C_2^D \rightarrow 0$ this slack is bounded below by $Y_2^D/2 > 0$ for n large, so C_2^D can be increased by any $\varepsilon \in (0, Y_2^D/2)$ without changing C_1^D , preserving every constraint. Since $(x_n)_{C_2^D} \rightarrow 0$ along the subsequence and u is continuous, $u((x_n)_{C_2^D} + \varepsilon) - u((x_n)_{C_2^D}) \rightarrow u(\varepsilon) - u(0) > 0$, so the gain is bounded below by $\frac{1}{2} [u(\varepsilon) - u(0)] > 0$ for n large, again contradicting maximization.

Hence every maximizing sequence has both coordinates eventually bounded away from zero by some $\delta > 0$. Since Ω is also bounded above, a subsequence converges in $\Omega \cap [\delta, M]^2$, a compact subset of $\mathbb{R}_{++}^2$. The objective is continuous on this compact set, so the limit point achieves the supremum. Strict concavity of the objective gives uniqueness of the daughter's consumption maximizer in both continuous subproblems.

Step 1: the non-pregnant daughter's consumption problem. For each $T \in [0, Y^P]$, the daughter's problem (18) is to maximize $u(C_1) + u(C_2)$ over

$$\mathcal{B}^{D, NP}(T) = \{(C_1, C_2) \in \mathbb{R}_+^2 : C_1 + C_2 \leq Y_1^D + Y_2^D + T, C_1 \leq Y_1^D + T, C_1 \geq \phi T\}.$$

This set is non-empty, closed, and bounded: the closed-form projection $C_1^*(T) = \min\{\max\{(Y_1^D + Y_2^D + T)/2, \phi T\}, Y_1^D + T\}$ paired with $C_2^*(T) = Y_1^D + Y_2^D + T - C_1^*(T)$, derived in the notation paragraph above, satisfies the in-kind floor, the liquidity cap, and the lifetime cap by construction for any $T \geq 0$. By the technical preliminary, the relevant feasible set is a compact subset of $\mathbb{R}_{++}^2$; the objective is continuous-real and strictly concave; Weierstrass and strict concavity yield a unique maximizer $(C_1^{D-NP*}(T), C_2^{D-NP*}(T)) \in \mathbb{R}_{++}^2$. The maximum value coincides with $v(T)$ defined in (23), and is jointly continuous in T by the maximum theorem.

Step 2: the pregnant daughter's stage-(iv) problem. Fix (ξ^D, k, S) with $T = T_{\text{NP}}^*$. For each $A \in \{0, 1\} \cap \mathcal{F}(T_{\text{NP}}^*, k; S)$, the consumption sub-problem (20) conditional on A is again continuous strictly concave maximization on a compact subset of $\mathbb{R}_{++}^2$. It yields a unique conditional consumption maximizer with value $\tilde{U}(A; T_{\text{NP}}^*, k, S)$, where (as noted) $\tilde{U}(0; T_{\text{NP}}^*, k, S) = v(T_{\text{NP}}^*)$ and the $A = 1$ case gives $v(T_{\text{NP}}^*)$ under $k = 1$ (parents cover) and $v_p(T_{\text{NP}}^*)$ under $k = 0$ (daughter pays out of cash, when funding-feasible). If $1 \notin \mathcal{F}(T_{\text{NP}}^*, k; S)$, set $\tilde{U}(1; T_{\text{NP}}^*, k, S) = -\infty$.

The daughter then picks $A^* \in \arg \max_{A \in \{0, 1\}} [\tilde{U}(A; T_{\text{NP}}^*, k, S) - \xi^D A]$ — a binary maximization with value $\max\{\tilde{U}(0; T_{\text{NP}}^*, k, S), \tilde{U}(1; T_{\text{NP}}^*, k, S) - \xi^D\}$. Uniqueness of A^* fails only on the indifference locus $\xi^D = \tilde{U}(1; T_{\text{NP}}^*, k, S) - \tilde{U}(0; T_{\text{NP}}^*, k, S)$ (a singleton in ξ^D for each (k, S)); the tie-breaking convention $A = 1$ over $A = 0$ gives a unique A^* .

Step 3: parents' stage-(iv) coverage choice. In the TELL branch, parents choose $k \in \{0, 1\}$ from equation (21). Because T is pre-committed to T_{NP}^* in stage (i), the only optimization in stage (iv) is over the binary k . A maximizer exists because k is binary: $k = 0$ is always feasible and assigned its real-valued payoff in (21), and $k = 1$ is either feasible (when $Y^P - T_{\text{NP}}^* > p$) and assigned its payoff in (21), or infeasible under the strict-positive-consumption convention and assigned value $-\infty$. This step establishes existence only. The coverage comparison has the scalar form in Lemma 2 on the complier-relevant region; off that region, coverage choices may be payoff-irrelevant for the abortion outcome and are handled in Section F.4. The selection is uniquely determined outside the indifference locus in ξ^P (a singleton for each ξ^D on the complier-relevant region; see Lemma 2). In the HIDE branch, $k = 0$ by convention.

Step 4: the daughter's stage-(iii) disclosure choice. Given the parental coverage selection from Step 3 and the daughter's own stage-(iv) best response from Step 2, the daughter's continuation values $V^{D,T}(\xi^D, \xi^P)$ and $V^{D,H}(\xi^D)$ are well-defined real numbers. She picks $R^* \in \arg \max_R V^{D,R}$. By the tie-breaking convention ($R = 1$ at indifference if $\xi^P \leq 0$, $R = 0$ if $\xi^P > 0$), R^* is uniquely selected.

Step 5: parents' stage-(i) pre-commitment transfer. The behavioral assumption sets the allowance to T_{NP}^* , where T_{NP}^* solves the non-pregnancy problem $\max G_H(T)$ with $G_H(T) \equiv u(Y^P - T) + \alpha^P v(T)$, as in (19). Under the strict-positive-consumption convention the feasible set is $T \in [0, Y^P)$ ($T = Y^P$ would force $C^P = 0$), which is not compact, so we cannot apply Weierstrass on the closed interval $[0, Y^P]$ directly. We use the same maximizing-sequence argument as the technical preliminary.

The map $T \mapsto v(T)$ is continuous and strictly concave by the regime case-analysis in the notation paragraph above; hence G_H is strictly concave on $[0, Y^P)$. Let $\{T_n\} \subset [0, Y^P)$ be a maximizing sequence. Suppose, for contradiction, $T_n \rightarrow Y^P$, so $C_n^P = Y^P - T_n \downarrow 0$. Each of the three regime

formulas for $v'(T)$ is bounded above on bounded T , so the supergradient $s_n \in \partial v(T_n)$ is bounded by some $M < \infty$ on a neighborhood of Y^P . Since $u'(0^+) = +\infty$ by Inada, for each n large choose $\varepsilon_n > 0$ small enough that $T_n - \varepsilon_n \geq 0$ and $u'(C_n^P + \varepsilon_n) > \alpha^P M$. By the mean-value lower bound for concave u ,

$$u(C_n^P + \varepsilon_n) - u(C_n^P) \geq u'(C_n^P + \varepsilon_n) \varepsilon_n > \alpha^P M \varepsilon_n,$$

while the daughter term loses at most $\alpha^P M \varepsilon_n$. The parental gain therefore strictly exceeds the daughter loss for n large, contradicting maximization. Hence every maximizing sequence is bounded away from Y^P , lying eventually in $[0, Y^P - \varepsilon]$ for some $\varepsilon > 0$. On this compact subset G_H is continuous, so Weierstrass yields existence; strict concavity yields uniqueness. The lower corner $T = 0$ is ruled out by the interior-transfer regularity assumption in Appendix E.1, which guarantees a strictly positive baseline transfer. Equivalently, one may state the assumption as the one-sided condition $\alpha^P v'_+(0) > u'(Y^P)$, where $v'_+(0)$ denotes the right derivative of v at zero, so that G_H has a strictly positive right derivative at $T = 0$. Since at $T = 0$ the daughter's problem is liquidity-bound (the smoothing target $(Y_1^D + Y_2^D)/2 > Y_1^D$ exceeds the liquidity cap Y_1^D , given $Y_1^D < Y_2^D$), $v'_+(0) = u'(Y_1^D)$ and the condition reduces to the explicit form $\alpha^P u'(Y_1^D) > u'(Y^P)$. (Note that Inada on u does not apply at $T = 0$, because parental consumption $Y^P > 0$ and the daughter's own incomes $Y_1^D, Y_2^D > 0$ keep all arguments of u strictly positive there.) Under this maintained assumption, the unique maximizer T_{NP}^* is interior ($0 < T_{NP}^* < Y^P$) and satisfies the KKT condition $u'(Y^P - T_{NP}^*) = \alpha^P s_H$ for some supergradient $s_H \in \partial v(T_{NP}^*)$.

This completes the construction of an equilibrium. $\square$

F.2 Parental-coverage single-crossing (proof of Lemma 2)

We work under $S = 0$, on the complier-relevant region $\xi^D \in (\xi_3^D, 0]$, and assume parental coverage is feasible ($Y^P - T_{NP}^* > p$); the infeasible case is handled by setting $\bar{\xi}^P = -\infty$. By definition of ξ_3^D , the inequality $v_p(T_{NP}^*) - \xi^D < v(T_{NP}^*)$ holds on this region, so the daughter does not self-fund in the hide branch under $k = 0$ —either because funding is literally infeasible or because paying p out of pocket is too costly.

Parental value at $k = 1$. With $A^{D-P*}(1) = 1$ (the procedure is free to the daughter; she aborts since $\xi^D \leq 0$):

$$V^{P,k=1} = u(Y^P - T_{NP}^* - p) + \alpha^P v(T_{NP}^*) - \xi^P.$$

Parental value at $k = 0$. With $A^{D-P*}(0) = 0$ (no abortion on the complier-relevant region):

$$V^{P,k=0} = u(Y^P - T_{NP}^*) + \alpha^P v(T_{NP}^*).$$

The $-\xi^P$ term enters V^P only when $A = 1$ and $R = 1$ (the abortion is both performed and disclosed);

under $k = 0$ on the complier-relevant region there is no abortion, so it drops out.

Difference.

$$D_P(\xi^P) \equiv V^{P,k=1} - V^{P,k=0} = u(Y^P - T_{NP}^* - p) - u(Y^P - T_{NP}^*) - \xi^P.$$

The daughter's consumption term $\alpha^P v(T_{NP}^*)$ cancels because the pre-committed allowance is the same under both k , so the daughter's consumption value is $v(T_{NP}^*)$ in both branches (no abortion under $k = 0$; under $k = 1$ parents pay and her consumption is unaffected). D_P is therefore strictly decreasing and affine in ξ^P , independent of ξ^D on the complier-relevant region.

Threshold. $D_P(\xi^P) \geq 0$ iff $\xi^P \leq \bar{\xi}^P$, where

$$\bar{\xi}^P \equiv u(Y^P - T_{NP}^* - p) - u(Y^P - T_{NP}^*).$$

By strict monotonicity of u and $p > 0$, $u(Y^P - T_{NP}^* - p) < u(Y^P - T_{NP}^*)$, so $\bar{\xi}^P < 0$. Parents prefer $k = 1$ iff $\xi^P \leq \bar{\xi}^P$; at $\xi^P = \bar{\xi}^P$ the tie-breaking convention assigns $k = 1$. The threshold depends only on environmental primitives via T_{NP}^* (itself a function of $(Y_1^D, Y_2^D, Y^P, \phi, \alpha^P, u)$). $\square$

F.3 Disclosure characterization (proof of Lemma 3)

We work under $S = 0$. Fix $\xi^D \in (\underline{\xi}_3^D, 0]$ (complier-relevant region). The daughter compares disclosing vs hiding, knowing the subsequent parental and own best responses.

If $\xi^P \leq \bar{\xi}^P$ (parents willing). By Lemma 2, upon disclosure parents play $k = 1$ and the daughter aborts. Her continuation value is $V^{D,T} = v(T_{NP}^*) - \xi^D$. On hide, she does not self-fund (since $\xi^D > \underline{\xi}_3^D = v_p(T_{NP}^*) - v(T_{NP}^*)$ implies $v_p(T_{NP}^*) - \xi^D < v(T_{NP}^*)$ —either self-funding is literally infeasible or paying p out of pocket is too costly relative to her abortion benefit), so she does not abort and her continuation value is $V^{D,H} = v(T_{NP}^*)$. The difference is $V^{D,T} - V^{D,H} = -\xi^D \geq 0$, with equality only at $\xi^D = 0$. By tie-breaking ($R = 1$ at indifference for willing parents), she discloses for all $\xi^D \in (\underline{\xi}_3^D, 0]$.

If $\xi^P > \bar{\xi}^P$ (parents unwilling). By Lemma 2, upon disclosure parents play $k = 0$. On the complier-relevant region the daughter does not self-fund, so she does not abort; her continuation value is $V^{D,T} = v(T_{NP}^*)$. On hide, also $V^{D,H} = v(T_{NP}^*)$. The two are equal; by tie-breaking ($R = 1$ iff $\xi^P \leq 0$), she discloses when $\xi^P \leq 0$ and hides when $\xi^P > 0$. The abortion outcome is $A^* = 0$ either way.

Hence on the complier-relevant region the daughter obtains an abortion iff $\xi^P \leq \bar{\xi}^P$ and discloses iff $\xi^P \leq 0$; the disclosure boundary $\xi^P = 0$ is distinct from the coverage threshold $\bar{\xi}^P$, and no separate disclosure threshold in ξ^D arises.

Off the complier-relevant region. For $\xi^D > 0$ (never-takers): the daughter does not abort in any branch under any subsidy regime, so $V^{D,T} = V^{D,H} = v(T_{NP}^*)$, and her disclosure choice does not affect

the abortion outcome. For $\xi^D \leq \underline{\xi}_3^D$ (self-funders): she can self-fund; $v_p(T_{NP}^*) - \xi^D \geq v(T_{NP}^*)$, so on hide she aborts. On disclosure, parents compare $k = 1$ (utility $u(Y^P - T_{NP}^* - p) + \alpha^P v(T_{NP}^*) - \xi^P$) with $k = 0$ (utility $u(Y^P - T_{NP}^*) + \alpha^P v_p(T_{NP}^*) - \xi^P$; the daughter still aborts via self-funding, so $A = R = 1$ and the $-\xi^P$ term is present in both branches). The $-\xi^P$ term cancels in the comparison, so parents prefer $k = 1$ iff condition (22) of Lemma 3 holds, independent of ξ^P . When (22) holds, parents cover on disclosure; the daughter then strictly prefers tell (consumption $v(T_{NP}^*) > v_p(T_{NP}^*)$) and disclosure is uniquely pinned by strict incentives. When (22) fails, parents do not cover; hide and tell yield the daughter the same payoff $v_p(T_{NP}^*) - \xi^D$, and the tie-breaking rule selects $R^* = 1$ iff $\xi^P \leq 0$. Either way the abortion outcome is $A^* = 1$. $\square$

F.4 Uniqueness of equilibrium outcomes (proof of Lemma 4, part b)

Combining Lemma 4 part a with Lemmas 2 and 3, the equilibrium outcomes on the complier-relevant region $\xi^D \in (\underline{\xi}_3^D, 0] \setminus \mathcal{N}_0$ (where $\mathcal{N}_0$ is the null set defined below) are characterized:

- For $\xi^P \leq \bar{\xi}^P$ (willing parents): $R^* = 1$, $k^* = 1$, $A^* = 1$, uniquely pinned down by strict preference outside the null locus $\{\xi^P = \bar{\xi}^P\}$.
- For $\xi^P > \bar{\xi}^P$ (unwilling parents): hiding and disclosing are payoff-equivalent (Lemma 3); $R^* = 1[\xi^P \leq 0]$ is uniquely selected by the disclosure tie-breaking convention but is not strictly preferred. No coverage arises on path ($k^* = 0$ whenever parents are reached). The abortion outcome $A^* = 0$ under $S = 0$ ($A^* = 1$ under $S = 1$) is invariant to the disclosure selection.

The abortion outcome is therefore strictly pinned down on the entire complier-relevant region (independently of the disclosure tie-breaking rule); disclosure itself is strictly pinned down only where it secures coverage ($\xi^P \leq \bar{\xi}^P$), and pinned down by selection otherwise. Equilibrium consumption is uniquely determined by the daughter's stage-(iv) best response (Lemma 4, Step 2), which is unique by strict concavity.

For $\xi^D > 0$ (never-takers), disclosure and coverage selections are payoff-irrelevant: the daughter does not abort in any branch and $A^* = 0$. For $\xi^D \leq \underline{\xi}_3^D$ (self-funders), Lemma 3 pins down parents' stage-(iv) coverage choice by the primitive inequality (22), independent of ξ^P . When the inequality holds, parents cover on disclosure and the daughter strictly prefers to disclose; both are uniquely pinned by strict incentives. When the inequality fails, parents would not cover and hide and tell are payoff-equivalent for the daughter; the disclosure tie-breaking convention selects one branch, but the abortion outcome $A^* = 1$ is invariant. The parental-coverage threshold $\bar{\xi}^P$ is region-specific (see Appendix E.4) and does not apply off the complier-relevant region.

Outside the null set

$$\mathcal{N}_0 = \{\xi^D = 0\} \cup \{\xi^D = \underline{\xi}_3^D : \xi_3^D \in \mathbb{R}\} \cup \{\xi^P = \bar{\xi}^P : \bar{\xi}^P \in \mathbb{R}\},$$

the abortion outcome A^* is pinned down by strict inequalities. On $\mathcal{N}_0$, A^* is selected by the tie-breaking convention (daughter chooses $A = 1$ at attitude indifference; parents choose $k = 1$ at coverage indifference). Each threshold locus is included only when the threshold is finite (infinite thresholds correspond to empty subsets of $\mathbb{R}^2$). Each finite-threshold locus is a horizontal or vertical line in $\mathbb{R}^2$, with two-dimensional Lebesgue measure zero. Since f is atomless with a density, $\mathcal{N}_0$ has f -measure zero, so population aggregates are unaffected by the tie-breaking convention.

All population aggregates—the complier mass and the abortion ratios $A(S \mid h)$ defined in Section 5—are therefore unique. $\square$

F.5 Decision thresholds (proof of Lemma 1)

We prove the four parts in turn. The notation $\bar{\xi}_l^D$ refers to the never-taker threshold in Layer l and $\underline{\xi}_l^D$ to the always-taker threshold. We use the explicit bar/underbar typography throughout. Throughout we adopt the *strict-positive-consumption convention* (Appendix E.1): any action requiring $C_t^D = 0$ in some consumed period is infeasible.

Proof of part 1(a) (never-taker thresholds). In each layer $l \in \{1, 2, 3\}$, under $S = 1$ the procedure carries no out-of-pocket cost. In Layer 1, the lifetime budget becomes $C_1^D + C_2^D \leq Y_1^D + Y_2^D$, independent of A . The consumption maximum $W(Y_1^D + Y_2^D) = 2u((Y_1^D + Y_2^D)/2)$ is the same whether the daughter aborts or not, so her choice reduces to maximizing $-\xi^D A$: $A^* = 1$ iff $\xi^D \leq 0$ (with tie-breaking at $\xi^D = 0$). Hence $\bar{\xi}_1^D = 0$.

In Layer 2, $S = 1$ leaves the liquidity constraint as $C_1^D \leq Y_1^D$, binding at the optimum, so $C_1^D = Y_1^D$ and $C_2^D = Y_2^D$ regardless of A . Same argument: $\bar{\xi}_2^D = 0$.

In Layer 3 with $S = 1$, the funding constraint (17) is slack (since $(1 - S)(1 - k)p = 0$ for either k), so the daughter's budget set is independent of A . Her consumption maximum equals $v(T_{NP}^*)$ regardless of A (whether tell-cover or hide), so she aborts iff $\xi^D \leq 0$: $\bar{\xi}_3^D = 0$.

Proof of part 1(b) (liquidity). *Layer 1.* Under $S = 0$, the daughter trades off ξ^D against the utility loss of lifetime resources falling from $Y_1^D + Y_2^D$ to $Y_1^D + Y_2^D - p$. With smoothing, the loss is $W(Y_1^D + Y_2^D) - W(Y_1^D + Y_2^D - p)$. She aborts iff $\xi^D \leq W(Y_1^D + Y_2^D - p) - W(Y_1^D + Y_2^D)$. When $Y_1^D + Y_2^D \leq p$ (post-abortion lifetime resources non-positive), the strict-positive-consumption convention makes abortion infeasible and $\underline{\xi}_1^D = -\infty$. This is exactly (11).

Layer 2. Under $S = 0$ with the liquidity constraint binding at $C_1^D = Y_1^D - pA$, consumption value is $U_2(A, 0) = u(Y_1^D - pA) + u(Y_2^D)$. The daughter aborts iff $\xi^D \leq U_2(1, 0) - U_2(0, 0)$, when feasible. Self-funding requires strictly positive period-1 consumption, which requires $Y_1^D > p$; at $Y_1^D \leq p$, the convention makes self-funding infeasible and $\xi_2^D = -\infty$. This matches (13).

Layer 1 vs Layer 2. We compare $\xi_1^D = W(Y_1^D + Y_2^D - p) - W(Y_1^D + Y_2^D)$ and $\xi_2^D = u(Y_1^D - p) - u(Y_1^D)$. Both are negative (utility losses), and we want $\xi_2^D \leq \xi_1^D$ (i.e., the Layer-2 loss is larger in magnitude).

Setting $m \equiv (Y_1^D + Y_2^D)/2$, we have $Y_1^D < m$ (using $Y_1^D < Y_2^D$). The difference of utility losses is, by integration,

$$\begin{aligned} [u(Y_1^D) - u(Y_1^D - p)] - 2[u(m) - u(m - p/2)] &= \int_0^p u'(Y_1^D - t) dt - \int_0^{p/2} 2u'(m - s) ds \\ &= \int_0^p u'(Y_1^D - t) dt - \int_0^p u'(m - t/2) dt \\ &= \int_0^p [u'(Y_1^D - t) - u'(m - t/2)] dt. \end{aligned}$$

For every $t \in [0, p]$, $Y_1^D - t < m - t/2$ iff $Y_1^D - m < t/2$, which holds since $Y_1^D - m < 0 \leq t/2$. By strict monotonicity of u' (a consequence of strict concavity), $u'(Y_1^D - t) > u'(m - t/2)$ for every $t \in [0, p]$. The integrand is strictly positive on $(0, p]$, so the integral is strictly positive: $[u(Y_1^D) - u(Y_1^D - p)] > 2[u(m) - u(m - p/2)]$, i.e., the Layer-2 loss is strictly larger. Therefore $\xi_2^D < \xi_1^D$ strictly when $Y_1^D < Y_2^D$ and $Y_1^D > p$. When $Y_1^D \leq p$, $\xi_2^D = -\infty \leq \xi_1^D$ weakly.

Proof of part 1(c) (family transfers). The Layer-3 hide-branch funding constraint at $S = 0$ is $p \leq Y_1^D + (1 - \phi)T_{\text{NP}}^*$, weakly more permissive than the Layer-2 constraint $p \leq Y_1^D$ (since $(1 - \phi)T_{\text{NP}}^* \geq 0$). When Layer-3 self-funding is infeasible—either $Y_1^D + (1 - \phi)T_{\text{NP}}^* < p$, or $\phi = 0$ and $Y_1^D + T_{\text{NP}}^* = p$ —both cases imply $Y_1^D < p$ (in the first since $(1 - \phi)T_{\text{NP}}^* \geq 0$; in the second since $Y_1^D = p - T_{\text{NP}}^* < p$ given the interior allowance $T_{\text{NP}}^* > 0$). So Layer-2 self-funding is also infeasible: $\xi_3^D = \xi_2^D = -\infty$ and the inequality $\xi_3^D \geq \xi_2^D$ holds with equality. We now treat the case in which Layer-3 self-funding is feasible, so $\xi_3^D = v_p(T_{\text{NP}}^*) - v(T_{\text{NP}}^*)$ is the indifference utility cost in the hide branch.

We prove $\xi_3^D \geq \xi_2^D$ by an envelope argument with a case-split on the binding consumption regime. Define for $T \geq 0$ and $\tau \in [0, p]$ the daughter's value at allowance T when she pays τ out of pocket,

$$\hat{v}(T, \tau) \equiv \max_{C_1, C_2 \geq 0} u(C_1) + u(C_2) \quad \text{s.t.} \quad C_1 \leq Y_1^D + T - \tau, \quad C_1 \geq \phi T, \quad C_1 + C_2 \leq Y_1^D + Y_2^D + T - \tau.$$

Then $v(T) = \hat{v}(T, 0)$, $v_p(T) = \hat{v}(T, p)$, and the Layer-3 loss is $L(T) \equiv \hat{v}(T, 0) - \hat{v}(T, p)$.

For fixed (T, τ) in the feasible domain, the optimal period-1 consumption is given by the projection

formula

$$C_1^*(T, \tau) = \min \left\{ \max \left\{ \frac{Y_1^D + Y_2^D + T - \tau}{2}, \phi T \right\}, Y_1^D + T - \tau \right\}, \quad C_2^*(T, \tau) = Y_1^D + Y_2^D + T - \tau - C_1^*(T, \tau).$$

By the envelope theorem, for any $\tau \in [0, p]$ at which $\hat{v}$ is differentiable in τ , $-\partial \hat{v}(T, \tau) / \partial \tau = \rho(T, \tau)$, where $\rho(T, \tau) \geq 0$ is the marginal utility cost to the daughter of the out-of-pocket payment τ (which tightens both the liquidity cap and the lifetime budget by τ). Three exhaustive cases for which constraint binds at the daughter's optimum:

- *Smoothing* (neither liquidity nor in-kind floor binds): $C_1 = C_2 = (Y_1^D + Y_2^D + T - \tau)/2$, so $\rho(T, \tau) = u'((Y_1^D + Y_2^D + T - \tau)/2)$.
- *Liquidity-bound* (the liquidity cap binds, in-kind slack): $C_1 = Y_1^D + T - \tau$, $C_2 = Y_2^D$, so $\rho(T, \tau) = u'(Y_1^D + T - \tau)$.
- *In-kind-bound* (the in-kind floor binds): $C_1 = \phi T$, $C_2 = Y_1^D + Y_2^D + (1 - \phi)T - \tau$, so $\rho(T, \tau) = u'(Y_1^D + Y_2^D + (1 - \phi)T - \tau)$.

In each regime, $\rho(T, \tau)$ equals u' evaluated at the relevant bottleneck consumption component: $u'(C_1^*(T, \tau))$ in the liquidity-bound and smoothing regimes, $u'(C_2^*(T, \tau))$ in the in-kind-bound regime. From the projection formula above, the bottleneck argument of u' is weakly increasing in T in each regime (the liquidity cap and the smoothing target are linear in T with slope 1 and 1/2, and the in-kind-bound period-2 component is linear in T with slope $1 - \phi$); the bottleneck matches continuously at the liquidity/smoothing and smoothing/in-kind boundaries. (The one-sided slopes of C_1^* in T differ across regimes — they are 1, 1/2, and ϕ in the liquidity-bound, smoothing, and in-kind-bound regimes respectively — but the adjacent envelope expressions for $\rho(T, \tau)$ agree in value at the projection kinks; the kink itself is handled by the convex-hull superdifferential below.) By strict concavity of u , u' is strictly decreasing, hence $\rho(T, \tau)$ is weakly decreasing in T at every τ . The value function $\hat{v}$ is concave in τ at fixed T (the feasible set shrinks linearly in τ via the $-\tau$ shift of the liquidity and lifetime caps); at a switch between two regimes, the superdifferential $\partial_\tau \hat{v}(T, \tau)$ is the convex hull of the two adjacent one-sided derivatives. Since both adjacent derivatives are weakly decreasing in T , every measurable selection of $\rho(T, \tau)$ is weakly decreasing in T as well. Integrating,

$$L(T) = \hat{v}(T, 0) - \hat{v}(T, p) = \int_0^p \rho(T, \tau) d\tau,$$

which is weakly decreasing in T .

At $T = 0$ the in-kind floor is slack ($\phi \cdot 0 = 0$) and there is no allowance to consume; since $Y_1^D < Y_2^D$ the daughter's hide-branch problem is liquidity-bound throughout $\tau \in [0, p]$, so $\hat{v}(0, \tau) = u(Y_1^D - \tau) + u(Y_2^D)$ whenever $Y_1^D > p$ (the strict-positive-consumption condition at $\tau = p$). Hence $L(0) =$

$u(Y_1^D) - u(Y_1^D - p) = -\xi_2^D$. Monotonicity then gives $L(T_{NP}^*) \leq L(0) = -\xi_2^D$, i.e., $-\xi_3^D \leq -\xi_2^D$, equivalently $\xi_3^D \geq \xi_2^D$.

When self-funding was infeasible in Layer 2 ($Y_1^D \leq p$) but the parental allowance makes Layer-3 self-funding feasible, the threshold rises from $-\infty$ to a finite value, strictly above ξ_2^D .

Parental-coverage threshold. The parental-coverage threshold $\bar{\xi}^P = u(Y^P - T_{NP}^* - p) - u(Y^P - T_{NP}^*)$ is constructed in Lemma 2 and is independent of ξ^D on the complier-relevant region. The self-funding characterization follows from the hide-branch binary comparison: under $S = 0$, $k = 0$, the daughter chooses $A = 1$ iff $v_p(T_{NP}^*) - \xi^D \geq v(T_{NP}^*)$, i.e., iff $\xi^D \leq v_p(T_{NP}^*) - v(T_{NP}^*) = \xi_3^D$, whenever self-funding is feasible in the sense of Appendix E.1; otherwise $\xi_3^D = -\infty$ by convention.

Proof of part 1(d) (thresholds determined by the environment). $\bar{\xi}_l^D = 0$ in all layers (a constant). ξ_1^D is the value of $W(Y_1^D + Y_2^D - p) - W(Y_1^D + Y_2^D)$ (or $-\infty$), depending only on $(Y_1^D + Y_2^D, p)$ and u . ξ_2^D is the value of $u(Y_1^D - p) - u(Y_1^D)$ (or $-\infty$), depending on (Y_1^D, p, u) . $\xi_3^D = v_p(T_{NP}^*) - v(T_{NP}^*)$ depends on $(Y_1^D, Y_2^D, T_{NP}^*, \phi, p, u)$, where T_{NP}^* depends only on $(Y_1^D, Y_2^D, Y^P, \phi, \alpha^P, u)$ (proof of Lemma 4, Step 5). $\bar{\xi}^P = u(Y^P - T_{NP}^* - p) - u(Y^P - T_{NP}^*)$ depends only on (Y^P, T_{NP}^*, p, u) , hence only on the environmental primitives. None depends on realized attitudes. $\square$

F.6 The complier region

This subsection derives the complier region (5). The complier indicator ΔA_l equals one only if $\xi^D \in (\xi_l^D, \bar{\xi}_l^D]$. In Layers 1 and 2, $\Delta A_l(\xi^D) = \mathbf{1}[\xi^D \in (\xi_l^D, 0]]$; in Layer 3, this daughter-attitude interval is necessary but not sufficient, and the complier set is the mismatch rectangle $\mathcal{C} = (\xi_3^D, 0] \times (\bar{\xi}^P, \infty)$, so that within that interval the subsidy is decisive only when parents are unwilling to cover. The argument uses only the threshold structure of Lemma 1 together with the characterization of Layer-3 equilibrium outcomes (Lemma 4, part b).

The daughter is a complier iff she aborts under $S = 1$ but not under $S = 0$. By Lemma 1(a), she aborts under $S = 1$ iff $\xi^D \leq 0$.

Layer 1. She aborts under $S = 0$ iff $\xi^D \leq \xi_1^D$. The complier set is $(\xi_1^D, 0]$ in ξ^D , and $\Delta A_1(\xi^D) = \mathbf{1}[\xi^D \in (\xi_1^D, 0]]$.

Layer 2. Same argument with ξ_2^D in place of ξ_1^D : complier set $(\xi_2^D, 0]$, indicator $\Delta A_2(\xi^D) = \mathbf{1}[\xi^D \in (\xi_2^D, 0]]$.

Layer 3. The complier set in ξ^D depends on ξ^P . By the equilibrium characterization (Lemma 4, part b):

- For $\xi^P > \bar{\xi}^P$ (unwilling parents): on the complier-relevant interval $\xi^D \in (\xi_3^D, 0]$, parents refuse coverage (Lemma 3) and the daughter does not self-fund, so $A = 0$ under $S = 0$ (her disclosure, $R = \mathbf{1}[\xi^P \leq 0]$, does not affect the outcome). On the self-fund region $\xi^D \leq \xi_3^D$, parental coverage

on disclosure follows (22) (Lemma 3), so the daughter aborts at $S = 0$ via either hidden self-funding or disclosed parental coverage; either way $A = 1$. On the never-taker region $\xi^D > 0$, $A = 0$ under both S . Hence the complier set is $(\underline{\xi}_3^D, 0]$, with indicator $\Delta A_3 = \mathbf{1}[\xi^D \in (\underline{\xi}_3^D, 0]]$.

- For $\xi^P \leq \bar{\xi}^P$ (willing parents): split by ξ^D . On the complier-relevant interval $\xi^D \in (\underline{\xi}_3^D, 0]$, Lemma 3 gives disclosure, parental coverage, and abortion under both $S = 0$ and $S = 1$ (always-taker, $\Delta A_3 = 0$). On the self-fund region $\xi^D \leq \underline{\xi}_3^D$, parental coverage on disclosure follows (22) (Lemma 3) and the daughter aborts at $S = 0$ via either route (always-taker, $\Delta A_3 = 0$); the parental-coverage threshold $\bar{\xi}^P$ does not apply on this region (see Appendix F.7, Case C, below), but the abortion outcome is invariant. Thus $\Delta A_3 = 0$ whenever $\xi^P \leq \bar{\xi}^P$ and $\xi^D \leq 0$; on the self-fund subregion this is an outcome classification, not a financing classification.

The Layer-3 outcomes are summarized in the table below.

Region in ξ^D	Parents	$S = 0$	$S = 1$	Type
$\xi^D > 0$	any	no	no	never-taker
$\xi^D \leq \underline{\xi}_3^D$	any	yes	yes	always-taker
$\xi^D \in (\underline{\xi}_3^D, 0]$	willing	yes (covered)	yes	always-taker
$\xi^D \in (\underline{\xi}_3^D, 0]$	unwilling	no	yes	complier

“Yes/no” denote whether the daughter aborts in equilibrium; “willing” means $\xi^P \leq \bar{\xi}^P$ and “unwilling” means $\xi^P > \bar{\xi}^P$.

F.7 The subsidy’s effect (proof of Proposition 1)

We prove the two parts in turn, using the complier characterization derived in Appendix F.6 together with the threshold structure of Lemma 1.

Proof of part 1(a) (monotonicity in ξ^D). In Layers 1 and 2, the complier indicator derived in Appendix F.6 $\Delta A_l = \mathbf{1}[\xi^D \in (\underline{\xi}_l^D, 0]]$ is zero outside the interval and one on it: non-monotone in ξ^D when $\underline{\xi}_l^D$ is finite. When self-funding is infeasible ($\underline{\xi}_l^D = -\infty$), the indicator reduces to $\mathbf{1}[\xi^D \leq 0]$, weakly decreasing. In Layer 3, the same dichotomy holds by case on ξ^P : in the unwilling-parent region ($\xi^P > \bar{\xi}^P$) the indicator in ξ^D is $\mathbf{1}[\xi^D \in (\underline{\xi}_3^D, 0]]$, non-monotone when $\underline{\xi}_3^D$ is finite and weakly decreasing when $\underline{\xi}_3^D = -\infty$; in the willing-parent region ($\xi^P \leq \bar{\xi}^P$) the indicator is identically zero on $\xi^D \leq 0$ (always-takers), hence trivially monotone.

Proof of part 1(b) (parents’ distaste). We show $\xi^P \mapsto \Delta A_3(\xi^D, \xi^P)$ is weakly increasing for each ξ^D .

Case A: $\xi^D > 0$. Never-taker; $\Delta A_3 = 0$ for all ξ^P (constant, hence weakly increasing).

Case B: $\xi^D \in (\underline{\xi}_3^D, 0]$. For $\xi^P \leq \bar{\xi}^P$: she discloses, parents cover, she aborts under $S = 0$ (always-taker, $\Delta A_3 = 0$). For $\xi^P > \bar{\xi}^P$: parents refuse coverage and she does not self-fund (since $\xi^D > \underline{\xi}_3^D$), so $A = 0$ under $S = 0$; complier ($\Delta A_3 = 1$). The function jumps from 0 to 1 at $\bar{\xi}^P$ — weakly increasing.

Case C: $\xi^D \leq \underline{\xi}_3^D$. Self-funder. Under $S = 0$ the daughter aborts regardless of the parental coverage choice: if she hides, hidden self-funding is feasible and she aborts; if she discloses, the parental coverage comparison differs from Lemma 2 (whose threshold $\bar{\xi}^P$ governs only the complier-relevant region where the daughter would not self-fund under $k = 0$). On the self-fund region the direct abortion-distaste term $-\xi^P$ is present under both $k = 0$ and $k = 1$ (since $A = R = 1$ in both), so the parents' choice depends on whether the altruistic benefit of sparing the daughter p exceeds the cost of paying p , not on $\bar{\xi}^P$. Either selection delivers $A = 1$. Hence $A = 1$ for all ξ^P under both $S = 0$ and $S = 1$: always-taker, $\Delta A_3 = 0$ constant.

In each case, ΔA_3 is weakly increasing in ξ^P . □

G A model with involvement stigma

In the baseline model of Section 4, disclosure carries no direct cost, so a daughter whose parents would not cover the procedure is indifferent between hiding and telling. The main text resolves this indifference through a tie-breaking convention: the daughter tells parents who weakly approve of the abortion and hides from those who oppose it (see Appendix E.4). This appendix shows that the convention is not arbitrary—it is the strict best response once involving one’s parents carries a cost proportional to their abortion distaste—and that the theoretical results in Section 4.3 are unchanged.

G.1 Setup

The model is exactly that of Appendix E, with one addition: when the daughter discloses her pregnancy she incurs a utility term $g \xi^P$. Her payoff (3) becomes

$$\tilde{V}^D = u(C_1^D) + u(C_2^D) - \xi^D \cdot A - g \xi^P \cdot R, \quad g > 0, \quad (26)$$

where $R \in \{0, 1\}$ is the disclosure indicator. Involving one’s parents is thus costly when they oppose the abortion ($\xi^P > 0$) and beneficial when they support it ($\xi^P < 0$): a daughter with supportive parents would rather have them know, a daughter with opposed parents would rather avoid the conflict, and the effect is stronger the more strongly the parents feel. The parents’ payoff (4) is unchanged: consistent with the limited-altruism formulation, parents do not internalize the daughter’s involvement term (any cost they perceive is already folded into ξ^P). All other primitives are those of Appendix E.1.

The disclosure term enters the daughter’s payoff additively and appears in no budget constraint, in no parental payoff, and not in the non-pregnancy problem (where $R = 0$). The following objects are therefore literally the same as in Appendix F, and their constructions apply with no change:

- the consumption value functions v and v_p of (23)–(24);
- the pre-committed parental allowance T_{NP}^* (Lemma 4, Step 5);
- the Layer-3 self-funding threshold $\xi_3^D = v_p(T_{NP}^*) - v(T_{NP}^*)$ (a hide-branch object, where $R = 0$);
- the parental-coverage threshold $\bar{\xi}^P = u(Y^P - T_{NP}^* - p) - u(Y^P - T_{NP}^*)$ (Lemma 2), which compares the parents’ own payoffs and so does not involve g .

Only the daughter’s stage-(iii) disclosure decision changes. We re-derive it, then verify the summary results.

G.2 Disclosure with stigma

Lemma 5 (Disclosure under stigma). *Fix the stigma model (26) with $g > 0$.*

- (a) Outcomes are unchanged. The thresholds $\bar{\xi}^P, \underline{\xi}_3^D$, the parents' coverage rule, and the abortion outcome A^* are exactly those of the baseline model (Lemma 3): the term $g \xi^P R$ enters the daughter's payoff only through R and changes neither her abortion decision nor the parents' willingness to cover.
- (b) The tie-break is strict. Wherever the baseline model leaves disclosure payoff-equivalent—that is, wherever disclosing secures no coverage, because the parents would refuse or because the procedure is free ($S = 1$)—the daughter strictly discloses when $\xi^P < 0$ and hides when $\xi^P > 0$; the locus $\xi^P = 0$ is null and is resolved by the convention $R = 1$. The baseline tie-break $R = \mathbf{1}[\xi^P \leq 0]$ is thus the strict best response wherever it is invoked, for every $g > 0$.

Proof. (a) The term $g \xi^P R$ is additively separable in R , vanishes at $R = 0$, and appears in no budget constraint, no parental payoff, and not in the non-pregnancy problem. Hence $v, v_p, T_{\text{NP}}^*, \underline{\xi}_3^D, \bar{\xi}^P$, and the parents' coverage rule are exactly as in the baseline (Appendix G.1). Given her disclosure choice, the stigma term is a constant in the daughter's abortion problem, so her abortion best response—maximize consumption net of $\xi^D A$ —is the baseline one. Disclosure can move the *realized* abortion only through coverage, i.e. only where telling secures coverage; there $\xi^P \leq \bar{\xi}^P < 0$ (Lemma 2), so the stigma term $-g \xi^P > 0$ rewards disclosure and the daughter discloses exactly as in the baseline; the abortion outcome is therefore unchanged. Thus $A^*(\xi^D, \xi^P)$ and the thresholds coincide with Lemma 3.

(b) Where disclosing secures no coverage, consumption and the abortion outcome are identical under $R = 1$ and $R = 0$, so the two branches differ only by the disclosure term: the daughter's payoff is higher under $R = 1$ by $-g \xi^P$. She therefore discloses iff $-g \xi^P \geq 0$, i.e. iff $\xi^P \leq 0$, strictly for $\xi^P \neq 0$; the tie at the null locus $\xi^P = 0$ is fixed by the convention. $\square$

G.3 Robustness of the summary results

Because the abortion outcome A^* is unchanged (Lemma 5), the Layer-3 complier indicator is identical to the baseline:

$$\Delta A_3(\xi^D, \xi^P) = \mathbf{1}[\xi^D \in (\underline{\xi}_3^D, 0]] \cdot \mathbf{1}[\xi^P > \bar{\xi}^P], \quad (27)$$

the mismatch rectangle $(\underline{\xi}_3^D, 0] \times (\bar{\xi}^P, \infty)$.

Proposition 5 (Robustness to involvement stigma). *Under the stigma model (26) with $g > 0$, Proposition 1 and Lemma 1 hold verbatim: the thresholds, the complier region, and the subsidy's effect on abortion decisions are exactly those of the baseline model.*

Proof. Every threshold in Lemma 1 is determined without reference to disclosure: the never-taker thresholds $\bar{\xi}_i^D = 0$ are pinned down under $S = 1$, where the daughter aborts irrespective of R ; the always-taker/self-funding thresholds $\underline{\xi}_i^D$ are hide-branch objects ($R = 0$); and $\bar{\xi}^P$ and T_{NP}^* are parental objects unaffected by g (Appendix G.1). Hence Lemma 1 carries over unchanged. By Lemma 5 the

abortion outcome, and therefore the complier indicator (27), coincide with the baseline; Proposition 1, a statement about that indicator, holds verbatim. $\square$

Stigma thus converts the parsimonious model's disclosure indifference into strict behavior without disturbing any abortion outcome. In particular, the concealment pattern of Proposition 6—the subsidy's abortions are hidden precisely from the parents who oppose them—is strict equilibrium behavior under any $g > 0$, not an artifact of the tie-breaking convention.

H Proof of Proposition 2

We prove the three claims in turn: (i-a) the post-subsidy ratio $A(1 \mid h)$ is weakly lower in the strict community; (i-b) the subsidy-induced share $\Delta A(h)/A(1 \mid h)$ is weakly higher in the strict community; and (ii) without further assumptions, the level $\Delta A(h)$ can move in either direction. Throughout, we use the coordinate order (ξ^P, ξ^D) of Assumption 1. Here $\Delta A(h) \equiv \Pr_{f_h}(\mathcal{C})$ is the community-level subsidy effect, the population analogue of the family-level indicator $\Delta A_3(\xi^D, \xi^P)$, whose complier set $\mathcal{C}$ is the mismatch rectangle (5). Assumption 3 is imposed community by community on f_h .

Preliminary: MLR and upper-set probabilities. Assumption 1 states the lattice inequality $f_1(x)f_0(y) \leq f_1(x \vee y)f_0(x \wedge y)$ for all x, y in the common support, and requires the common support to be a sublattice of $\mathbb{R}^2$. This is the bivariate MLR order of Karlin and Rinott (1980); see also Shaked and Shanthikumar (2007, Section 6.E). We work with the lattice form throughout rather than likelihood ratios f_1/f_0 , which sidesteps undefined ratios at zero-density points within the common support, and treat f_h as densities on $\mathbb{R}^2$ extended by zero outside their common support. Because the common support $\mathcal{S}$ is a sublattice (Assumption 1) and $\mathcal{N} = \{\xi^D \leq 0\}$ is also a sublattice, the conditional common support $\mathcal{S} \cap \mathcal{N}$ is a sublattice. Thus conditioning on $\mathcal{N}$ preserves the lattice structure used in Step 2 below.

For comparable pairs $x \leq y$ (coordinatewise), $x \vee y = y$ and $x \wedge y = x$, so the inequality reduces to $f_1(x)f_0(y) \leq f_1(y)f_0(x)$ — the standard univariate MLR ordering of f_1 over f_0 along any chain. For non-comparable pairs the inequality extends the ordering to the lattice.

The standard consequence is that the bivariate likelihood-ratio order implies the usual multivariate stochastic order: for any *upper set* $U \subseteq \mathbb{R}^2$ (i.e., a set such that $(x, y) \in U$ and $(x', y') \geq (x, y)$ coordinatewise imply $(x', y') \in U$), $\Pr_{f_1}(U) \geq \Pr_{f_0}(U)$; dually $\Pr_{f_1}(L) \leq \Pr_{f_0}(L)$ for lower sets L . This is Karlin and Rinott (1980, Theorem 2.2); see also Shaked and Shanthikumar (2007, Theorem 6.E.8).

We additionally use the assumed regularity $A(1 \mid h) = \Pr_{f_h}(\xi^D \leq 0) > 0$ for both h , which ensures the conditional probabilities $\Pr_{f_h}(\cdot \mid \mathcal{N})$ are well-defined.

Step 1: post-subsidy ratio. By Lemma 1(a), $A(1 \mid h) = \Pr_{f_h}(\xi^D \leq 0)$. The set $\{(\xi^P, \xi^D) : \xi^D \leq 0\}$ is a lower set in $\mathbb{R}^2$ (coordinate-decreasing in both coordinates is preserved). The MLR consequence above gives $\Pr_{f_1}(\xi^D \leq 0) \leq \Pr_{f_0}(\xi^D \leq 0)$, i.e., $A(1 \mid 1) \leq A(1 \mid 0)$.

Step 2: subsidy-induced share. By the complier characterization (5), the Layer-3 complier set is the mismatch rectangle

$$\mathcal{C} \equiv \{(\xi^P, \xi^D) : \xi^P > \bar{\xi}^P, \xi^D \in (\underline{\xi}_3^D, 0]\}.$$

Let $\mathcal{N} \equiv \{(\xi^P, \xi^D) : \xi^D \leq 0\}$ be the post-subsidy abortion set. Then $\mathcal{C} \subseteq \mathcal{N}$ and $A(1 | h) = \Pr_{f_h}(\mathcal{N})$, so the subsidy-induced share is the conditional probability

$$\frac{\Delta A(h)}{A(1 | h)} = \Pr_{f_h}(\mathcal{C} | \mathcal{N}).$$

We show $\Pr_{f_1}(\mathcal{C} | \mathcal{N}) \geq \Pr_{f_0}(\mathcal{C} | \mathcal{N})$.

Fact (a). $\mathcal{C}$ is an upper set within $\mathcal{N}$. For any $\xi^D \leq 0$, the ξ^P -section of $\mathcal{C}$ is:

- empty if $\xi^D \leq \underline{\xi}_3^D$;
- $(\bar{\xi}^P, \infty)$ if $\xi^D \in (\underline{\xi}_3^D, 0]$.

Each ξ^P -section is either empty or an upper interval, so the indicator $\mathbf{1}[(\xi^P, \xi^D) \in \mathcal{C}]$ is weakly increasing in ξ^P . As ξ^D increases (from $-\infty$ to 0), the ξ^P -section weakly grows (from empty to upper interval), so the indicator is also weakly increasing in ξ^D . Hence $\mathcal{C}$ is an upper set in (ξ^P, ξ^D) within $\mathcal{N}$ (the upper-set property of $\mathcal{C}$ may fail outside $\mathcal{N}$ — at $\xi^D > 0$, the daughter is a never-taker and the complier indicator drops to zero — but we condition on $\mathcal{N}$, where this is moot).

Fact (b). *MLR is preserved under conditioning on a sublattice.* The set $\mathcal{N} = \{(\xi^P, \xi^D) : \xi^D \leq 0\}$ is closed under coordinatewise maximum and minimum, so it is a sublattice of $\mathbb{R}^2$ (and a lower set). Let Q_h denote the conditional distribution of (ξ^P, ξ^D) given $\mathcal{N}$ under f_h , with density $q_h \equiv f_h \cdot \mathbf{1}_{\mathcal{N}} / \Pr_{f_h}(\mathcal{N})$ on $\mathcal{N}$. The positive normalizing constant $1 / \Pr_{f_h}(\mathcal{N})$ is independent of (ξ^P, ξ^D) (and finite by the assumed regularity $A(1 | h) > 0$), so for any $x, y \in \mathcal{N}$, $x \vee y, x \wedge y \in \mathcal{N}$ as well (sublattice), and

$$q_1(x)q_0(y) \propto f_1(x)f_0(y) \leq f_1(x \vee y)f_0(x \wedge y) \propto q_1(x \vee y)q_0(x \wedge y),$$

with the same positive proportionality constant on both sides. So Q_1 MLR-dominates Q_0 on $\mathcal{N}$ in the lattice sense. The lattice form (rather than ratios) makes the argument robust to zero-density points within $\mathcal{N}$. In particular, for any upper set $V \subseteq \mathcal{N}$, $Q_1(V) \geq Q_0(V)$.

Combining (a) and (b). Applying Fact (b) to the upper set $\mathcal{C} \subseteq \mathcal{N}$ (Fact a):

$$Q_1(\mathcal{C}) \geq Q_0(\mathcal{C}),$$

i.e.,

$$\Pr_{f_1}(\mathcal{C} | \mathcal{N}) \geq \Pr_{f_0}(\mathcal{C} | \mathcal{N}),$$

which is the second claim.

Step 3: ambiguity of the level $\Delta A(h)$. The level is the product $\Delta A(h) = A(1 | h) \cdot [\Delta A(h) / A(1 | h)]$. By Steps 1 and 2, the first factor weakly falls and the second weakly rises as h increases from 0 to 1.

Their product can go either way. We exhibit two MLR-ordered pairs with opposite level comparisons.

Nondegeneracy. For the examples to deliver strict comparisons, fix an economic environment with $\bar{\xi}^P$ finite (so the event $\xi^P > \bar{\xi}^P$ is non-trivial) and nonempty mismatch interval $(\underline{\xi}_3^D, 0]$. We take g and ψ in both constructions to be strictly positive log-concave densities on $\mathbb{R}$ (e.g., the standard normal), so that f_0 and f_1 have common support $\mathbb{R}^2$, the regularity $A(1 | h) > 0$ holds for both h , and the mismatch rectangle carries positive mass.

Level rises with strictness. Let ψ be any density on $\mathbb{R}$ with positive mass on $(\underline{\xi}_3^D, 0]$ (so the complier region has positive baseline mass for some ξ^P), and let g be a log-concave density on $\mathbb{R}$ (e.g., the standard normal). For $\delta > 0$, define $f_0(\xi^P, \xi^D) = g(\xi^P) \cdot \psi(\xi^D)$ and $f_1(\xi^P, \xi^D) = g(\xi^P - \delta) \cdot \psi(\xi^D)$. The ratio $f_1/f_0 = g(\xi^P - \delta)/g(\xi^P)$ is increasing in ξ^P (by log-concavity of g , see Bagnoli and Bergstrom 2005) and constant in ξ^D ; hence f_1/f_0 is weakly increasing on the lattice, so Assumption 1 holds. The daughter marginal is unchanged across h , so $A(1 | 1) = A(1 | 0)$. The mismatch region $\{\xi^P > \bar{\xi}^P, \xi^D \in (\underline{\xi}_3^D, 0]\}$ shifts mass to the right in ξ^P under f_1 , so $\Pr_{f_1}(\text{mismatch}) > \Pr_{f_0}(\text{mismatch})$, i.e., $\Delta A(1) > \Delta A(0)$.

Level falls with strictness. Hold the parental marginal fixed across h (e.g., $g_0 = g_1 = g$, any log-concave density). Choose the daughter marginal as a location family $\psi_h(\xi^D) = \psi(\xi^D - \delta_h)$ with $\delta_0 = 0$ and δ_1 large. With ψ log-concave (e.g., standard normal), the ratio $\psi_1/\psi_0 = \psi(\xi^D - \delta_1)/\psi(\xi^D)$ is increasing in ξ^D , so f_1/f_0 is weakly increasing on the lattice and Assumption 1 holds. For any baseline f_0 with $\Delta A(0) > 0$, as $\delta_1 \rightarrow \infty$: $A(1 | 1) = \Pr_{\psi_1}(\xi^D \leq 0) \rightarrow 0$ (mass shifts right out of $(-\infty, 0]$), and $\Delta A(1) \leq A(1 | 1) \rightarrow 0$. Meanwhile $\Delta A(0)$ is fixed and positive. So $\Delta A(1) < \Delta A(0)$ for δ_1 large enough.

Both constructions satisfy Assumption 1. Hence the level effect is sign-ambiguous under the assumption alone. $\square$

Remark on the strength of Assumption 1. The proof uses two consequences of bivariate MLR: (i) lower-set probability ordering for the marginal in ξ^D (Step 1), and (ii) preservation of MLR under conditioning on the lower set $\mathcal{N} = \{\xi^D \leq 0\}$ (Step 2). Both follow directly from the definition of MLR. We retain bivariate MLR in the statement because it is verifiable from the joint density alone and is closed under the operations we use.

I Model calibration

This appendix details the calibration summarized in Section 5.2.

Road map. As illustrated in Figure 3, Panel (c), the abortion ratios with and without the subsidy, $A(1 | h)$ and $A(0 | h)$, are given by:

- $A(1 | h)$: the share of families in the always-taker or complier regions;
- $A(0 | h)$: the share of families in the always-taker region.

Therefore, $A(1 | h)$ and $A(0 | h)$ depend only on the joint distribution of abortion attitudes, (ξ^D, ξ^P) , and on the thresholds ξ_3^D , $\bar{\xi}_3^D$, and $\bar{\xi}^P$.

We use this framework to validate the model in five steps:

- **Step 1.** Estimate the distribution of abortion attitudes using the WVS survey.
- **Step 2.** Recover abortion ratios from the difference-in-differences estimates in Table 1.
- **Step 3.** Find the thresholds ξ_3^D , $\bar{\xi}_3^D$, and $\bar{\xi}^P$ that rationalize the observed abortion ratios for the strict group, given the observed distribution of abortion attitudes in that group.
- **Step 4.** Using these thresholds and the observed distribution of abortion attitudes in the lenient group, predict the lenient-group subsidy effect.
- **Step 5.** Compare the predicted lenient-group subsidy effect from Step 4 with the effect observed in the data.

Step 1: Attitude distributions. For each generation $g \in \{D, P\}$ —daughters (D), proxied by women aged 18–23, and parents (P), proxied by men and women aged 45–75—and each community $h \in \{\text{strict}, \text{lenient}\}$, define the empirical attitude distribution

$$\hat{F}_{g,h}(c) \equiv \Pr(\text{score} < c \mid g, h),$$

where score is the respondent’s answer to the WVS question on whether abortion is justifiable, measured on a scale from 1 (never justifiable) to 10 (always justifiable). We evaluate $\hat{F}_{g,h}$ at each integer score using the WVS survey and linearly interpolate between adjacent integers.

We use these empirical distributions to discipline the shape of the model distributions of abortion attitudes. Specifically, we assume that ξ^D and ξ^P are distributed according to the daughter and parent WVS distributions, respectively, within each community, up to a location shift. This shift allows the

WVS scale, which is ordinal and measured in survey units, to differ from the model's utility scale; the calibration below determines where the model thresholds fall on these distributions.

Step 2: Observed abortion ratios. The difference-in-differences estimates in Table 1 deliver, for each community h , two quantities: the abortion ratio absent the subsidy and the effect of the subsidy on that ratio. These map into the model as follows. The pre-subsidy ratio is $A(0 \mid h)$, the share of pregnancies ending in abortion when the procedure is not subsidized. The difference-in-differences estimate is the subsidy effect, $\Delta A(h) \equiv A(1 \mid h) - A(0 \mid h)$. The post-subsidy ratio is then recovered as $A(1 \mid h) = A(0 \mid h) + \Delta A(h)$.

Remark I.1 (Survey-scale interpretation of the recovered thresholds). The calibration recovers cutoffs on the WVS survey scale, which maps into the model's utility scale up to an unknown location shift. For notational simplicity, we continue to denote the survey-scale cutoffs by $\underline{\xi}_3^D$, $\bar{\xi}_3^D$, and $\bar{\xi}^P$. Accordingly, the recovered value of $\bar{\xi}_3^D$ need not equal zero in survey units: it is the point on the WVS daughter distribution corresponding to the model cutoff $\bar{\xi}_3^D = 0$ in utility units.

Step 3: Thresholds. The model has three unknown thresholds: $\underline{\xi}_3^D$, $\bar{\xi}_3^D$, and $\bar{\xi}^P$. In addition, we do not observe the copula linking parents' and daughters' attitudes, since the survey reports only marginal distributions. By contrast, the strict-group abortion data provide only two moments, $A(0 \mid h)$ and $A(1 \mid h)$, with which to discipline these objects. We therefore consider several approaches to this underidentification problem. In all cases, the output is a set of thresholds, $\underline{\xi}_3^D$, $\bar{\xi}_3^D$, and $\bar{\xi}^P$, estimated using only data from the strict community. Steps 4 and 5 then use these thresholds in the same way, regardless of the approach used to estimate them. We first describe Steps 4 and 5, and then return to the alternative implementations of Step 3.

Step 4: Predicted abortion ratios. The thresholds $\underline{\xi}_3^D$, $\bar{\xi}_3^D$, and $\bar{\xi}^P$ do not depend on the distribution of abortion attitudes (Lemma 1(d)). Therefore, the thresholds recovered in Step 3 using data from the strict community also apply to the lenient community. We combine these thresholds with the lenient-group attitude distribution from Step 1 to compute

$$\begin{aligned}\widehat{A}(1 \mid \text{lenient}) &= \Pr(\text{always-taker or complier} \mid \text{lenient}), \\ \widehat{A}(0 \mid \text{lenient}) &= \Pr(\text{always-taker} \mid \text{lenient}),\end{aligned}$$

and the predicted lenient-group subsidy effect,

$$\widehat{\Delta A}(\text{lenient}) = \widehat{A}(1 \mid \text{lenient}) - \widehat{A}(0 \mid \text{lenient}).$$

Step 5: Model fit. We compare the predicted lenient-group subsidy effect from Step 4 with the effect observed in the data, as reported in Table 1. Because the thresholds used in Step 4 are disciplined using only strict-group data, the predicted lenient-group effect is a genuine out-of-sample prediction: no lenient-group abortion ratio enters its construction. A close match therefore indicates that the observed differences in abortion-attitude distributions are sufficient to account for the cross-community difference in subsidy effects.

I.1 Step 3: Alternative approaches

We now describe four approaches to recovering the model's three unknown thresholds: ξ_3^D , $\bar{\xi}_3^D$, and $\bar{\xi}^P$. Two observations guide the exercise. First, $A(1 | h)$ is the share of daughters who would abort if the procedure were free. It therefore depends only on the daughter's marginal attitude distribution and the never-taker threshold $\bar{\xi}_3^D$. Thus, the strict-group moment $A(1 | \text{strict})$ pins down $\bar{\xi}_3^D$ directly in all four approaches.

Second, Approaches 1–3 assume that daughters cannot self-finance. Under this assumption, the self-financing threshold drops out, $\xi_3^D \rightarrow -\infty$, and the always-taker region consists only of families in which the daughter would abort and her parents are willing to cover the procedure. The remaining strict-group moment, $A(0 | \text{strict})$, must therefore discipline two objects: the parental-coverage threshold $\bar{\xi}^P$ and the within-family dependence of attitudes. We summarize this dependence using a Gaussian copula with correlation $\rho \geq 0$, where $\rho = 0$ denotes independent attitudes and larger values of ρ denote more concordant parent-daughter attitudes. Because one moment cannot separately identify both $\bar{\xi}^P$ and ρ , Approaches 1–3 differ in how they treat ρ : Approach 1 imposes independence, $\rho = 0$, while Approaches 2–3 vary ρ over alternative values. Approach 4 then relaxes the no-self-financing assumption.

Approach 1: point prediction under independence. We set $\rho = 0$. Under independence, the joint distribution factorizes, so the always-taker share is

$$A(0 | h) = A(1 | h) \Pr(\xi^P \leq \bar{\xi}^P | h).$$

The strict-group moments therefore identify the parental-coverage threshold as the $\bar{\xi}^P$ satisfying

$$\Pr(\xi^P \leq \bar{\xi}^P | \text{strict}) = \frac{A(0 | \text{strict})}{A(1 | \text{strict})}.$$

Approach 2: common correlation. We allow $\rho > 0$ but impose the same value of ρ in both communities. For each $\rho \in [0, 0.99]$, we calibrate $\bar{\xi}^P$ as follows. The strict-group moment $A(1 | \text{strict})$ continues

to pin down $\bar{\xi}_3^D$. Holding this threshold fixed, and given a candidate value of ρ , the Gaussian copula implies a predicted value of $A(0 \mid \text{strict})$ for each parental-coverage threshold $\bar{\xi}^P$. Because this mapping is nonlinear, we use a numerical solver to find the value of $\bar{\xi}^P$ that matches the observed $A(0 \mid \text{strict})$ for each ρ .

Approach 3: asymmetric correlation. We relax the assumption that the two communities share the same copula and allow separate correlations, ρ_{strict} and ρ_{lenient} . We fix $\rho_{\text{strict}} = 0.5$ and recover $\bar{\xi}^P$ using the same procedure as in Approach 2. We then predict the lenient-group subsidy effect by sweeping over $\rho_{\text{lenient}} \in [0, 0.99]$, obtaining a separate prediction for each value of ρ_{lenient} .

Approach 4: set-identified bounds. We drop the assumption that daughters cannot self-finance and reintroduce a finite self-financing threshold, ξ_3^D . We impose an additional restriction: a share π^{NS} of daughters cannot self-finance the procedure, regardless of their abortion attitudes. We restrict π^{NS} to be at least as large as the share of daughters in our sample who do not have stable employment. For each $\rho \in [0, 0.99]$, we find all $(\xi_3^D, \bar{\xi}^P)$ that match the strict-group moments, $A(0 \mid \text{strict})$ and $A(1 \mid \text{strict})$, and satisfy the self-finance restriction. We assume a single ρ common to both communities. We then compute the implied lenient-group subsidy effect for each admissible $(\rho, \xi_3^D, \bar{\xi}^P)$. This yields a set of admissible predictions—and therefore bounds on the lenient-group subsidy effect—rather than a single point estimate.


I.2 Results

We now report the results of each of the four approaches. In every case the thresholds are disciplined using only strict-group data and the lenient-group effect is predicted out of sample.

Approach 1. Under independence, the calibrated thresholds imply a point prediction for the lenient group of $\widehat{\Delta A}(\text{lenient}) = 0.050$. This lies within about a percentage point of the estimated $\Delta A(\text{lenient}) = 0.056$ and well below the strict-group effect $\Delta A(\text{strict}) = 0.148$. The difference in marginal attitude distributions alone is therefore large enough to account for the cross-community difference in subsidy effects.

Approaches 2 and 3. Figure [I.1](#) reports the predicted lenient-group effects across alternative values of ρ . Panel [\(a\)](#) reports Approach 2, and Panel [\(b\)](#) reports Approach 3. In both approaches, the predicted lenient-group effect declines as ρ rises: stronger within-family correlation in abortion attitudes places daughters who are more willing to abort in families with parents who are more willing to cover the procedure, shrinking the mass of families in the mismatch region and therefore the predicted subsidy

Figure I.1: Predicted lenient subsidy effect from the survey attitude marginals



Notes: Predicted lenient-group subsidy effect as a function of the copula correlation ρ between daughter and parent abortion attitudes. Panel (a) reports Approach 2; the x-axis represents a single value of ρ common to both communities. Panel (b) reports Approach 3; the x-axis represents ρ for the lenient community, with ρ for the strict community fixed at 0.5. The black curve is the model's out-of-sample prediction. The dashed lines are the empirical difference-in-differences estimates from Table 1.

effect. Despite this variation, the predicted lenient-group effect is at most about half of the strict-group effect across the full range of ρ . These results reinforce the conclusion that differences in marginal attitude distributions alone are large enough to account for the difference in subsidy effects between strict and lenient communities.

Approach 4. Relaxing the assumption that daughters cannot self-finance replaces the point prediction with a set. Across all admissible configurations and all $\rho \in [0, 0.99]$, the predicted lenient effect lies in $\widehat{\Delta A}(\text{lenient}) \in [0.002, 0.131]$. The estimated lenient effect $\Delta A(\text{lenient}) = 0.056$ falls inside this band, whereas the estimated strict effect $\Delta A(\text{strict}) = 0.148$ falls outside it: even at its most permissive, the calibrated model cannot rationalize a lenient effect as large as the strict one.

Summary. The four approaches deliver the same conclusion. Disciplined only by the strict-group moments and the observed attitude distributions, the model predicts a much smaller subsidy effect in the lenient community—close to the estimated lenient effect and far below the strict-group effect. This conclusion is robust to allowing the within-family correlation in attitudes to vary and to relaxing the assumption that daughters cannot self-finance.

J Proofs for Section 5.3

Throughout this appendix we maintain the model of Section 4 and Assumption 2. We first state and prove an auxiliary lemma on the divergence of the pre-commitment allowance, then derive parts 3(a), 3(b), and 3(c) of Proposition 3 in turn.

Lemma 6 (Allowance and consumption divergence). *Fix $\phi \in [0, 1]$ and (Y_1^D, p, α^P) . Then, as $Y^P \rightarrow \infty$, the pre-commitment allowance $T_{NP}^*(Y^P, Y_2^D) \rightarrow \infty$ uniformly in Y_2^D . Consequently, as $Y^P \rightarrow \infty$, the daughter's period-1 consumption $C_1^D \rightarrow \infty$, and her period-2 consumption $C_2^D \rightarrow \infty$ via either of two channels:*

- (i) *Transfer-cash channel. If $\phi < 1$, $C_2^D \rightarrow \infty$ as $Y^P \rightarrow \infty$, regardless of Y_2^D .*
- (ii) *Own-income channel. If $Y_2^D \rightarrow \infty$, $C_2^D \rightarrow \infty$, regardless of ϕ .*

Proof. We prove the allowance divergence and the consumption claims in turn.

$T_{NP}^* \rightarrow \infty$. Parents solve $\max_{T \in [0, Y^P]} G_H(T) = u(Y^P - T) + \alpha^P v(T)$, where v is defined in (23). By the proof of Lemma 4, Step 5, and the interior-transfer regularity assumption in Appendix E.1, G_H is strictly concave with unique interior maximum $T_{NP}^* > 0$ characterized by the KKT condition $u'(Y^P - T_{NP}^*) = \alpha^P s_H$ for some supergradient $s_H \in \partial v(T_{NP}^*)$.

We rule out $T_{NP}^* \leq \bar{T}$ for any fixed $\bar{T} < \infty$ as $Y^P \rightarrow \infty$. Suppose, for contradiction, $T_{NP}^*(Y_n^P, Y_2^D) \leq \bar{T}$ along $Y_n^P \rightarrow \infty$. We show that every supergradient of v at T_{NP}^* is bounded below by a positive constant depending only on $(Y_1^D, \bar{T}, \phi)$, not on Y_2^D .

The daughter's value $v(T)$ at $T \in [0, \bar{T}]$ is determined by her optimal consumption split, subject to liquidity, in-kind, and non-negative-savings constraints. Three exhaustive cases for the binding constraints:

- *Liquidity-bound* (large Y_2^D , $s = 0$): $C_1^D = Y_1^D + T$, $C_2^D = Y_2^D$, $v'(T) = u'(Y_1^D + T) \geq u'(Y_1^D + \bar{T}) > 0$.
- *Smoothing* (small Y_2^D , no constraint binding): $C_1^D = C_2^D = (Y_1^D + Y_2^D + T)/2$, $v'(T) = u'((Y_1^D + Y_2^D + T)/2)$. Since smoothing requires $Y_2^D \leq Y_1^D + T \leq Y_1^D + \bar{T}$, the argument of u' is at most $Y_1^D + \bar{T}$, so $v'(T) \geq u'(Y_1^D + \bar{T}) > 0$.
- *In-kind-bound* (large ϕT relative to lifetime resources): $C_1^D = \phi T$, $C_2^D = Y_1^D + Y_2^D + (1 - \phi)T$ (cash all goes to period 2). The supergradient of v at T includes $\phi u'(\phi T) + (1 - \phi)u'(Y_1^D + Y_2^D + (1 - \phi)T)$. For $T \leq \bar{T}$ with $\phi > 0$, $\phi u'(\phi T) \geq \phi u'(\phi \bar{T}) > 0$, bounded below uniformly in Y_2^D .

At a regime-switch point, $\partial v(T)$ is the convex hull of the adjacent regime derivatives. Since each adjacent derivative is bounded below by the same positive constant on $[0, \bar{T}]$, every supergradient at the kink is also bounded below by that constant. Combining the interior-regime formulas with the

convex-hull characterization at kinks, every supergradient of v at any $T \leq \bar{T}$ is bounded below by a positive constant $\underline{v} \equiv \min\{u'(Y_1^D + \bar{T}), \phi u'(\phi \bar{T})\} > 0$ (with the second term omitted when $\phi = 0$), independent of Y_2^D .

The parental KKT condition then gives $u'(Y_n^P - T_{NP}^*) = \alpha^P s_H \geq \alpha^P \underline{v} > 0$, which bounds $Y_n^P - T_{NP}^*$ above by $(u')^{-1}(\alpha^P \underline{v}) < \infty$. But $Y_n^P - T_{NP}^* \geq Y_n^P - \bar{T} \rightarrow \infty$, a contradiction. Hence $T_{NP}^*(Y^P, Y_2^D) \rightarrow \infty$ as $Y^P \rightarrow \infty$, uniformly in Y_2^D .

$C_1^D \rightarrow \infty$. Given $T_{NP}^* \rightarrow \infty$: if $\phi > 0$, the in-kind floor $C_1^D \geq \phi T_{NP}^* \rightarrow \infty$ forces $C_1^D \rightarrow \infty$; if $\phi = 0$, the daughter's optimum at allowance T_{NP}^* satisfies $C_1^D = \min\{Y_1^D + T_{NP}^*, (Y_1^D + Y_2^D + T_{NP}^*)/2\}$ (liquidity-bound when the smoothing target exceeds $Y_1^D + T_{NP}^*$; smoothing otherwise). As $T_{NP}^* \rightarrow \infty$, both expressions diverge to infinity, so $C_1^D \rightarrow \infty$. In either case, $C_1^D \rightarrow \infty$.

Channel (i): $C_2^D \rightarrow \infty$ via the transfer-cash channel ($\phi < 1$). For $\phi < 1$, the daughter's period-1 cash is $Y_1^D + (1 - \phi)T_{NP}^*$, of which the cash share $(1 - \phi)T_{NP}^* \rightarrow \infty$. Her optimum allocates the cash between period-1 consumption (above the in-kind floor) and period-2 savings. We show $C_2^D \rightarrow \infty$ along any sequence with $Y^P \rightarrow \infty$, by binding regime at the daughter's optimum:

- *Liquidity-bound.* If the liquidity cap binds, then $C_2^D = Y_2^D$. Liquidity binding requires the smoothing target to exceed the cap, i.e., $(Y_1^D + Y_2^D + T_{NP}^*)/2 > Y_1^D + T_{NP}^*$, which simplifies to $Y_2^D > Y_1^D + T_{NP}^*$. So $Y_2^D \rightarrow \infty$ along this subsequence, hence $C_2^D = Y_2^D \rightarrow \infty$.
- *Smoothing.* If smoothing holds, then $C_2^D = (Y_1^D + Y_2^D + T_{NP}^*)/2 \rightarrow \infty$ because $T_{NP}^* \rightarrow \infty$.
- *In-kind-bound.* If the in-kind floor binds, then $C_1^D = \phi T_{NP}^*$ and the daughter saves the residual cash to period 2: $C_2^D = Y_1^D + Y_2^D + (1 - \phi)T_{NP}^* \rightarrow \infty$ because $(1 - \phi)T_{NP}^* \rightarrow \infty$.

The three regimes exhaust the possibilities, so $C_2^D \rightarrow \infty$ in every case, regardless of the path of Y_2^D . (For abortion, subtract p from C_2^D ; the limit is unchanged.)

Channel (ii): $C_2^D \rightarrow \infty$ via the own-income channel ($Y_2^D \rightarrow \infty$). Write $q \equiv (1 - S)(1 - k)pA \geq 0$ for the actual out-of-pocket payment out of the daughter's period-1 cash (zero under $S = 1$ or parental coverage $k = 1$, equal to pA in the hide-self-fund case under $S = 0$). By the liquidity constraint $C_1^D + q \leq Y_1^D + T_{NP}^*$ (equation (14)), the daughter's period-1 expenditure is bounded by her period-1 resources. Hence the daughter cannot borrow from period 2 to period 1: any period-2 income Y_2^D remains in period 2. Formally, the lifetime budget plus the liquidity constraint imply $C_2^D = Y_1^D + Y_2^D + T_{NP}^* - C_1^D - q \geq Y_2^D$ (using $C_1^D \leq Y_1^D + T_{NP}^* - q$). So $C_2^D \geq Y_2^D \rightarrow \infty$ as $Y_2^D \rightarrow \infty$, regardless of ϕ and regardless of the subsidy regime or coverage choice. $\square$

The proof of Proposition 3 reduces to showing convergence of the Layer-3 thresholds. By Lemma 6,

$T_{\text{NP}}^* \rightarrow \infty$ in the joint limit $\min(Y^P, Y_2^D) \rightarrow \infty$, and the daughter's consumption components diverge as well (under Channel i or ii). We use these inputs to establish:

Auxiliary: parental consumption diverges. By the parental KKT condition $u'(Y^P - T_{\text{NP}}^*) = \alpha^P s_H$ for some $s_H \in \partial v(T_{\text{NP}}^*)$. To bound s_H , we list the possible supergradients of v by regime, using the same regime case-analysis as the proof of Lemma 6. Depending on which constraint binds at the daughter's optimum, the supergradient takes one of three forms:

- *Liquidity-bound regime:* $v'(T) = u'(Y_1^D + T)$.
- *Smoothing regime:* $v'(T) = u'((Y_1^D + Y_2^D + T)/2)$.
- *In-kind-bound regime:* $v'(T) = \phi u'(\phi T) + (1 - \phi) u'(Y_1^D + Y_2^D + (1 - \phi)T)$.

At regime-switch points the superdifferential is a convex combination of the adjacent expressions. In every case the supergradient is a convex combination of u' evaluated at the daughter's optimal consumption components, with coefficients in $[0, 1]$.

By Lemma 6, $T_{\text{NP}}^* \rightarrow \infty$. Under the joint limit considered in Proposition 3, each consumption argument appearing in the relevant regime's formula diverges (Channels i and ii of Lemma 6), so each $u'(\cdot)$ vanishes by Assumption 2. Hence $s_H \rightarrow 0$, so $u'(Y^P - T_{\text{NP}}^*) \rightarrow 0$, which by strict monotonicity of u' and Assumption 2 forces $Y^P - T_{\text{NP}}^* \rightarrow \infty$ — parental period-1 consumption diverges.

Convergence of $\bar{\xi}^P \rightarrow 0$ (always). Since the auxiliary argument gives $Y^P - T_{\text{NP}}^* \rightarrow \infty$, parental coverage is eventually strictly feasible ($Y^P - T_{\text{NP}}^* > p$ for all sufficiently large Y^P), so the finite-threshold closed form of Lemma 2 applies. By Lemma 2, $\bar{\xi}^P = u(Y^P - T_{\text{NP}}^* - p) - u(Y^P - T_{\text{NP}}^*)$. By the mean-value theorem and the bound $|(Y^P - T_{\text{NP}}^* - p) - (Y^P - T_{\text{NP}}^*)| = p$,

$$|\bar{\xi}^P| = |u(Y^P - T_{\text{NP}}^* - p) - u(Y^P - T_{\text{NP}}^*)| \leq p \cdot u'(C_{\min}^P) \rightarrow 0,$$

where $C_{\min}^P = Y^P - T_{\text{NP}}^* - p \rightarrow \infty$ by the auxiliary above, and $u'(C_{\min}^P) \rightarrow 0$ by Assumption 2. Hence $\bar{\xi}^P \rightarrow 0$.

Convergence of ξ_3^D (case split). By the closed-form characterization of the Layer-3 self-funding threshold (Appendix E.4; derived in the proof of Lemma 1, part (c), at Appendix F.5), $\xi_3^D = v_p(T_{\text{NP}}^*) - v(T_{\text{NP}}^*)$ when self-funding is feasible at T_{NP}^* in the sense of the operational feasibility convention (Appendix E.4), and $-\infty$ otherwise.

Case A: $\phi < 1$ or $Y_1^D \geq p$. Self-funding is eventually feasible: under $\phi < 1$, $(1 - \phi)T_{\text{NP}}^* \rightarrow \infty$ (Lemma 6) makes the funding constraint slack regardless of Y_1^D ; under $Y_1^D \geq p$ (with $\phi = 1$ since otherwise the previous condition applies), the in-kind floor $\phi T_{\text{NP}}^* = T_{\text{NP}}^* > 0$ keeps C_1^D strictly positive

at the boundary, and the funding constraint $Y_1^D + (1 - \phi)T_{NP}^* = Y_1^D \geq p$ holds at every $T_{NP}^* \geq 0$. The difference $v_p(T_{NP}^*) - v(T_{NP}^*)$ equals the utility loss of paying p from an optimum with diverging consumption. The bound $|v(T_{NP}^*) - v_p(T_{NP}^*)| \rightarrow 0$ requires showing that paying p is feasible from the no-payment optimum with $o(1)$ utility loss. Let (C_1^0, C_2^0) be the optimizer for $v(T_{NP}^*)$, with both components diverging by Lemma 6 (eventually exceeding p). Construct a feasible allocation for $v_p(T_{NP}^*)$ by regime:

- *Liquidity-bound at the no-payment optimum* ($C_1^0 = Y_1^D + T_{NP}^*$). The post-payment liquidity cap is $C_1 \leq Y_1^D + T_{NP}^* - p$, so use $(C_1^0 - p, C_2^0)$. This satisfies the in-kind floor whenever $C_1^0 - p \geq \phi T_{NP}^*$, i.e., $Y_1^D + (1 - \phi)T_{NP}^* \geq p$, which is exactly the self-funding feasibility condition assumed in Case A.
- *In-kind-bound at the no-payment optimum* ($C_1^0 = \phi T_{NP}^*$). Self-funding feasibility gives $Y_1^D + (1 - \phi)T_{NP}^* \geq p$, so the unchanged $C_1^0 = \phi T_{NP}^*$ remains feasible after payment: the in-kind floor holds with equality, and the post-payment liquidity cap $C_1 \leq Y_1^D + T_{NP}^* - p$ follows from the feasibility inequality rearranged as $\phi T_{NP}^* \leq Y_1^D + T_{NP}^* - p$. Use $(C_1^0, C_2^0 - p)$.
- *Smoothing/interior at the no-payment optimum* ($C_1^0 = C_2^0 = (Y_1^D + Y_2^D + T_{NP}^*)/2$). Three sub-configurations by which post-payment constraint binds:
 - If the post-payment liquidity cap is slack and the post-payment smoothed allocation satisfies the in-kind floor ($\phi T_{NP}^* \leq (Y_1^D + Y_2^D + T_{NP}^* - p)/2$), use $C_1 = C_2 = (Y_1^D + Y_2^D + T_{NP}^* - p)/2$.
 - If the post-payment smoothed allocation violates the in-kind floor (i.e., $(Y_1^D + Y_2^D + T_{NP}^* - p)/2 < \phi T_{NP}^* \leq (Y_1^D + Y_2^D + T_{NP}^*)/2$), use $C_1 = \phi T_{NP}^*$ and $C_2 = Y_1^D + Y_2^D + T_{NP}^* - p - \phi T_{NP}^*$. The in-kind floor holds with equality; the post-payment liquidity cap $C_1 \leq Y_1^D + T_{NP}^* - p$ follows from self-funding feasibility $Y_1^D + (1 - \phi)T_{NP}^* \geq p$. Because the floor ϕT_{NP}^* lies between the post-payment and no-payment smoothing targets, the first component differs from C_1^0 by at most $p/2$, and the second component differs from C_2^0 by at most p .
 - If the post-payment liquidity cap binds, reduce C_1 to exactly satisfy the cap and put the residual in C_2 ; the adjustment is bounded by p in each component.

Define $m_n \equiv \min\{C_1^{0,n}, C_2^{0,n}\}$ for the no-payment optimizer along the sequence. By Lemma 6, $m_n \rightarrow \infty$. In each of the constructed post-payment allocations above, every component is at least $m_n - p$ (the deviation is at most p in each component, as established in the three subcases). By the mean-value theorem applied to u on consumption levels at least $m_n - p$, the utility loss in each component is at most $p \cdot u'(m_n - p)$, so the total loss satisfies $|v(T_{NP}^*) - v_p(T_{NP}^*)| \leq 2p \cdot u'(m_n - p)$, which tends to 0 by Assumption 2. Hence $\xi_3^D \rightarrow 0$.

Case B: $\phi = 1$ AND $Y_1^D < p$. The daughter's hide-branch period-1 cash equals $Y_1^D + 0 \cdot T_{NP}^* = Y_1^D < p$ for every Y^P : the funding constraint $Y_1^D + (1 - \phi)T_{NP}^* \geq p$ is never satisfied, so self-funding is never feasible under the strict-positive-consumption convention. Hence $\xi_3^D = -\infty$ for every Y^P and remains $-\infty$ in the limit.

We collect: in the joint limit $\min(Y^P, Y_2^D) \rightarrow \infty$,

$$\bar{\xi}_3^D = 0, \quad \bar{\xi}^P \rightarrow 0, \quad \underline{\xi}_3^D \rightarrow \underline{\xi}_3^{D,\infty} \in \{0, -\infty\}.$$

The limiting always-taker threshold equals 0 in Case A and $-\infty$ in Case B. The dichotomy is summarized below.

Case	$\underline{\xi}_3^{D,\infty}$	$\bar{\xi}^{P,\infty}$	Limiting complier set $\mathcal{C}^\infty$
A: $\phi < 1$ or $Y_1^D \geq p$	0	0	empty (up to boundary)
B: $\phi = 1$ and $Y_1^D < p$	$-\infty$	0	$\{\xi^P > 0, \xi^D \leq 0\}$

In Case A, $\Delta A^\infty = 0$; in Case B, ΔA^∞ is the generically positive mass of the half-plane above.

Proof of Proposition 3(a). The population subsidy effect is $\Delta A(h, Y^P, Y_2^D, Y_1^D, \phi) = \Pr_{f_h}(\mathcal{C}(Y^P, Y_2^D))$, where $\mathcal{C}(Y^P, Y_2^D)$ is the complier set at the given parental and daughter incomes (eq. 5). $\mathcal{C}$ is the mismatch rectangle, with boundaries given by the three Layer-3 thresholds.

By the threshold-convergence statement above, each threshold converges to a limit (possibly $-\infty$). Define the limiting complier set $\mathcal{C}^\infty$ by these limiting thresholds. The form of $\mathcal{C}^\infty$ depends on the case:

Case A ($\phi < 1$ or $Y_1^D \geq p$): all three limiting thresholds equal 0. The limiting mismatch interval $(\underline{\xi}_3^{D,\infty}, 0] = (0, 0] = \emptyset$. $\mathcal{C}^\infty$ reduces to a measure-zero boundary, so $\Delta A^\infty = \Pr_{f_h}(\mathcal{C}^\infty) = 0$.

Case B ($\phi = 1$ AND $Y_1^D < p$): $\underline{\xi}_3^{D,\infty} = -\infty$, while $\bar{\xi}^{P,\infty} = 0$. The limiting mismatch interval becomes $(\underline{\xi}_3^{D,\infty}, 0] = (-\infty, 0]$ (full lower half-plane in ξ^D), giving mismatch region $\{\xi^P > \bar{\xi}^{P,\infty} = 0, \xi^D \in (-\infty, 0]\}$. So $\mathcal{C}^\infty = \{\xi^P > 0, \xi^D \leq 0\}$, and $\Delta A^\infty = \Pr_{f_h}(\xi^P > 0, \xi^D \leq 0)$, generically positive.

Convergence in both cases. The indicator $1[(\xi^P, \xi^D) \in \mathcal{C}(Y^P, Y_2^D)]$ converges pointwise to $1[(\xi^P, \xi^D) \in \mathcal{C}^\infty]$ outside the limiting threshold loci. Because f_h is absolutely continuous with respect to two-dimensional Lebesgue measure (Assumption 3, imposed community by community on f_h as noted in Appendix H), the limiting boundary sets $\{\xi^D = 0\}$, $\{\xi^P = 0\}$, and any finite moving-threshold limits have f_h -measure zero. Therefore the weak/strict inequality convention at the limiting boundary does not affect the probability limit. Since the indicator is bounded by 1 and f_h is a density, dominated convergence gives $\Pr_{f_h}(\mathcal{C}(Y^P, Y_2^D)) \rightarrow \Pr_{f_h}(\mathcal{C}^\infty)$, i.e., $\Delta A \rightarrow \Delta A^\infty(h, Y_1^D, \phi)$.

Saturation. For any $b > 0$, $\Delta A(h, bY^P, bY_2^D, Y_1^D, \phi) - \Delta A(h, Y^P, Y_2^D, Y_1^D, \phi) \rightarrow 0$ as $\min(Y^P, Y_2^D) \rightarrow \infty$, because both terms converge to the same limit ΔA^∞ . The case-B persistence (where $\Delta A^\infty > 0$) does not violate saturation: the limit is invariant to proportional scaling of adult incomes. $\square$

Proof of Proposition 3(b). Fix $h, \phi \in [0, 1]$, and assume $Y_1^D \geq p$. We use the joint limit $\min(Y^P, Y_2^D) \rightarrow \infty$. This is Case A of the threshold-convergence statement (the disjunct $Y_1^D \geq p$).

Feasibility of self-funding. Since $Y_1^D \geq p$, the funding constraint $Y_1^D + (1 - \phi)T_{\text{NP}}^* \geq p$ holds at every $T_{\text{NP}}^* \geq 0$. At the boundary $Y_1^D = p$ with $\phi = 1$, the in-kind floor $\phi T_{\text{NP}}^* = T_{\text{NP}}^* > 0$ keeps C_1^D strictly positive (Remark on the operational feasibility convention, Appendix E.4); at $\phi < 1$ the cash slack $(1 - \phi)T_{\text{NP}}^* > 0$ already makes the funding constraint hold strictly. Hence hidden self-funding is feasible throughout (in particular, eventually feasible in the joint limit).

By the threshold-convergence statement (Case A), all three thresholds converge to 0: $\xi_3^D \rightarrow 0$, $\bar{\xi}^P \rightarrow 0$, $\bar{\xi}_3^D = 0$. The limiting complier set (mismatch region with limiting thresholds) is the measure-zero curve $\{\xi^P > 0, \xi^D = 0\}$. By the proof of Proposition 3(a) (dominated convergence + atomless f_h), $\Delta A^\infty(h, Y_1^D, \phi) = 0$ for any $\phi \in [0, 1]$. $\square$

Proof of Proposition 3(c). Fix h , Y_1^D , and $\phi \in [0, 1)$. We use $Y^P \rightarrow \infty$ (the result holds under this weaker condition than the joint limit, since the transfer-cash channel does not require $Y_2^D \rightarrow \infty$).

Threshold convergence under $Y^P \rightarrow \infty$ alone. The threshold-convergence statement above was framed for the joint limit $\min(Y^P, Y_2^D) \rightarrow \infty$; we restate the chain under $\phi < 1$ and $Y^P \rightarrow \infty$ alone, uniformly in Y_2^D . By Lemma 6, $T_{\text{NP}}^* \rightarrow \infty$ as $Y^P \rightarrow \infty$ uniformly in Y_2^D . By Channel (i) of Lemma 6, $C_1^D \rightarrow \infty$ and $C_2^D \rightarrow \infty$ as $Y^P \rightarrow \infty$, both independent of Y_2^D . Hence every consumption argument appearing in the regime-specific formulas for $v'(T_{\text{NP}}^*)$ diverges, so each $u'(\cdot)$ in the supergradient vanishes by Assumption 2, and $s_H \rightarrow 0$. The parental KKT condition $u'(Y^P - T_{\text{NP}}^*) = \alpha^P s_H$ then forces $u'(Y^P - T_{\text{NP}}^*) \rightarrow 0$, which by strict monotonicity of u' and Assumption 2 gives $Y^P - T_{\text{NP}}^* \rightarrow \infty$. Parental coverage is therefore eventually strictly feasible, and the mean-value bound from before yields $\bar{\xi}^P \rightarrow 0$. Self-funding is also eventually feasible (under $\phi < 1$, $(1 - \phi)T_{\text{NP}}^* \rightarrow \infty$ makes the funding constraint slack regardless of Y_1^D), and the regime-by-regime post-payment construction from Case A above yields $\xi_3^D \rightarrow 0$. The never-taker threshold $\bar{\xi}_3^D = 0$ identically.

All three thresholds therefore converge to 0 under $Y^P \rightarrow \infty$ alone. As in the proof of Proposition 3(b), the complier set in the limit reduces to a measure-zero curve; by dominated convergence, $\Delta A^\infty(h, Y_1^D, \phi) = 0$ for any $\phi < 1$ and any Y_1^D . Since this conclusion uses only $Y^P \rightarrow \infty$, it also implies the joint-limit version (the joint limit $\min(Y^P, Y_2^D) \rightarrow \infty$ includes $Y^P \rightarrow \infty$). $\square$

K Welfare: Derivations and Proofs

This appendix proves Proposition 4 and Corollary 1, states and proves the concealment result referenced in Section 6.3 (Proposition 6), collects the full welfare-incidence table, and records two supporting remarks. All objects are the Layer-3 equilibrium objects of Appendix E.4; we abbreviate $T \equiv T_{\text{NP}}^*$, $v \equiv v(T)$, $v_p \equiv v_p(T)$ (equations (23)–(24)), and recall

$$\bar{\xi}^P = u(Y^P - T - p) - u(Y^P - T) < 0, \quad \xi_3^D = v_p - v < 0.$$

We assume the interior configuration in which the regions of Figure 3, Panel (c), are non-empty: parental coverage is feasible ($Y^P - T > p$, so $\bar{\xi}^P$ is finite) and hidden self-funding is feasible ($Y_1^D + (1 - \phi)T \geq p$, so ξ_3^D is finite). When either feasibility fails, the corresponding region is empty (Lemma 2; Appendix E.4, “Operational form of the feasibility convention”) and the remaining statements are unaffected.

Welfare objects. As in Section 6.1, family welfare is $W = V^D + w^P$, where V^D is the daughter’s utility (3) and w^P is the parents’ utility net of the altruism term: $w^P = u(C^P) - \xi^P AR$ under the *knowledge* benchmark, exactly as in (4), and $w^P = u(C^P) - \xi^P A$ under the *action* benchmark. Parental consumption is the residual $C^P = Y^P - T - (1 - S)p k A$. The daughter’s realized consumption value equals v whenever she pays nothing out of pocket (no abortion, a covered abortion, or any abortion under $S = 1$) and v_p when she self-finances the procedure ($A = 1$, $k = 0$, $S = 0$). We write ΔW for the change in W from the move $S : 0 \rightarrow 1$, holding T fixed across regimes—which is immediate, since T solves the non-pregnancy problem (19) and does not involve S . The two benchmarks induce identical equilibrium behavior, as asserted in Section 6.1: the allowance is chosen on the non-pregnancy branch (Remark E.1); the parents’ only remaining decision—coverage—arises on the TELL branch, where $AR = A$, so the parents’ objective is identical under the two benchmarks; and the daughter’s payoff does not involve ξ^P . Every equilibrium object below is therefore common to both benchmarks.

The canonical coverage case. Off the complier-relevant region, parental coverage of a daughter who would abort anyway is governed by the primitive condition (22) (Lemma 3), which is independent of ξ^P . Throughout this appendix we maintain the *canonical case* in which condition (22) fails: parents do not pay for a procedure their daughter can and would fund herself. The knife-edge opposite case is treated in Remark K.2; it changes no sign in any result below.

K.1 Equilibrium outcomes under each regime

The following lemma reads the equilibrium outcomes off the characterizations established in Appendices E–F and is the only input to the proofs.

Lemma 7 (Regime-by-regime equilibrium outcomes). *Fix a family (ξ^D, ξ^P) , hold T fixed across regimes, and maintain the canonical coverage case.*

- (i) Subsidized regime ($S = 1$). *The procedure is free, so the daughter's consumption value is v irrespective of A , and she aborts iff $\xi^D \leq 0$; parents pay nothing, $C^P = Y^P - T$. Disclosure is payoff-equivalent for the daughter and is selected by the tie-breaking rule of Appendix E.4: $R = 1$ iff $\xi^P \leq 0$.*
- (ii) Unsubsidized regime ($S = 0$).
 - (a) Never-taker ($\xi^D > 0$): $A = 0$; the daughter consumes v and $C^P = Y^P - T$; disclosure is payoff-irrelevant ($AR = 0$ either way).
 - (b) Covered abortion ($\xi_3^D < \xi^D \leq 0$, $\xi^P \leq \bar{\xi}^P$): the daughter discloses, parents cover, and she aborts, $A = R = k = 1$; she consumes v and $C^P = Y^P - T - p$ (Lemmas 2 and 3).
 - (c) Self-funded abortion ($\xi^D \leq \xi_3^D$): the daughter aborts and pays for the procedure herself, $A = 1$, $k = 0$; she consumes v_p and $C^P = Y^P - T$. Hiding and disclosing are payoff-equivalent for her (parents refuse coverage either way, by the canonical case), and the tie-breaking rule selects $R = 1$ iff $\xi^P \leq 0$.
 - (d) Complier ($\xi_3^D < \xi^D \leq 0$, $\xi^P > \bar{\xi}^P$): no abortion, $A = 0$; her disclosure is $R = 1[\xi^P \leq 0]$, payoff-irrelevant here since $A = 0$; she consumes v and $C^P = Y^P - T$ (Lemma 3).

Proof. Part (i). With $S = 1$ the out-of-pocket term $(1 - S)(1 - k)pA$ vanishes, so the daughter's budget set is independent of A and her consumption value is v in every branch; she maximizes $-\xi^D A$, so $A = 1$ iff $\xi^D \leq 0$ (Lemma 1(a) and its proof). Parents pay $(1 - S)p k A R = 0$ in every branch, so $C^P = Y^P - T$. Because her payoff does not depend on R , the disclosure selection is the maintained tie-breaking rule (Appendix E.4): $R = 1$ iff $\xi^P \leq 0$.

Part (ii). (a) For $\xi^D > 0$ the daughter never aborts (Lemma 1(a)), consumes v , and no payment arises. (b) On the complier-relevant region with willing parents, Lemma 3 gives disclosure and Lemma 2 gives coverage; the daughter aborts (the procedure is free to her and $\xi^D \leq 0$), consuming v , while parents pay p , so $C^P = Y^P - T - p$. (c) For $\xi^D \leq \xi_3^D$ the daughter aborts under $k = 0$ by the definition of ξ_3^D ($v_p - \xi^D \geq v$). Upon disclosure, parents compare covering with refusing while anticipating that she aborts either way; by the canonical case (condition (22) fails) they refuse (Lemma 3). Her payoff is therefore $v_p - \xi^D$ whether she hides or discloses, and the tie-breaking rule selects $R = 1$

iff $\xi^P \leq 0$. (d) This is Lemma 3: the daughter cannot profitably self-fund ($\xi^D > \underline{\xi}_3^D$), parents would refuse ($\xi^P > \bar{\xi}^P$), so she does not abort, consuming v ; her disclosure $R = \mathbf{1}[\xi^P \leq 0]$ is payoff-irrelevant since $A = 0$. $\square$

K.2 Proof of Proposition 4

We compute ΔW as the $S = 1$ value of W minus its $S = 0$ value, using Lemma 7; under $S = 1$ every family has daughter consumption value v and $C^P = Y^P - T$. The attitude terms can change only if A changes (action benchmark) or AR changes (knowledge benchmark).

Never-takers (part 4(a)). By Lemma 7(i),(ii.a), $A = 0$ under both regimes, the daughter consumes v , and $C^P = Y^P - T$ under both. All consumption terms are unchanged and $A = AR = 0$ in both regimes, so $\Delta W = 0$ under either benchmark.

Always-takers (part 4(b)). Part 4(b) asserts only the sign; the proof also derives the magnitudes collected in Table K.1. *Self-funded* ($\xi^D \leq \underline{\xi}_3^D$). By Lemma 7(i),(ii.c), $A = 1$ under both regimes and the disclosure selection is the same in both ($R = 1$ iff $\xi^P \leq 0$), so both $-\xi^P A$ and $-\xi^P AR$ are unchanged: the attitude terms cancel under either benchmark. The daughter's consumption value rises from v_p to v and C^P is unchanged, so

$$\Delta W = v - v_p > 0$$

under either benchmark (strict because $p > 0$ and u is strictly increasing)—the utility value to the daughter, who pays for the procedure under $S = 0$, of no longer paying p .

Covered ($\underline{\xi}_3^D < \xi^D \leq 0$, $\xi^P \leq \bar{\xi}^P$). By Lemma 7(i),(ii.b), $A = R = 1$ under both regimes, so the attitude terms again cancel under either benchmark. The daughter consumes v under both regimes, while parental consumption rises from $Y^P - T - p$ to $Y^P - T$, so

$$\Delta W = u(Y^P - T) - u(Y^P - T - p) = -\bar{\xi}^P > 0$$

under either benchmark—the utility value to the parents, who pay under $S = 0$, of no longer paying p . In the knife-edge coverage case, $\Delta W > 0$ continues to hold on the self-fund region under either benchmark (Remark K.2), so part 4(b) holds without qualification.

Compliers (part 4(c)). By Lemma 7(ii.d), the $S = 0$ outcome is $A = 0$ (so $AR = 0$), daughter value v , $C^P = Y^P - T$; by Lemma 7(i), the $S = 1$ outcome is a free abortion with $A = 1$ and $R = \mathbf{1}[\xi^P \leq 0]$, daughter value $v - \xi^D$, $C^P = Y^P - T$. All consumption terms are unchanged, so ΔW consists of attitude terms only. The daughter's term contributes $-\xi^D \cdot \Delta A = -\xi^D$. The parents' term differs by

benchmark. Under the action benchmark $-\xi^P A$ moves from 0 to $-\xi^P$, contributing $-\xi^P$, so $\Delta W = -\xi^D - \xi^P$. Under the knowledge benchmark $-\xi^P AR$ moves from 0 to $-\xi^P \cdot \mathbf{1}[\xi^P \leq 0] = -\min\{\xi^P, 0\}$, so $\Delta W = -\xi^D - \min\{\xi^P, 0\}$, which is ≥ 0 since $-\xi^D \geq 0$ and $-\min\{\xi^P, 0\} \geq 0$. Equivalently, the two benchmarks coincide on the approving sub-region ($\bar{\xi}^P < \xi^P \leq 0$, where the abortion is disclosed) and differ on the opposed sub-region ($\xi^P > 0$, where it is concealed): there the knowledge value is $-\xi^D$ and the action value is $-\xi^D - \xi^P$. $\square$

K.3 Proof of Corollary 1

By Lemma 7(i),(ii.d), on the open complier region the abortion outcome is the same for every family—no abortion under $S = 0$, an abortion under $S = 1$. Any two attitude pairs in the region therefore generate the same abortion behavior: no data on abortion outcomes can distinguish them. Membership in the region bounds attitudes only by $\xi^D \in (\underline{\xi}_3^D, 0]$ and $\xi^P > \bar{\xi}^P$ (Appendix F.6), and the second bound is one-sided: ξ^P is unbounded above.

Under the action benchmark, $\Delta W = -\xi^D - \xi^P$ (Proposition 4(c)) is positive iff $\xi^P < -\xi^D$ and negative iff $\xi^P > -\xi^D$. The boundary $\xi^P = -\xi^D$ intersects the interior of the complier rectangle: for $\xi^D \in (\underline{\xi}_3^D, 0]$ we have $-\xi^D \in [0, -\underline{\xi}_3^D)$, while ξ^P ranges over $(\bar{\xi}^P, \infty) \ni -\xi^D$ since $-\xi^D \geq 0 > \bar{\xi}^P$. Because families on the two sides of the boundary generate identical abortion behavior, neither an individual complier's sign nor the complier mass on each side is identified from abortion data: for any candidate joint attitude distribution f , any distribution $\tilde{f}$ that agrees with f on the masses of the partition's regions—however it allocates complier mass across the boundary—generates an identical distribution of abortion outcomes. $\square$

K.4 The subsidy's abortions are concealed

Section 6.3 asserts that the subsidy's abortions are concealed precisely from the parents who oppose them. This subsection states and proves that result. Let $H(S \mid h)$ denote the share of pregnancies in community h that end in a *hidden* abortion ($A = 1, R = 0$) under regime S , with probabilities taken under the attitude density f_h (Assumption 3).

Proposition 6 (The subsidy's abortions are concealed from opposed parents). *Consider the move $S : 0 \rightarrow 1$, in the canonical coverage case.*

- (a) *The subsidy induces abortion exactly among compliers. Among them, the abortion is concealed ($R = 0$) when parents oppose it ($\xi^P > 0$) and disclosed ($R = 1$) when they approve but would not pay ($\bar{\xi}^P < \xi^P \leq 0$).*

(b) Hidden abortions rise by the mass of opposed compliers, which is at most the rise in total abortions.⁵⁴

$$H(1 | h) - H(0 | h) = \Pr(\xi_3^D < \xi^D \leq 0, \xi^P > 0 | h) \leq A(1 | h) - A(0 | h),$$

with equality iff the approving-complier band $\{\xi_3^D < \xi^D \leq 0, \bar{\xi}^P < \xi^P \leq 0\}$ has zero mass.

Proof. Part 6(a). The families whose abortion outcome the subsidy changes are exactly the compliers (Appendix F.6), defined by $\xi^D \in (\xi_3^D, 0]$ and $\xi^P > \bar{\xi}^P$. By Lemma 7(i), each complier's abortion under $S = 1$ carries $R = 1[\xi^P \leq 0]$: concealed ($R = 0$) iff $\xi^P > 0$ and disclosed ($R = 1$) iff $\bar{\xi}^P < \xi^P \leq 0$.

Part 6(b). A hidden abortion is a pair $(A, R) = (1, 0)$, i.e. an abortion by a family with $\xi^P > 0$. Under $S = 1$ every daughter with $\xi^D \leq 0$ aborts, so the $S = 1$ hidden set is $\{\xi^D \leq 0, \xi^P > 0\}$. Under $S = 0$ the only abortions are by self-funders ($\xi^D \leq \xi_3^D$) and covered always-takers ($\xi^P \leq \bar{\xi}^P$, hence not opposed), so the $S = 0$ hidden set is $\{\xi^D \leq \xi_3^D, \xi^P > 0\}$. Subtracting,

$$H(1 | h) - H(0 | h) = \Pr(\xi_3^D < \xi^D \leq 0, \xi^P > 0 | h).$$

This is a subset of the complier region $\{\xi_3^D < \xi^D \leq 0, \xi^P > \bar{\xi}^P\}$, whose mass is $A(1 | h) - A(0 | h)$ (Section 5); since $\bar{\xi}^P < 0$, the inclusion is strict exactly when the approving-complier band $\{\bar{\xi}^P < \xi^P \leq 0\}$ carries positive mass. $\square$

K.5 The incidence table

Table K.1 collects the results of Proposition 4 region by region, under each benchmark, together with the change in the public outlay ΔG : under $S = 1$ the government finances one procedure, at cost p , in every family that aborts ($\xi^D \leq 0$), and none for never-takers.

K.6 Supporting remarks

Remark K.1 (Generation weights and double-counting: sign-equivalence). Consider the weighted objective $\mathcal{W}_\lambda = \lambda_D V^D + \lambda_P w^P$ with arbitrary weights $\lambda_D, \lambda_P > 0$ (W is the case $\lambda_D = \lambda_P = 1$). By the proof of Proposition 4, in each region $\Delta \mathcal{W}_\lambda$ is: never-takers, 0; self-funding always-takers, $\lambda_D(v - v_p) > 0$; covered always-takers, $\lambda_P(-\bar{\xi}^P) > 0$; compliers, $-\lambda_D \xi^D - \lambda_P \min\{\xi^P, 0\} \geq 0$ under the knowledge benchmark and $-\lambda_D \xi^D - \lambda_P \xi^P$ under the action benchmark. Every sign statement in Proposition 4 and Corollary 1 is therefore invariant to the weights; under the action benchmark the sign boundary rotates to $\xi^P = (\lambda_D/\lambda_P)(-\xi^D)$, but it still emanates from $(\xi^D, \xi^P) = (0, 0)$, still crosses the interior of the complier rectangle, and its location within the rectangle remains unidentified for

⁵⁴In the opposite knife-edge coverage case the subsidy in addition converts opposed self-funders from disclosed to hidden (Remark K.2), reinforcing part 6(a).

Table K.1: Welfare incidence of the move $S : 0 \rightarrow 1$, by region

Region	Pays ($S=0$)	ΔW (knowledge)	ΔW (action)	ΔG
Never-taker ($\xi^D > 0$)	none	0	0	0
Always-taker, self-funded ^a	daughter	$v - v_p$	$v - v_p$	$-p$
Always-taker, covered	parents	$-\bar{\xi}^P$	$-\bar{\xi}^P$	$-p$
Complier	none	$-\xi^D - \min\{\xi^P, 0\}$	$-\xi^D - \xi^P$	$-p$

^aEntries assume the canonical coverage case (condition (22) fails), in which parents do not pay for a procedure their daughter would obtain anyway. The knife-edge opposite case shifts the self-funded row's gain to the parents and leaves every sign unchanged (Remark K.2). All knowledge-benchmark entries are ≥ 0 ; the two benchmarks agree except on the complier row, and there only for opposed parents ($\xi^P > 0$)—exactly the concealed abortions—where the knowledge value is $-\xi^D \geq 0$ and the action value $-\xi^D - \xi^P$ is unsigned (Corollary 1). ΔG is the change in public funds (an outlay of p per procedure financed).

the same reason as before. The same equivalence holds for the transparent sum $V^D + V^P$, which adds $\alpha^P \Delta[u(C_1^D) + u(C_2^D)]$ to ΔW : this addition is $\alpha^P(v - v_p) > 0$ for self-funding always-takers and zero in every other region (the daughter's consumption is unchanged there), so it can only reinforce a weakly positive sign and never affects the complier region, where the contested comparison lives.

Remark K.2 (The knife-edge coverage case). Suppose condition (22) holds: parents are altruistic enough to pay for a procedure their daughter can fund herself. Then under $S = 0$ every self-funding always-taker discloses and is covered—disclosure is her strict best response, since coverage raises her consumption value from v_p to v (Lemma 3)—so her outcome is $(A, R, k) = (1, 1, 1)$, she consumes v , and $C^P = Y^P - T - p$. Under $S = 1$ she consumes v , $C^P = Y^P - T$, and $R = 1$ iff $\xi^P \leq 0$ by the tie-break. Recomputing ΔW on the self-fund region: the gain shifts to the parents, who no longer pay. For approving parents ($\xi^P \leq 0$), A and R are unchanged across regimes, so $\Delta W = -\bar{\xi}^P > 0$ under either benchmark. For opposed parents ($\xi^P > 0$), the abortion leaves the parents' view under the subsidy, $R : 1 \rightarrow 0$: under the action benchmark the attitude terms still cancel ($A = 1$ in both regimes) and $\Delta W = -\bar{\xi}^P > 0$; under the knowledge benchmark the parents' term changes by $+\xi^P$, so $\Delta W = -\bar{\xi}^P + \xi^P > 0$. This cell is the only place in the model where the two benchmarks assign different welfare changes outside the complier region. The sign in Proposition 4(b) is unchanged; what shifts are the magnitudes in the self-funded row of Table K.1: the gain moves to the parents and, under the knowledge benchmark, acquires the $+\xi^P$ term above. For Proposition 6, the conversions $R : 1 \rightarrow 0$ occur only for opposed parents ($\xi^P > 0$), adding opposed self-funders to the concealed set and so reinforcing part 6(a).

L Financing, Insurance, and the Fiscal Margin

The welfare analysis of Section 6 treats the subsidy’s fiscal cost as a separate object: Proposition 4 evaluates private family welfare, and Table K.1 reports the public outlay alongside it. This appendix closes the public budget explicitly and asks whether doing so changes anything. It does not: with the natural financing scheme, every conclusion of Section 6 survives up to terms that vanish with the pregnancy probability, and the financed subsidy acquires an additional, weight-free interpretation—it completes a missing credit market. We emphasize at the outset that this appendix is a *benchmark*, not a *load-bearing part of the analysis*: no result in Sections 4–6 relies on it.

L.1 The fiscal scale and the financing scheme

Under $S = 1$ the government pays p for every abortion performed (Table K.1). In a community h , the expected outlay per daughter is

$$\bar{\tau} \equiv \pi \cdot A(1 \mid h) \cdot p :$$

the probability of pregnancy, times the share of pregnancies that end in abortion when the procedure is free, times the price. Each factor is small; in particular π is on the order of 4.7% in the relevant population, so $\bar{\tau}$ is second order relative to the already-small price p .

We finance $\bar{\tau}$ with lump-sum taxes levied on every family, irrespective of pregnancy or abortion status: τ^P on parental income in period 1 and τ^D on the daughter’s income in period 2, with $\tau^P + \tau^D = \bar{\tau}$. The two bases are chosen deliberately. Both are *low marginal-utility* margins: parental resources are large relative to the procedure cost (the saturation logic of Proposition 3), and the daughter’s period-2 income exceeds her period-1 income by assumption. What the scheme deliberately does *not* tax is the daughter’s period-1 income—the constrained margin that the subsidy exists to relax. We do not take a stand on the split (τ^P, τ^D) ; every statement below holds for any split across these two bases.

Proposition 7 (Financing benchmark; non-load-bearing). *Fix the Layer-3 environment and let the financed subsidy be $S = 1$ with endowments $(Y_1^D, Y_2^D - \tau^D, Y^P - \tau^P)$, where $\tau^P + \tau^D = \bar{\tau} = \pi A(1 \mid h)p$. Compare it with the no-program baseline ($S = 0$, no taxes). Then:*

- (a) The never-taker margin is unchanged. Under the financed subsidy the daughter aborts iff $\xi^D \leq 0$: the taxes do not depend on A , so her budget set is independent of A exactly as in Lemma 1(a).
- (b) All other objects move continuously, and their deviations vanish with π . The allowance T_{NP}^* and the thresholds ξ_3^D and $\bar{\xi}^P$ computed in the taxed environment converge to their untaxed values as $\pi \rightarrow 0$ (equivalently $\bar{\tau} \rightarrow 0$).

(c) Welfare conclusions are preserved. For every family off the threshold loci, the welfare change ΔW from the no-program baseline to the financed subsidy converges to the value in Proposition 4 as $\pi \rightarrow 0$, under either benchmark.⁵⁵ In particular, every strict sign in Proposition 4 and Table K.1 is preserved for π small enough, and the non-identification result (Corollary 1) is unaffected: the taxes carry no information about ξ^P , which behavior still bounds only from below.

Proof. Part 7(a). Under the financed subsidy the out-of-pocket procedure cost is zero and the taxes enter the budget constraints as constants, independent of (A, R, k) . The daughter's budget set is therefore independent of A , her consumption value is common across $A \in \{0, 1\}$, and she maximizes $-\xi^D A$: $A = 1$ iff $\xi^D \leq 0$, exactly as in the proof of Lemma 1(a).

Part 7(b). The allowance T_{NP}^* is the unique maximizer of the strictly concave, continuous problem (19) (Lemma 4, Step 5); by the maximum theorem, applied to the compact-valued, continuous constraint correspondence indexed by the endowment vector, the maximizer is continuous in (Y^P, Y_2^D) on the region where the interiority condition holds—which it does for $\bar{\tau}$ small, since it holds strictly at $\bar{\tau} = 0$. The thresholds are continuous compositions: $\bar{\xi}^P = u(Y^P - \tau^P - T_{NP}^* - p) - u(Y^P - \tau^P - T_{NP}^*)$ and $\xi_3^D = v_p(T_{NP}^*) - v(T_{NP}^*)$, where v and v_p are continuous in the allowance and the endowments (equations (23)–(24); continuity by the maximum theorem, as established in Appendix F). Since $\bar{\tau} = \pi A(1 | h)p \leq \pi p \rightarrow 0$ as $\pi \rightarrow 0$, the taxed objects converge to the untaxed ones.

Part 7(c). Fix a family off the threshold loci of both the untaxed and the limiting partition. By part 7(b), for π small the family occupies the same region of the partition under the taxed regime as under the untaxed one, so its equilibrium outcome vector (A, R, k) in each regime is the one in Lemma 7. Its payoffs under the financed subsidy differ from the untaxed $S = 1$ payoffs only through the taxes and the induced allowance shift, all of which vanish continuously as $\pi \rightarrow 0$ (part 7(b) and continuity of u, v, v_p). Hence ΔW converges, region by region and under either benchmark, to the value in Proposition 4, and any strict inequality there holds for π small enough. On the complier region the taxes carry no attitude content, so under the action benchmark the comparison still turns on ξ^P relative to $-\xi^D$, up to vanishing consumption terms, and the non-identification argument of Corollary 1 applies unchanged. $\square$

⁵⁵The convergence is pointwise, not uniform. The taxes shift the coverage and self-fund thresholds by $O(\pi)$ (part 7(b); the never-taker threshold $\xi^D = 0$ does not move, by part 7(a)), so at each π a band of families of mass $O(\pi)$ is classified differently than in the untaxed model—for instance, a covered always-taker becomes a complier. On this band ΔW need not equal its Proposition 4 value, but it stays bounded, so the band's contribution to aggregate welfare is $O(\pi)$ and vanishes. Every fixed family off the untaxed loci eventually leaves the band, so its ΔW converges as stated, and the aggregate welfare conclusion is preserved as well.

L.2 Market completion and insurance

The financed subsidy has an interpretation that requires no stand on generation weights or on the benchmarks of Section 6.1. Consider the polar case $\tau^P = 0$, in which the program is financed entirely by the daughters' generation: every daughter pays $\bar{\tau}$ out of period-2 income, and every pregnant daughter who chooses an abortion receives a procedure worth p in period 1.

This package is precisely the contract that the missing credit market fails to deliver. The binding friction in Layer 2 is that a young woman cannot borrow against her future income during the two-month window in which the abortion decision must be made: the loan of p in period 1, repaid out of period-2 income, is the trade the borrowing constraint rules out. The financed subsidy supplies it—with the repayment pooled across the pregnancy shock, so that each daughter repays the actuarial value $\pi A(1 | h)p$ in every state rather than p in the pregnancy state. Financing the program within the daughters' generation is thus mutual insurance among daughters, funded from their own slack future income: it moves resources from a low marginal-utility margin (period-2 income, and only $\bar{\tau}$ of it) to the highest marginal-utility margin in the model (the liquidity-constrained period-1 pregnant state), relaxing a binding constraint at second-order cost. This is an efficiency gain in the usual incomplete-markets sense, and it is the fiscal-side counterpart of the paper's central positive finding that the subsidy operates through liquidity rather than price.

To the extent the tax instead falls on parents ($\tau^P > 0$), the program adds an intergenerational transfer: resources move from the unconstrained, saturated parental margin (Proposition 3) toward the constrained daughters. Under limited altruism ($\alpha^P < 1$), parents undervalue their daughters' felicity relative to a planner who weights the generations equally, so such a planner values this transfer positively; a planner using the family's own weights is indifferent at the margin. We do not pin the split (τ^P, τ^D) , and none of our conclusions requires doing so: both bases are low-marginal-utility margins, and the qualitative statements above hold for any mix.

L.3 Caveats

Four caveats bound what this appendix claims. First, the taxes are lump sum; with distortionary taxation, public funds carry a marginal cost above one, and the inframarginal outlay on always-takers—a pure transfer—would carry the usual deadweight consideration. Second, the marginal procedures themselves consume real resources of p per complier; this is small by the same magnitudes that make the price channel weak in Section 4, and it is dwarfed by the family-side stakes (ξ^D, ξ^P) that the welfare question actually turns on. Third, never-takers pay $\bar{\tau}$ and never use the program; their stake is the utility value of an $O(\pi)$ tax, second order by construction—and any concern that never-takers have systematically higher marginal utility (for instance, if child-rearing costs entered the budget con-

straint rather than ξ^D) is bounded by the same rare-pregnancy magnitude that justifies the allowance pre-commitment (Remark E.1). Fourth, nothing behavioral changes: by Proposition 7, the financed subsidy shifts thresholds only by amounts that vanish with π , so the positive analysis of Sections 4–5 is untouched.

M Parental Consent versus the Subsidy

Section 6.3 argues that the subsidy shifts decision-making authority within the family, inviting a comparison with the policy instrument that allocates such authority explicitly: parental-consent requirements. This appendix formalizes the comparison. The message is that the two instruments are related but not equivalent: they act on *disjoint* regions of the attitude space, because a consent requirement withholds permission while the subsidy provides money. Removing a consent requirement is therefore not equivalent to granting a subsidy, and vice versa (Figure M.1). Throughout, we maintain the canonical coverage case of Appendix K (condition (22) fails); Remark M.1 treats the knife-edge opposite case.

M.1 The consent regime

The model of Section 4 describes a world without a parental consent requirement: a daughter who can pay can obtain the procedure without involving her parents. A consent requirement removes exactly that route. Formally, we modify stage (iv) of the Layer-3 game (Appendix E.4) as follows: an abortion requires disclosure and parental authorization. In the HIDE branch, $A = 0$ is imposed. In the TELL branch, parents choose an authorization $\gamma \in \{0, 1\}$ together with coverage $k \in \{0, 1\}$, $k \leq \gamma$, to maximize V^P , and the daughter's abortion choice is constrained to $A \leq \gamma$. All other elements of the game—the pre-committed allowance, preferences, budget sets, and tie-breaking conventions—are unchanged. Note that under a consent regime every abortion is disclosed ($A = 1$ implies $R = 1$), so the knowledge and action benchmarks of Section 6.1 coincide; the wedge between them is created by the hidden route that consent removal opens.

Parents' authorization decision compares three continuation payoffs. Fix $S = 0$ and $T \equiv T_{\text{NP}}^*$, and consider a pregnant daughter who has disclosed and would abort if permitted. Refusing ($\gamma = 0$) yields $u(Y^P - T) + \alpha^P v$: no abortion. Authorizing with coverage ($\gamma = k = 1$) yields $u(Y^P - T - p) + \alpha^P v - \xi^P$, provided the daughter then aborts. Authorizing without coverage ($\gamma = 1, k = 0$) yields $u(Y^P - T) + \alpha^P v_p - \xi^P$ when the daughter self-funds, and coincides with refusal when she cannot. Comparing refusal with the better of the two authorizations gives a single authorization threshold on each region:

$$\tilde{\xi}^P \equiv \max \{ \bar{\xi}^P, \alpha^P \xi_3^D \},$$

where $\bar{\xi}^P = u(Y^P - T - p) - u(Y^P - T)$ is the coverage threshold and $\alpha^P \xi_3^D = \alpha^P (v_p - v)$ is the parents' altruistic valuation of letting the daughter pay for the procedure herself. In the canonical case, $\alpha^P \xi_3^D > \bar{\xi}^P$, so $\tilde{\xi}^P = \alpha^P \xi_3^D \in (\bar{\xi}^P, 0)$.

Proposition 8 (Consent removal and the subsidy reach different populations). *Fix $S = 0$ and the canonical coverage case, and let $B \equiv \{\xi^D \leq \xi_3^D\} \times \{\xi^P > \tilde{\xi}^P\}$.*

- (a) Outcomes under consent. Under the consent regime, the daughter aborts iff either $\xi_3^D < \xi^D \leq 0$ and $\xi^P \leq \bar{\xi}^P$ (a covered, authorized abortion) or $\xi^D \leq \xi_3^D$ and $\xi^P \leq \tilde{\xi}^P$ (an authorized abortion, self-funded in the canonical case). Every abortion is disclosed.
- (b) Consent removal. Removing the consent requirement (moving to the model of Section 4, still at $S = 0$) changes the abortion outcome exactly on B : daughters with strong abortion preferences ($\xi^D \leq \xi_3^D$, so they can and would pay) and strongly opposed parents ($\xi^P > \tilde{\xi}^P$), who were legally blocked, now obtain self-funded abortions.
- (c) Disjoint populations. The subsidy (the move $S : 0 \rightarrow 1$ in the no-consent world) changes the abortion outcome exactly on the complier rectangle $(\xi_3^D, 0] \times (\bar{\xi}^P, \infty)$, which is disjoint from B . Consent removal reaches daughters who can pay but lack permission; the subsidy reaches daughters who cannot pay. Neither policy affects the other's margin.
- (d) Welfare of consent removal. On B , consent removal changes family welfare by

$$\Delta W = \begin{cases} \xi_3^D - \xi^D - \min\{\xi^P, 0\} \geq 0, & \text{under the knowledge benchmark,} \\ \xi_3^D - \xi^D - \xi^P, & \text{under the action benchmark,} \end{cases}$$

and $\Delta W = 0$ elsewhere. The daughter gains her revealed-preference surplus net of the price p she now pays, $\xi_3^D - \xi^D \geq 0$ (the term $\xi_3^D = v_p - v < 0$); as with the subsidy, the approving parents she tells ($\xi^P \leq 0$) add their gain $-\xi^P \geq 0$ under the knowledge benchmark, while the opposed ones she hides from ($\xi^P > 0$) add nothing. Under the action benchmark the sign flips across $\xi^P = \xi_3^D - \xi^D$, and the same non-identification as in Corollary 1 applies: behavior bounds ξ^P on B only from below.

Proof. Part 8(a). A daughter with $\xi^D > 0$ does not abort in any branch (Lemma 1(a)), so consider $\xi^D \leq 0$. Since hiding cannot produce an abortion, her abortion prospects require disclosure; disclosure is costless, and hiding yields v , so she discloses whenever some authorized continuation weakly improves on v (tie-breaking as in Appendix E.4). On the region $\xi_3^D < \xi^D \leq 0$, she aborts under authorization only if covered ($v - \xi^D \geq v$ iff $\xi^D \leq 0$; self-funding is not profitable, $v_p - \xi^D < v$, by the definition of ξ_3^D). Parents therefore compare refusal, $u(Y^P - T) + \alpha^P v$, with authorization plus coverage, $u(Y^P - T - p) + \alpha^P v - \xi^P$: they authorize iff $\xi^P \leq \bar{\xi}^P$ (Lemma 2; authorization without coverage induces no abortion and duplicates refusal). On the region $\xi^D \leq \xi_3^D$, the daughter aborts under any authorization (self-funding if uncovered, since $v_p - \xi^D \geq v$). Parents' payoffs are: refusal, $u(Y^P - T) + \alpha^P v$; authorization without coverage, $u(Y^P - T) + \alpha^P v_p - \xi^P$; authorization with coverage, $u(Y^P - T - p) + \alpha^P v - \xi^P$. In the canonical case the second dominates the third (condition (22) fails), and it beats refusal iff $\xi^P \leq \alpha^P(v_p - v) = \alpha^P \xi_3^D = \tilde{\xi}^P$. Every abortion so obtained follows disclosure, $R = 1$.

Part 8(b). In the no-consent model at $S = 0$, the abortion set is $\{\xi_3^D < \xi^D \leq 0, \xi^P \leq \bar{\xi}^P\} \cup \{\xi^D \leq \underline{\xi}_3^D\}$ (Appendix F.6). Subtracting the consent-regime abortion set of part 8(a) leaves exactly $\{\xi^D \leq \underline{\xi}_3^D, \xi^P > \tilde{\xi}^P\} = B$. On B the no-consent outcome is a self-funded abortion (Lemma 7(ii.c)), disclosed iff $\xi^P \leq 0$; since $\tilde{\xi}^P < 0$ in the canonical case, B contains both disclosing ($\tilde{\xi}^P < \xi^P \leq 0$) and concealing ($\xi^P > 0$) families.


Part 8(c). The subsidy's margin is the complier rectangle (Appendix F.6), whose ξ^D -projection is $(\underline{\xi}_3^D, 0]$; the ξ^D -projection of B is $(-\infty, \underline{\xi}_3^D]$. The two sets are disjoint. A complier family is unaffected by consent removal (she cannot fund the procedure either way at $S = 0$), and a B family is unaffected by the subsidy in the no-consent world (she aborts at $S = 0$ already: an always-taker).

Part 8(d). On B , the consent-regime outcome is $A = 0$, daughter's value v , $C^P = Y^P - T$; the no-consent outcome is a self-funded abortion (daughter's value $v_p - \xi^D$, $C^P = Y^P - T$) disclosed iff $\xi^P \leq 0$. The daughter's change is $(v_p - \xi^D) - v = \underline{\xi}_3^D - \xi^D \geq 0$ on B (since $\xi^D \leq \underline{\xi}_3^D$): her revealed-preference surplus net of the price p she now pays ($v_p - v = \underline{\xi}_3^D < 0$). The parents' consumption is unchanged. Under the action benchmark the term $-\xi^P A$ contributes $-\xi^P$, so $\Delta W = \underline{\xi}_3^D - \xi^D - \xi^P$, positive iff $\xi^P < \underline{\xi}_3^D - \xi^D$. Under the knowledge benchmark $-\xi^P AR$ contributes $-\xi^P \cdot \mathbf{1}[\xi^P \leq 0] = -\min\{\xi^P, 0\}$ (the abortion is disclosed only when $\xi^P \leq 0$, and $A = 0$ under consent), so $\Delta W = \underline{\xi}_3^D - \xi^D - \min\{\xi^P, 0\} \geq 0$. Off B , part 8(b) gives identical outcomes under both regimes, so all payoffs are unchanged. Behavior on B reveals only the threshold crossings $\xi^D \leq \underline{\xi}_3^D$ and $\xi^P > \tilde{\xi}^P$, with ξ^P unbounded above, by the same argument as Corollary 1. $\square$


Remark M.1 (The knife-edge coverage case). When condition (22) holds, parents who authorize a self-funding daughter prefer to cover her, so $\tilde{\xi}^P = \bar{\xi}^P$, and the region on which consent removal changes the abortion outcome becomes $\{\xi^D \leq \underline{\xi}_3^D, \xi^P > \bar{\xi}^P\}$. After removal, those daughters disclose and are covered (Remark K.2), so the newly enabled abortions are disclosed and parent-funded: the parents pay p and bear ξ^P on a disclosed abortion, so $\Delta W = \bar{\xi}^P - \xi^D - \xi^P$ under *both* benchmarks. The populations remain disjoint from the subsidy's margin, and every qualitative conclusion of Proposition 8 is unchanged.

Remark M.2 (A subsidy does not undo a consent requirement). The wedge runs in both directions. Keep the consent regime and introduce the subsidy ($S : 0 \rightarrow 1$). Upon authorization the procedure is free, so coverage is costless and the parents' authorization comparison reduces to $-\xi^P$ against 0: they authorize iff $\xi^P \leq 0$. Under consent-plus-subsidy, then, a daughter aborts iff $\xi^D \leq 0$ and $\xi^P \leq 0$ —whereas the subsidy without a consent requirement delivers an abortion to every daughter with $\xi^D \leq 0$ (Lemma 1(a)). Families with $\xi^D \leq 0$ and $\xi^P > 0$ remain blocked by the legal veto no matter how generous the subsidy.

Figure M.1: Parental consent versus the subsidy



(a) Behavioral margins of the two policies



(b) Sign of ΔW from consent removal (action benchmark)

Notes: Both panels use the attitude partition of Figure 3, Panel (c); axes are schematic. **Panel (a):** the region where removing a parental-consent requirement changes the abortion outcome (blue: self-funders with $\xi^P > \xi^P$, who can pay but were legally blocked; Proposition 8(b)) and the region where the subsidy changes it (orange: the complier rectangle, daughters who cannot pay and whose parents refuse; Section 4). The regions are disjoint: permission and money relax different constraints. Self-funders with $\xi^P \in (\xi^P, \xi^P]$ are authorized under consent (parents prefer letting them pay to blocking them), so consent removal does not change their outcome. **Panel (b):** sign of the welfare change ΔW of consent removal under the action benchmark (Proposition 8(d)): $\Delta W = \xi_3^D - \xi^D - \xi^P$ nets the daughter's gain, net of the price p she now pays, against the parents' distaste. It is positive below the dashed locus $\xi^P = \xi_3^D - \xi^D$, which enters the region at $(\xi^D, \xi^P) = (\xi_3^D, 0)$, and negative above it. Under the knowledge benchmark $\Delta W = \xi_3^D - \xi^D - \min\{\xi^P, 0\} \geq 0$ throughout the region, so the red region vanishes. The structure parallels the subsidy's complier margin (Figure 5), with one difference: consent removal is not a transfer—the newly enabled daughters pay for the procedure themselves, which is why the locus is shifted down by the price term ξ_3^D relative to $\xi^P = -\xi^D$.